

# CS4.301 Data & Applications

Ponnurangam Kumaraguru ("PK")  
#ProfGiri @ IIIT Hyderabad



pk.profgiri



/in/ponguru



@ponguru



Ponnurangam.kumaraguru

# Weak Entity Types

An entity that does not have a key attribute and that is identification-dependent on another entity type.

A weak entity must participate in an identifying relationship type with an owner or identifying entity type

Entities are identified by the combination of:

A partial key of the weak entity type

The particular entity they are related to in the identifying relationship type

## Example:

A DEPENDENT entity is identified by the dependent's first name, *and* the specific EMPLOYEE with whom the dependent is related

Name of DEPENDENT is the *partial key*

DEPENDENT is a *weak entity type*

EMPLOYEE is its identifying entity type via the identifying relationship type DEPENDENT\_OF

# Weak Entity Types

Concept	Symbol	Meaning
Strong Entity	Single rectangle	Independent existence
Weak Entity	Double rectangle	Depends on strong entity
Identifying Relationship	Double diamond	Defines weak entity's identity

# Attributes of Relationship types

A relationship type can have attributes:

For example, HoursPerWeek of WORKS\_ON

Its value for each relationship instance describes the number of hours per week that an EMPLOYEE works on a PROJECT.

A value of HoursPerWeek depends on a particular (employee, project) combination

Most relationship attributes are used with M:N relationships

A **relationship type** can have an **attribute** when that attribute **belongs to the association itself** — not to any of the entities individually.

Think of it as information that **only makes sense because of the relationship** between the entities.

# Attributes of Relationship types

Add attributes to a relationship type **only when**:

1. The information applies **to the link itself**, not to either entity.
2. It's **not derivable** from the entities' existing attributes.
3. It helps capture **real-world meaning** of that association.

# Attributes of Relationship types

When an ER relationship has attributes, those attributes are stored as columns in the relationship table (also called a junction or association table) that represents the relationship in the relational database.


Employee(Emp\_ID, Name)

Project(Proj\_ID, Title)

Works\_On(Emp\_ID, Proj\_ID, hours\_per\_week)

Relationship attributes become **columns in the relationship (junction) table** when the ER model is converted to the relational model.

# The (min,max) notation for relationship constraints




Read the min,max numbers next to the entity type and looking **away from** the entity type

**Figure 3.16**

The COMPANY conceptual schema in UML class diagram notation.

## UML class diagram for COMPANY database schema



# DBMS Programming Language Interfaces

Programmer interfaces for embedding DML in a programming languages:

**Embedded Approach:** e.g embedded SQL (for C, C++, etc.), SQLJ (for Java)

**Procedure Call Approach:** e.g. JDBC for Java, ODBC (Open Database Connectivity) for other programming languages as API's (application programming interfaces)

**Database Programming Language Approach:** e.g. ORACLE has PL/SQL, a programming language based on SQL; language incorporates SQL and its data types as integral components

**Scripting Languages:** PHP (client-side scripting) and Python (server-side scripting) are used to write database programs.

# User-Friendly DBMS Interfaces

Menu-based (Web-based), popular for browsing on the web

Forms-based, designed for naïve users used to filling in entries on a form

Graphics-based

- Point and Click, Drag and Drop, etc.

- Specifying a query on a schema diagram

Natural language: requests in written English

Combinations of the above:

- For example, both menus and forms used extensively in Web database interfaces

# Database System Utilities

To perform certain functions such as:

Loading data stored in files into a database. Includes data conversion tools.

Backing up the database periodically on tape.

Reorganizing database file structures.

Performance monitoring utilities.

Report generation utilities.

Other functions, such as sorting, user monitoring, data compression, etc.

# The Relational Data Model and Relational Database Constraints

# Informal Definitions

## Key of a Relation:

Each row has a value of a data item (or set of items) that uniquely identifies that row in the table

Called the *key*

In the STUDENT table, SSN is the key

Sometimes row-ids or sequential numbers are assigned as keys to identify the rows in a table

Called *artificial key* or *surrogate key*

# Activity

Create 3 recursive relationship and show them in ER diagram

Create 3 different weak entities in 3 different mini-world and show them in ER diagram

Create 3 different sets of relationships in Infinium showing cardinality / min-max

# Formal Definitions - Schema

The **Schema** (or description) of a Relation:

Denoted by  $R(A_1, A_2, \dots, A_n)$

$R$  is the **name** of the relation

The **attributes** of the relation are  $A_1, A_2, \dots, A_n$

Example:

CUSTOMER (Cust-id, Cust-name, Address, Phone#)

CUSTOMER is the relation name

Defined over the four attributes: Cust-id, Cust-name, Address, Phone#

Each attribute has a **domain** or a set of valid values.

For example, the domain of Cust-id is 6 digit numbers.

# Formal Definitions - Tuple

A **tuple** is an ordered set of values (enclosed in angled brackets '< ... >')

Each value is derived from an appropriate *domain*.

A row in the CUSTOMER relation is a 4-tuple and would consist of four values, for example:

<632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">

This is called a 4-tuple as it has 4 values

A tuple (row) in the CUSTOMER relation.

A relation is a **set** of such tuples (rows)

# Formal Definitions - Domain

A **domain** has a logical definition:

Example: “USA\_phone\_numbers” are the set of 10 digit phone numbers valid in the U.S.

A domain also has a data-type or a format defined for it.

The USA\_phone\_numbers may have a format: (ddd)ddd-dddd where each d is a decimal digit.

Dates have various formats such as year, month, date formatted as yyyy-mm-dd, or as dd mm,yyyy etc.

The attribute name designates the role played by a domain in a relation:

Used to interpret the meaning of the data elements corresponding to that attribute

Example: The domain Date may be used to define two attributes named “Invoice-date” and “Payment-date” with different meanings

# Definition Summary

<u>Informal Terms</u>	<u>Formal Terms</u>
Table	Relation
Column Header	Attribute
All possible Column Values	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	State of the Relation

# Characteristics Of Relations

Ordering of tuples in a relation  $r(R)$ :

The tuples are *not considered to be ordered*, even though they appear to be in the tabular form.

Ordering of attributes in a relation schema  $R$  (and of values within each tuple):

We will consider the attributes in  $R(A_1, A_2, \dots, A_n)$  and the values in  $t = <v_1, v_2, \dots, v_n>$  to be ordered .

(However, a more general alternative definition of relation does not require this ordering.  
It includes both the name and the value for each of the attributes ).

Example:  $t = \{ <\text{name}, \text{"John"} \gt, <\text{SSN}, 123456789 \gt \}$

This representation may be called as “self-describing”.

# This Lecture

# Characteristics Of Relations

## Values in a tuple:

All values are considered atomic (indivisible).

Each value in a tuple must be from the domain of the attribute for that column

If tuple  $t = \langle v_1, v_2, \dots, v_n \rangle$  is a tuple (row) in the relation state  $r$  of  $R(A_1, A_2, \dots, A_n)$

Then each  $v_i$  must be a value from  $\text{dom}(A_i)$

A special **null** value is used to represent values that are unknown or not available or inapplicable in certain tuples.

# Characteristics Of Relations

Notation:

We refer to **component values** of a tuple  $t$  by:

$t[A_i]$  or  $t.A_i$

This is the value  $v_i$  of attribute  $A_i$  for tuple  $t$

Similarly,  $t[A_u, A_v, \dots, A_w]$  refers to the subtuple of  $t$  containing the values of attributes  $A_u, A_v, \dots, A_w$ , respectively in  $t$

# CONSTRAINTS

Constraints determine which values are permissible and which are not in the database. They are of three main types:

1. **Inherent or Implicit Constraints:** These are based on the data model itself. (E.g., relational model does not have tuples to be duplicates)
2. **Schema-based or Explicit Constraints:** They are expressed in the schema by using the facilities provided by the model. (E.g., max. cardinality ratio constraint in the ER model)
3. **Application based or semantic constraints:** These are beyond the expressive power of the model and must be specified and enforced by the application programs.

# Inherent or Implicit Constraints

Constraints that exist naturally due to the meaning of data or real-world logic, not explicitly declared in the schema

Inherent constraints are enforced by logic, not the DBMS—unless we model them

Example	Explanation
A person's date of birth cannot be in the future	True by nature of time
A student's age must be positive	Logical and real-world rule
Arrival time > Departure time for the same flight	Temporal reality
A bridge can't connect to itself	Semantic impossibility

# Schema-based or Explicit Constraints

Constraints that are formally declared in the database schema using SQL. These are explicitly stated in the schema and automatically checked by the DBMS.

Type	Example	Enforced by
Domain Constraint	age INT CHECK (age >= 18)	Column definition
Key Constraint	PRIMARY KEY (student_id)	Table-level rule
Entity Integrity	Primary key cannot be NULL	DBMS automatically
Referential Integrity	FOREIGN KEY (dept_id) REFERENCES Department(dept_id)	Enforced by DBMS
Unique Constraint	UNIQUE (email)	Ensures no duplicates

# Application based or semantic constraints

Constraints that depend on business logic, process, or context — not easily enforced in schema. These constraints live in the application layer, triggers, or stored procedures — not the schema itself.

Example	Explanation
A professor cannot teach more than 3 courses per semester	Depends on institutional policy
Discount applies only if order value > ₹5000	Business rule enforced by application code
Loan approval requires credit score > 700	External data & logic-based
Employee salary increase ≤ 20% per year	Policy constraint, not structural
Two users cannot book the same seat	Controlled through transactions/application logic

# Constraints summary

Type	Defined By	Example	Where Enforced
Inherent / Implicit	Nature of data, logic	Age > 0, DoB < Today	Reality / Semantics
Schema-based / Explicit	SQL schema definition	CHECK, PK, FK	DBMS
Application / Semantic	Business rules, process logic	Max 3 courses / semester	Application layer

# Relational Integrity Constraints

Constraints are **conditions** that must hold on **all** valid relation states.

There are three *main types* of (explicit schema-based) constraints that can be expressed in the relational model:

**Key** constraints

**Entity integrity** constraints

**Referential integrity** constraints

Another schema-based constraint is the **domain** constraint

Every value in a tuple must be from the *domain of its attribute* (or it could be **null**, if allowed for that attribute)

# Key Constraints

## Superkey of R:

Super Key is an attribute (or set of attributes) that is used to uniquely identifies all attributes in a relation.

Is a set of attributes SK of R with the following condition:

No two tuples in any valid relation state  $r(R)$  will have the same value for SK

That is, for any distinct tuples  $t_1$  and  $t_2$  in  $r(R)$ ,  $t_1[SK] \neq t_2[SK]$

This condition must hold in *any valid state*  $r(R)$

## Key of R:

A "minimal" superkey

A candidate key (or simply a key) is a set of one or more attributes that uniquely identify each tuple in a relation (or table)

# Key Constraints (continued)

Example: Student{ID, First\_name, Last\_name, Age, Sex, Phone\_no}

Two candidate keys: ID & {First\_name, Last\_name, Age, Sex, Phone\_no}

Consider the STUDENT relation schema:

Attribute Ssn is a key, because no 2 students tuples can have the same Ssn

Any set of attributes that include Ssn, like {Ssn, Name, Age} is a superkey it uniquely identifies all attributes in a relation

In general:

Any key is a *superkey* (but not vice versa)

Any set of attributes that *includes a key* is a *superkey*

A *minimal* superkey is also a key

# Example

Candidate keys?

RollNo	Email	Aadhaar	Name	Dept
101	alex@iiith.ac.in	1234	Alex	CSE
102	bria@iiith.ac.in	5678	Bria	ECE
103	chay@iiith.ac.in	9999	Chay	CSE

# Example

Candidate keys? {RollNo}, {Email}, {Aadhaar}

Why name can't be a candidate key?

RollNo	Email	Aadhaar	Name	Dept
101	alex@iiith.ac.in	1234	Alex	CSE
102	bria@iiith.ac.in	5678	Bria	ECE
103	chay@iiith.ac.in	9999	Chay	CSE

# Candidate key

Uniqueness: No two tuples (rows) in the table can have the same key value.

Minimality: No subset of that key can uniquely identify the tuple.

RollNo	Email	Aadhaar	Name	Dept
101	alex@iiith.ac.in	1234	Alex	CSE
102	bria@iiith.ac.in	5678	Bria	ECE
103	chay@iiith.ac.in	9999	Chay	CSE
104	alex2@iiith.ac.in	4321	Alex	ECE

# Candidate key

Feature	Superkey	Candidate Key
Definition	Any set of attributes that uniquely identifies tuples	Minimal set of attributes that uniquely identifies tuples
Uniqueness	Must hold	Must hold
Minimality	Not required	Must hold
Count per Table	Many	Subset of superkeys
Example	{RollNo, Name}	{RollNo}, {Email}

# Candidate key & Super key difference

Super Key	Candidate Key
Super Key is an attribute (or set of attributes) that is used to uniquely identifies all attributes in a relation.	Candidate Key is a subset of a super key.
All super keys can't be candidate keys.	But all candidate keys are super keys.
Various super keys together makes the criteria to select the candidate keys.	Various candidate keys together makes the criteria to select the primary keys.
In a relation, number of super keys is more than number of candidate keys.	While in a relation, number of candidate keys are less than number of super keys.
Super key attributes can contain NULL values.	Candidate key attributes can also contain NULL values.

# Key Constraints (continued)

If a relation has several **candidate keys**, one is chosen arbitrarily to be the **primary key**.

The primary key attributes are underlined.

Example: Consider the CAR relation schema:

License\_number & Engine\_serial\_number

We chose Engine\_serial\_number as the primary key

The primary key value is used to *uniquely identify* each tuple in a relation

Provides the tuple identity

Also used to *reference* the tuple from another tuple

General rule: Choose as primary key the smallest of the candidate keys (in terms of size)

Not always applicable – choice is sometimes subjective

# CAR table with two candidate keys – LicenseNumber chosen as Primary Key

**CAR**

<u>License_number</u>	Engine_serial_number	Make	Model	Year
Texas ABC-739	A69352	Ford	Mustang	02
Florida TVP-347	B43696	Oldsmobile	Cutlass	05
New York MPO-22	X83554	Oldsmobile	Delta	01
California 432-TFY	C43742	Mercedes	190-D	99
California RSK-629	Y82935	Toyota	Camry	04
Texas RSK-629	U028365	Jaguar	XJS	04

**Figure 5.4**

The CAR relation, with two candidate keys: License\_number and Engine\_serial\_number.

# Super & Primary key difference

S.NO	Super Key	Primary Key
1.	Super Key is an attribute (or set of attributes) that is used to uniquely identifies all attributes in a relation.	Primary Key is a minimal set of attribute (or set of attributes) that is used to uniquely identifies all attributes in a relation.
2.	All super keys can't be primary keys.	Primary key is a minimal super key.
3.	Various super keys together makes the criteria to select the candidate keys.	We can choose any of the minimal candidate key to be a primary key.
4.	In a relation, number of super keys are more than number of primary keys.	While in a relation, number of primary keys are less than number of super keys.
5.	Super key's attributes can contain NULL values.	Primary key's attributes cannot contain NULL values.

# Relational Database Schema

## **Relational Database Schema:**

A set  $S$  of relation schemas that belong to the same database.

$S$  is the name of the whole **database schema**

$S = \{R_1, R_2, \dots, R_n\}$  and a set  $IC$  of integrity constraints.

$R_1, R_2, \dots, R_n$  are the names of the individual **relation schemas** within the database  $S$

Following slide shows a COMPANY database schema with 6 relation schemas

## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

## DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

## PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

## WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

## DEPENDENT

<u>Essn</u>	Dependent_name	Sex	Bdate	Relationship
-------------	----------------	-----	-------	--------------

**Figure 5.5**

Schema diagram for  
the COMPANY  
relational database  
schema.

## EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	-----	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
-------	---------	---------	----------------

## DEPT\_LOCATIONS

Dnumber	Dlocation
---------	-----------

## PROJECT


Pname	Pnumber	Plocation	Dnum
-------	---------	-----------	------

## WORKS\_ON

Essn	Pno	Hours
------	-----	-------

## DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
------	----------------	-----	-------	--------------



# Relational Database State

A **relational database state** DB of  $S$  is a set of relation states  $\text{DB} = \{r_1, r_2, \dots, r_m\}$  such that each  $r_i$  is a state of  $R_i$  and such that the  $r_i$  relation states satisfy the integrity constraints specified in IC.

A relational database *state* is sometimes called a relational database *snapshot* or *instance*.

# Populated database state

Each *relation* will have many tuples in its current relation state

The *relational database state* is a union of all the individual relation states

Whenever the database is changed, a new state arises

Basic operations for changing the database:

- INSERT a new tuple in a relation

- DELETE an existing tuple from a relation

- MODIFY an attribute of an existing tuple

**Figure 5.6**

One possible database state for the COMPANY relational database schema.

**EMPLOYEE**

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

**DEPARTMENT**

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

**DEPT\_LOCATIONS**

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

**WORKS\_ON**

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

**PROJECT**

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

**DEPENDENT**

Essn	Dependent_name	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

# Populated DB state for COMPANY

# Entity Integrity

## Entity Integrity:

The *primary key attributes* PK of each relation schema R in S cannot have null values in any tuple of  $r(R)$ .

This is because primary key values are used to *identify* the individual tuples.

$t[PK] \neq \text{null}$  for any tuple  $t$  in  $r(R)$

If PK has several attributes, null is not allowed in any of these attributes

Note: Other attributes of R may be constrained to disallow null values, even though they are not members of the primary key.

# Referential Integrity

Tuples in the **referencing relation** R1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the **referenced relation** R2.

A tuple t1 in R1 is said to **reference** a tuple t2 in R2 if  $t1[FK] = t2[PK]$ .

A referential integrity constraint can be displayed in a relational database schema as a directed arc from R1.FK to R2.PK

# Referential Integrity (or foreign key) Constraint

## Statement of the constraint

The value in the foreign key column (or columns) FK of the **referencing relation R1** can be **either**:

- (1) a value of an existing primary key value of a corresponding primary key PK in the **referenced relation R2**, or
- (2) a **null**.

In case (2), the FK in R1 should **not** be a part of its own primary key.

# Displaying a relational database schema and its constraints

Each relation schema can be displayed as a row of attribute names

The name of the relation is written above the attribute names

The primary key attribute (or attributes) will be underlined

A foreign key (referential integrity) constraints is displayed as a directed arc (arrow) from the foreign key attributes to the referenced table

Can also point the the primary key of the referenced relation for clarity

Next slide shows the COMPANY **relational schema diagram with referential integrity constraints**

# Referential integrity constraints for Company

**Figure 5.7**

Referential integrity constraints displayed on the COMPANY relational database schema.

## EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	-----	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
-------	---------	---------	----------------

## DEPT\_LOCATIONS

Dnumber	Dlocation
---------	-----------

## PROJECT

Pname	Pnumber	Plocation	Dnum
-------	---------	-----------	------

## WORKS\_ON

Essn	Pno	Hours
------	-----	-------

## DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
------	----------------	-----	-------	--------------

# Activity: Infinium mini-world

Write down 3 examples for Inherent constraints

Write down 3 examples for Explicit constraints

Write down 3 examples for Application constraints

Generate a Relation, Super keys, Candidate keys, Primary keys

Write down at least 2 integrity constraints, and 2 referential integrity constraints

With multiple entities show at least 2 referential integrity constraints through relational schema

# Bibliography / Acknowledgements

Instructor materials from Elmasri & Navathe 7e



pk.profgiri



Ponnurangam.kumaraguru



/in/ponguru



ponguru



pk.guru@iiit.ac.in

Thank you  
for attending  
the class!!!