

CS4.301 Data & Applications

Ponnurangam Kumaraguru ("PK")
#ProfGiri @ IIIT Hyderabad



pk.profgiri



/in/ponguru



@ponguru



Ponnurangam.kumaraguru

```
mysql> SELECT
    -> D.Dname, E.Lname, E.Fname, P.Pname
    -> FROM
    -> DEPARTMENT AS D, EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P
    -> WHERE
    -> D.Dnumber = E.Dno AND E.Ssn = W.Essn AND W.Pno = P.Pnumber
    -> ORDER BY
    -> D.Dname, E.Lname, E.Fname;
+-----+-----+-----+-----+
| Dname      | Lname   | Fname   | Pname    |
+-----+-----+-----+-----+
| Administration | Jabbar | Ahmad | Computerization |
| Administration | Jabbar | Ahmad | Newbenefits |
| Administration | Wallace | Jennifer | Reorganization |
| Administration | Wallace | Jennifer | Newbenefits |
| Administration | Zelaya | Alicia | Computerization |
| Administration | Zelaya | Alicia | Newbenefits |
| Headquarters | Borg | James | Reorganization |
| Research | English | Joyce | ProductX |
| Research | English | Joyce | ProductY |
| Research | Narayan | Ramesh | ProductZ |
| Research | Smith | John | ProductX |
| Research | Smith | John | ProductY |
| Research | Wong | Franklin | ProductY |
| Research | Wong | Franklin | ProductZ |
| Research | Wong | Franklin | Computerization |
| Research | Wong | Franklin | Reorganization |
+-----+-----+-----+-----+
16 rows in set (0.02 sec)
```

The INSERT Command

Specify the relation name and a list of values for the tuple. All values including nulls are supplied.

```
U1:    INSERT INTO  EMPLOYEE  
          VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
                           Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

```
mysql> INSERT INTO EMPLOYEE  
      -> VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98 Oak Forest, Katy, TX', 'M', 37000, '653298653', 4);  
Query OK, 1 row affected (0.01 sec)
```

```
mysql> SELECT * FROM EMPLOYEE;  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
| Fname | Minit | Lname | Ssn   | Bdate  | Address           | Sex  | Salary | Super_ssn | Dno |  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
| John  | B     | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston TX | M   | 30000 | 333445555 | 5  |  
| Franklin | T    | Wong  | 333445555 | 1965-12-08 | 638 Voss, Houston TX  | M   | 40000 | 888665555 | 5  |  
| Joyce  | A    | English | 453453453 | 1972-07-31 | 5631 Rice, Houston TX | F   | 25000 | 333445555 | 5  |  
| Richard | K    | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M   | 37000 | 653298653 | 4  |  
| Ramesh  | K    | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble TX | M   | 38000 | 333445555 | 5  |  
| James   | E    | Borg   | 888665555 | 1937-11-10 | 450 Stone, Houston TX | M   | 55000 | NULL   | 1   |  
| Jennifer | S    | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire TX | F   | 43000 | 888665555 | 4  |  
| Ahmad   | V    | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston TX | M   | 25000 | 987654321 | 4  |  
| Alicia  | J    | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring TX | F   | 25000 | 987654321 | 4  |  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
9 rows in set (0.00 sec)
```

```
mysql> CREATE TABLE WORKS_ON_INFO (
    ->     Emp_name VARCHAR(150),
    ->     Proj_name VARCHAR(150),
    ->     Hours_per_week DECIMAL(3, 1)
    -> );
Query OK, 0 rows affected (0.02 sec)

mysql> INSERT INTO WORKS_ON_INFO ( Emp_name, Proj_name, Hours_per_week )
    -> SELECT E.Lname, P.Pname, W.Hours
    -> FROM PROJECT P, WORKS_ON W, EMPLOYEE E
    -> WHERE P.Pnumber = W.Pno AND W.Essn = E.Ssn;
Query OK, 16 rows affected (0.02 sec)
Records: 16  Duplicates: 0  Warnings: 0

mysql> SELECT * FROM WORKS_ON_INFO;
+-----+-----+-----+
| Emp_name | Proj_name      | Hours_per_week |
+-----+-----+-----+
| Wong     | Computerization |      10.0      |
| Jabbar   | Computerization |      35.0      |
| Zelaya   | Computerization |      10.0      |
| Wallace  | Newbenefits    |      20.0      |
| Jabbar   | Newbenefits    |       5.0      |
| Zelaya   | Newbenefits    |      30.0      |
| Smith    | ProductX       |      32.5      |
| English  | ProductX       |      20.0      |
| Smith    | ProductY       |       7.5      |
| Wong     | ProductY       |      10.0      |
| English  | ProductY       |      20.0      |
| Wong     | ProductZ       |      10.0      |
| Narayan  | ProductZ       |      40.0      |
| Wong     | Reorganization  |      10.0      |
| Borg     | Reorganization  |      16.0      |
| Wallace  | Reorganization  |      15.0      |
+-----+-----+-----+
16 rows in set (0.00 sec)
```

The DELETE Command

Removes tuples from a relation

Includes a WHERE clause to select the tuples to be deleted. The number of tuples deleted will vary.

| | | |
|-------------|--------------------|------------------|
| U4A: | DELETE FROM | EMPLOYEE |
| | WHERE | Lname='Brown'; |
| U4B: | DELETE FROM | EMPLOYEE |
| | WHERE | Ssn='123456789'; |
| U4C: | DELETE FROM | EMPLOYEE |
| | WHERE | Dno=5; |
| U4D: | DELETE FROM | EMPLOYEE; |

```
mysql> UPDATE PROJECT  
    -> SET PLOCATION = 'Bellaire', DNUM = 5  
    -> WHERE PNUMBER = 10;
```

```
Query OK, 1 row affected (0.01 sec)  
Rows matched: 1  Changed: 1  Warnings: 0
```

```
mysql> SELECT * FROM PROJECT;
```

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Bellaire | 5 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

6 rows in set (0.00 sec)

UPDATE (contd.)

Example: Give all employees in the 'Research' department a 10% raise in salary.

```
mysql> UPDATE EMPLOYEE  
-> SET SALARY = SALARY * 1.1  
-> WHERE DNO IN (  
->     SELECT DNUMBER  
->     FROM DEPARTMENT  
->     WHERE DNAME = 'Research'  
-> );
```

```
Query OK, 4 rows affected (0.02 sec)  
Rows matched: 4  Changed: 4  Warnings: 0
```

```
mysql> SELECT * FROM EMPLOYEE;
```

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|-------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston TX | M | 33000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1965-12-08 | 638 Voss, Houston TX | M | 44000 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston TX | F | 27500 | 333445555 | 5 |
| Richard | K | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M | 37000 | 653298653 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble TX | M | 41800 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston TX | M | 55000 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire TX | F | 43000 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston TX | M | 25000 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring TX | F | 25000 | 987654321 | 4 |

```
9 rows in set (0.00 sec)
```

Specifying Joined Tables in the FROM Clause of SQL

Joined table

Permits users to specify a table resulting from a join operation in the FROM clause of a query

The FROM clause in Q1A

Contains a single joined table. JOIN may also be called INNER JOIN

```
Select fname, lname, address  
from (employee join department  
on dno=dnumber) where  
dname='research';
```

```
mysql> Select fname, lname, address from (employee  
join department on dno=dnumber) where dname='resear  
ch';  
+-----+-----+-----+  
| fname | lname | address |  
+-----+-----+-----+  
| John  | Smith | 731 Fondren, Houston TX |  
| Franklin | Wong | 638 Voss, Houston TX |  
| Joyce | English | 5631 Rice, Houston TX |  
| Ramesh | Narayan | 975 Fire Oak, Humble TX |  
+-----+-----+-----+  
4 rows in set (0.04 sec)
```

EMPLOYEE

| emp_id | fname | lname | dno |
|--------|-------|-------|-----|
| 1 | John | Doe | 10 |
| 2 | Alice | Lee | 20 |
| 3 | Bob | Kim | 10 |

DEPARTMENT

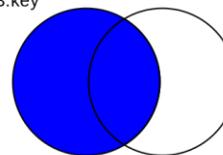
| dno | dname | location |
|-----|----------|----------|
| 10 | Research | Delhi |
| 20 | HR | Mumbai |
| 30 | Finance | Chennai |

```
SELECT fname, lname, dname, location  
FROM employee NATURAL JOIN department;
```

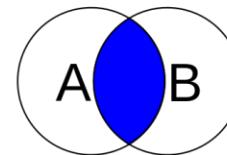
| fname | lname | dname | location |
|-------|-------|----------|----------|
| John | Doe | Research | Delhi |
| Bob | Kim | Research | Delhi |
| Alice | Lee | HR | Mumbai |

Joins differences

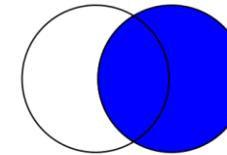
```
SELECT <fields>
  FROM TableA A
  LEFT JOIN TableB B
    ON A.key = B.key
```



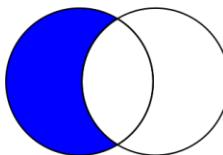
```
SELECT <fields>
  FROM TableA A
 INNER JOIN TableB B
    ON A.key = B.key
```



```
SELECT <fields>
  FROM TableA A
 RIGHT JOIN TableB B
    ON A.key = B.key
```

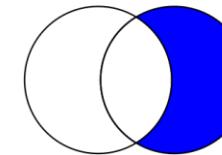


```
SELECT <fields>
  FROM TableA A
  LEFT JOIN TableB B
    ON A.key = B.key
 WHERE B.key IS NULL
```

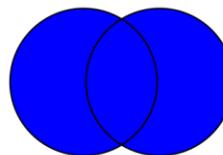


SQL JOINS

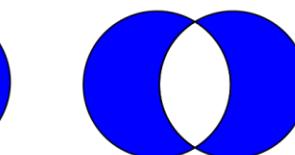
```
SELECT <fields>
  FROM TableA A
 RIGHT JOIN TableB B
    ON A.key = B.key
 WHERE A.key IS NULL
```



```
SELECT <fields>
  FROM TableA A
 FULL OUTER JOIN TableB B
    ON A.key = B.key
```



```
SELECT <fields>
  FROM TableA A
 FULL OUTER JOIN TableB B
    ON A.key = B.key
 WHERE A.key IS NULL
   OR B.key IS NULL
```



This work is licensed under a Creative Commons Attribution 3.0 Unported License.
Author: <http://commons.wikimedia.org/wiki/User:Arbeck>

Multiway JOIN in the FROM clause

Can nest JOIN specifications for a multiway join:

```
SELECT Pnumber, Dnum, Lname,  
Address, Bdate FROM ((PROJECT  
JOIN DEPARTMENT ON  
Dnum=Dnumber) JOIN EMPLOYEE  
ON Mgr_ssN=Ssn) WHERE  
Plocation='Stafford';
```

```
mysql> SELECT Pnumber, Dnum, Lname, Address, Bdate FROM ((PROJECT JOIN DEPARTMENT ON Dnum=Dnumber) JOIN EMPLOYEE ON Mgr_ssN=Ssn) WHERE Plocation='Stafford';  
+-----+-----+-----+-----+-----+  
| Pnumber | Dnum | Lname | Address | Bdate |  
+-----+-----+-----+-----+-----+  
| 10 | 4 | Wallace | 291 Berry, Bellaire TX | 1941-06-20 |  
| 30 | 4 | Wallace | 291 Berry, Bellaire TX | 1941-06-20 |  
+-----+-----+-----+-----+-----+  
2 rows in set (0.02 sec)
```

Activity

Try all the queries all yourself if you have not tried it

Write 3 join statements using any of the Company DB

Submit the query and results

Try the Multiway JOIN statement with any of the Company DB

Submit the query and results

This Lecture

CHAPTER 14

Basics of Functional Dependencies and Normalization for Relational Databases

Informal Design Guidelines for Relational Databases

We first discuss informal guidelines for good relational design

Then we discuss formal concepts of functional dependencies and normal forms

- 1NF (First Normal Form)
- 2NF (Second Normal Form)
- 3NF (Third Normal Form)
- BCNF (Boyce-Codd Normal Form)

1.1 Semantics of the Relational Attributes must be clear

GUIDELINE 1: Informally, each tuple in a relation should represent one entity or relationship instance. (Applies to individual relations and their attributes).

Attributes of different entities (EMPLOYEEs, DEPARTMENTs, PROJECTs) should not be mixed in the same relation

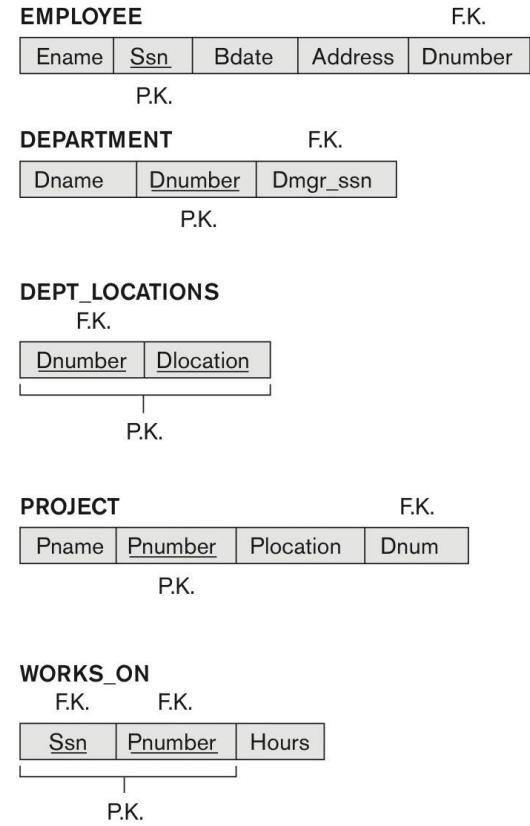
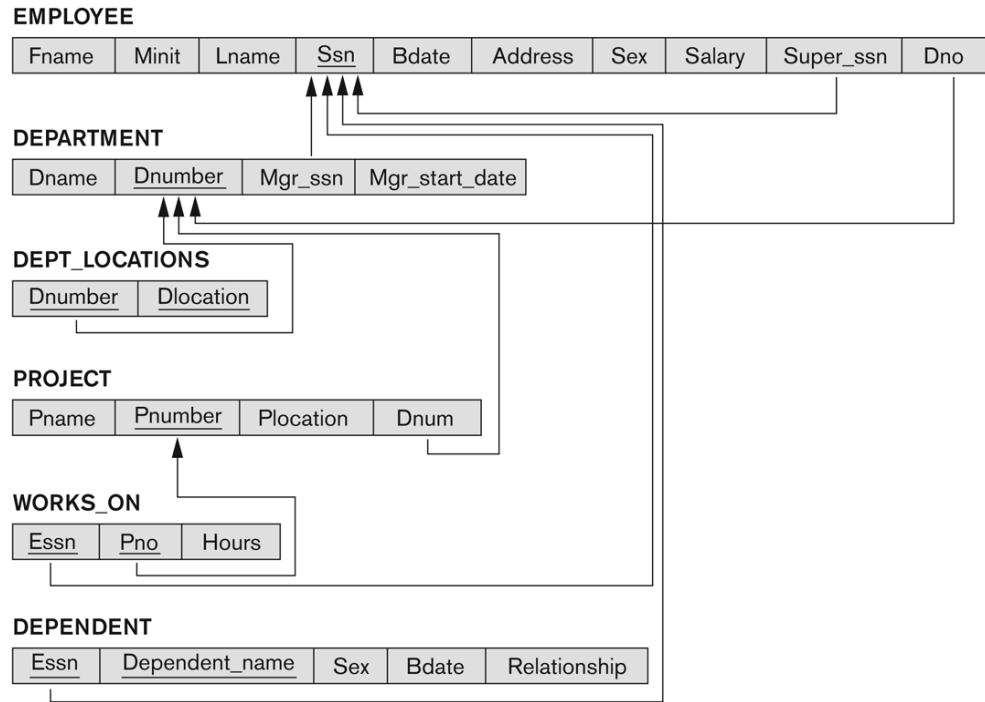
Only foreign keys should be used to refer to other entities

Bottom Line: *Design a schema that can be explained easily relation by relation. The semantics of attributes should be easy to interpret.*

Figure 14.1 A simplified COMPANY relational database schema

Figure 5.7

Referential integrity constraints displayed on the COMPANY relational database schema.



EMPLOYEE

| Ename | Ssn | Bdate | Address | Dnumber |
|----------------------|-----------|------------|--------------------------|---------|
| Smith, John B. | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | 5 |
| Wong, Franklin T. | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | 5 |
| Zelaya, Alicia J. | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX | 4 |
| Wallace, Jennifer S. | 987654321 | 1941-06-20 | 291Berry, Bellaire, TX | 4 |
| Narayan, Ramesh K. | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | 5 |
| English, Joyce A. | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | 5 |
| Jabbar, Ahmad V. | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | 4 |
| Borg, James E. | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | 1 |

DEPARTMENT

| Dname | Dnumber | Dmgr_ssn |
|----------------|---------|-----------|
| Research | 5 | 333445555 |
| Administration | 4 | 987654321 |
| Headquarters | 1 | 888665555 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| Ssn | Pnumber | Hours |
|-----------|---------|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | Null |

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

(a)

EMP_DEPT

| Ename | Ssn | Bdate | Address | Dnumber | Dname | Dmgr_ssn |
|-------|-----|-------|---------|---------|-------|----------|
| | | | | | | |

```
graph TD; Ssn[ ] --> SsnCell[ ]; Bdate[ ] --> BdateCell[ ]; Address[ ] --> AddressCell[ ]; Dnumber[ ] --> DnumberCell[ ]; Dname[ ] --> DnameCell[ ]; Dmgr_ssn[ ] --> Dmgr_ssnCell[ ]
```

Any concerns here?

(b)

EMP_PROJ

| Ssn | Pnumber | Hours | Ename | Pname | Plocation |
|-----|---------|-------|-------|-------|-----------|
| FD1 | | | | | |
| FD2 | | | | | |
| FD3 | | | | | |

```
graph TD; Ssn[FD1] --> SsnCell[ ]; Pnumber[FD1] --> PnumberCell[ ]; Hours[FD2] --> HoursCell[ ]; Ename[FD2] --> EnameCell[ ]; Pname[FD2] --> PnameCell[ ]; Plocation[FD2] --> PlocationCell[ ]; FD3[ ] --> PnumberCell[ ]; FD3[ ] --> HoursCell[ ]; FD3[ ] --> EnameCell[ ]; FD3[ ] --> PnameCell[ ]; FD3[ ] --> PlocationCell[ ]
```

(a)

EMP_DEPT

| Ename | Ssn | Bdate | Address | Dnumber | Dname | Dmgr_ssn |
|-------|-----|-------|---------|---------|-------|----------|
| | | | | | | |

```
graph TD; Ename[ ] --> Ename; Ssn[ ] --> Ssn; Bdate[ ] --> Bdate; Address[ ] --> Address; Dname[ ] --> Dname; Dmgr_ssn[ ] --> Dmgr_ssn;
```

Any concerns here?

(b)

EMP_PROJ

| Ssn | Pnumber | Hours | Ename | Pname | Plocation |
|-----|---------|-------|-------|-------|-----------|
| FD1 | | | | | |
| FD2 | | | | | |
| FD3 | | | | | |

```
graph TD; Ssn[FD1] --> Ssn; Pnumber[FD1] --> Pnumber; Hours[FD1] --> Hours; Ename[FD2] --> Ename; Pname[FD2] --> Pname; Plocation[FD2] --> Plocation; FD3[ ] --> Ssn; FD3[ ] --> Pnumber; FD3[ ] --> Hours; FD3[ ] --> Ename; FD3[ ] --> Pname; FD3[ ] --> Plocation;
```

EMP_DEPT: mixing attributes of employees & departments

EMP_PROJ: mixing attributes of employees, projects & works_on

| EMP_DEPT | | | | | | | Redundancy |
|----------------------|-----------|------------|--------------------------|---------|----------------|-----------|------------|
| Ename | Ssn | Bdate | Address | Dnumber | Dname | Dmgr_ssn | |
| Smith, John B. | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | 5 | Research | 333445555 | |
| Wong, Franklin T. | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | 5 | Research | 333445555 | |
| Zelaya, Alicia J. | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX | 4 | Administration | 987654321 | |
| Wallace, Jennifer S. | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | 4 | Administration | 987654321 | |
| Narayan, Ramesh K. | 666884444 | 1962-09-15 | 975 FireOak, Humble, TX | 5 | Research | 333445555 | |
| English, Joyce A. | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | 5 | Research | 333445555 | |
| Jabbar, Ahmad V. | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | 4 | Administration | 987654321 | |
| Borg, James E. | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | 1 | Headquarters | 888665555 | |

Figure 14.4

Sample states for EMP_DEPT and EMP_PROJ resulting from applying NATURAL JOIN to the relations in Figure 14.2.

| EMP_PROJ | | | | | | Redundancy | Redundancy |
|-----------|----------|-------|----------------------|-----------------|-----------|------------|------------|
| Ssn | Pnumber_ | Hours | Ename | Pname | Plocation | | |
| 123456789 | 1 | 32.5 | Smith, John B. | ProductX | Bellaire | | |
| 123456789 | 2 | 7.5 | Smith, John B. | ProductY | Sugarland | | |
| 666884444 | 3 | 40.0 | Narayan, Ramesh K. | ProductZ | Houston | | |
| 453453453 | 1 | 20.0 | English, Joyce A. | ProductX | Bellaire | | |
| 453453453 | 2 | 20.0 | English, Joyce A. | ProductY | Sugarland | | |
| 333445555 | 2 | 10.0 | Wong, Franklin T. | ProductY | Sugarland | | |
| 333445555 | 3 | 10.0 | Wong, Franklin T. | ProductZ | Houston | | |
| 333445555 | 10 | 10.0 | Wong, Franklin T. | Computerization | Stafford | | |
| 333445555 | 20 | 10.0 | Wong, Franklin T. | Reorganization | Houston | | |
| 999887777 | 30 | 30.0 | Zelaya, Alicia J. | Newbenefits | Stafford | | |
| 999887777 | 10 | 10.0 | Zelaya, Alicia J. | Computerization | Stafford | | |
| 987987987 | 10 | 35.0 | Jabbar, Ahmad V. | Computerization | Stafford | | |
| 987987987 | 30 | 5.0 | Jabbar, Ahmad V. | Newbenefits | Stafford | | |
| 987654321 | 30 | 20.0 | Wallace, Jennifer S. | Newbenefits | Stafford | | |
| 987654321 | 20 | 15.0 | Wallace, Jennifer S. | Reorganization | Houston | | |
| 888665555 | 20 | Null | Borg, James E. | Reorganization | Houston | | |

1.2 Redundant Information in Tuples and Update Anomalies

Information is stored redundantly

- Wastes storage

- Causes problems with update anomalies

 - Insertion anomalies

 - Deletion anomalies

 - Modification anomalies

EXAMPLE OF AN INSERT ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Insert Anomaly:

Cannot insert a project unless an employee is assigned to it

Conversely

Cannot insert an employee unless an he/she is assigned to a project

EXAMPLE OF A DELETE ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Delete Anomaly:

When a project is deleted, it will result in deleting all the employees who work on that project.

Alternately, if an employee is the sole employee on a project, deleting that employee would result in deleting the corresponding project.

EXAMPLE OF AN UPDATE ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Update Anomaly:

Changing the name of project number P1 from “Billing” to “Customer-Accounting” may cause this update to be made for all 100 employees working on project P1.

EXAMPLE OF AN UPDATE ANOMALY

| Emp_ID | Emp_Name | Dept_Name | Dept_Location |
|--------|----------|-----------|---------------|
| 101 | Ravi | HR | Delhi |
| 102 | Neha | HR | Delhi |
| 103 | Amit | IT | Mumbai |

EXAMPLE OF AN UPDATE ANOMALY

| Emp_ID | Emp_Name | Dept_Name | Dept_Location |
|--------|----------|-----------|---------------|
| 101 | Ravi | HR | Delhi |
| 102 | Neha | HR | Delhi |
| 103 | Amit | IT | Mumbai |

| Emp_ID | Emp_Name | Dept_Name | Dept_Location |
|--------|----------|-----------|---------------|
| 101 | Ravi | HR | Delhi ✗ |
| 102 | Neha | HR | Noida ✓ |
| 103 | Amit | IT | Mumbai |

```
UPDATE Employee  
SET Dept_Location = 'Noida'  
WHERE Emp_ID = 102;
```

EXAMPLE OF AN UPDATE ANOMALY

Partial or inconsistent updates (human or system error)

Multiple users or applications

Real-world DBs are complex

Guideline for Redundant Information in Tuples and Update Anomalies

GUIDELINE 2:

Design a schema that does not suffer from the insertion, deletion and update anomalies

If there are any anomalies present, then note them so that applications can be made to take them into account

1.3 Null Values in Tuples

GUIDELINE 3:

Relations should be designed such that their tuples will have as few NULL values as possible

Attributes that are NULL frequently could be placed in separate relations (with the primary key)

Reasons for nulls; different meanings for null:

Attribute not applicable or invalid [visa status to US students]

Attribute value unknown [DOB of an employee]

Value is known but absent; it has not been recorded yet [phone # of employee]

1.3 Null Values in Tuples

Poor design:

Employee(EmpID, Name, Department, Salary, Commission, Bonus)

Why? Better design?

1.3 Null Values in Tuples

Poor design:

Employee(EmpID, Name, Department, Salary, Commission, Bonus)

Why? Better design?

Employee(EmpID, Name, Department, Salary)

SalesCommission(EmpID, Commission)

BonusPayment(EmpID, Bonus)

1.4 Generation of Spurious Tuples – avoid at any cost

Bad designs for a relational database may result in erroneous results for certain JOIN operations

GUIDELINE 4:

No spurious tuples should be generated by doing a natural-join of any relations.

(a)

EMP_LOCS

| Ename | Plocation |
|-------|-----------|
| | |

P.K.

EMP_PROJ1

| Ssn | Pnumber | Hours | Pname | Plocation |
|-----|---------|-------|-------|-----------|
| | | | | |

P.K.

(b)

EMP_LOCS

| Ename | Plocation |
|----------------------|-----------|
| Smith, John B. | Bellaire |
| Smith, John B. | Sugarland |
| Narayan, Ramesh K. | Houston |
| English, Joyce A. | Bellaire |
| English, Joyce A. | Sugarland |
| Wong, Franklin T. | Sugarland |
| Wong, Franklin T. | Houston |
| Wong, Franklin T. | Stafford |
| Zelaya, Alicia J. | Stafford |
| Jabbar, Ahmad V. | Stafford |
| Wallace, Jennifer S. | Stafford |
| Wallace, Jennifer S. | Houston |
| Borg, James E. | Houston |

EMP_PROJ1

| Ssn | Pnumber | Hours | Pname | Plocation |
|-----------|---------|-------|-----------------|-----------|
| 123456789 | 1 | 32.5 | ProductX | Bellaire |
| 123456789 | 2 | 7.5 | ProductY | Sugarland |
| 666884444 | 3 | 40.0 | ProductZ | Houston |
| 453453453 | 1 | 20.0 | ProductX | Bellaire |
| 453453453 | 2 | 20.0 | ProductY | Sugarland |
| 333445555 | 2 | 10.0 | ProductY | Sugarland |
| 333445555 | 3 | 10.0 | ProductZ | Houston |
| 333445555 | 10 | 10.0 | Computerization | Stafford |
| 333445555 | 20 | 10.0 | Reorganization | Houston |
| 999887777 | 30 | 30.0 | Newbenefits | Stafford |
| 999887777 | 10 | 10.0 | Computerization | Stafford |
| 987987987 | 10 | 35.0 | Computerization | Stafford |
| 987987987 | 30 | 5.0 | Newbenefits | Stafford |
| 987654321 | 30 | 20.0 | Newbenefits | Stafford |
| 987654321 | 20 | 15.0 | Reorganization | Houston |
| 888665555 | 20 | NULL | Reorganization | Houston |

| Ssn | Pnumber | Hours | Pname | Plocation | Ename |
|-----------|-----------|-------|-----------------|-----------------------------|-----------------------------|
| 123456789 | 1 | 32.5 | ProductX | Bellaire | Smith, John B. |
| * | 123456789 | 1 | 32.5 | ProductX | Bellaire English, Joyce A. |
| 123456789 | 2 | 7.5 | ProductY | Sugarland | Smith, John B. |
| * | 123456789 | 2 | 7.5 | ProductY | Sugarland English, Joyce A. |
| * | 123456789 | 2 | 7.5 | ProductY | Sugarland Wong, Franklin T. |
| 666884444 | 3 | 40.0 | ProductZ | Houston | Narayan, Ramesh K. |
| * | 666884444 | 3 | 40.0 | ProductZ | Houston Wong, Franklin T. |
| * | 453453453 | 1 | 20.0 | ProductX | Bellaire Smith, John B. |
| 453453453 | 1 | 20.0 | ProductX | Bellaire English, Joyce A. | |
| * | 453453453 | 2 | 20.0 | ProductY | Sugarland Smith, John B. |
| 453453453 | 2 | 20.0 | ProductY | Sugarland English, Joyce A. | |
| * | 453453453 | 2 | 20.0 | ProductY | Sugarland Wong, Franklin T. |
| * | 333445555 | 2 | 10.0 | ProductY | Sugarland Smith, John B. |
| * | 333445555 | 2 | 10.0 | ProductY | Sugarland English, Joyce A. |
| 333445555 | 2 | 10.0 | ProductY | Sugarland Wong, Franklin T. | |
| * | 333445555 | 3 | 10.0 | ProductZ | Houston Narayan, Ramesh K. |
| 333445555 | 3 | 10.0 | ProductZ | Houston Wong, Franklin T. | |
| 333445555 | 10 | 10.0 | Computerization | Stafford | Wong, Franklin T. |
| * | 333445555 | 20 | 10.0 | Reorganization | Houston Narayan, Ramesh K. |
| 333445555 | 20 | 10.0 | Reorganization | Houston Wong, Franklin T. | |

*

*

*

Additional tuples that were not there in Emp_proj is here, they are called spurious tuples

2. Functional Dependencies

Functional dependencies (FDs)

Are used to specify *formal measures* of the "goodness" of relational designs

And keys are used to define **normal forms** for relations

Are **constraints** that are derived from the *meaning* and *interrelationships* of the data attributes

A set of attributes X *functionally determines* a set of attributes Y if the value of X determines a unique value for Y

2.1 Defining Functional Dependencies

$X \rightarrow Y$ holds if whenever two tuples have the same value for X , they *must have the same value for Y*

For any two tuples t_1 and t_2 in any relation instance $r(R)$: If $t_1[X]=t_2[X]$, then $t_1[Y]=t_2[Y]$

$X \rightarrow Y$ in R specifies a *constraint* on all relation instances $r(R)$

Written as $X \rightarrow Y$; can be displayed graphically on a relation schema as in Figures; denoted by the arrow \rightarrow

FDs are derived from the real-world constraints on the attributes

Examples of FD constraints (1)

Social security number determines employee name

$$\text{SSN} \rightarrow \text{ENAME}$$

Project number determines project name and location

$$\text{PNUMBER} \rightarrow \{\text{PNAME}, \text{PLOCATION}\}$$

Employee ssn and project number determines the hours per week
that the employee works on the project

$$\{\text{SSN}, \text{PNUMBER}\} \rightarrow \text{HOURS}$$

Examples of FD constraints (1)

Social security number determines employee name

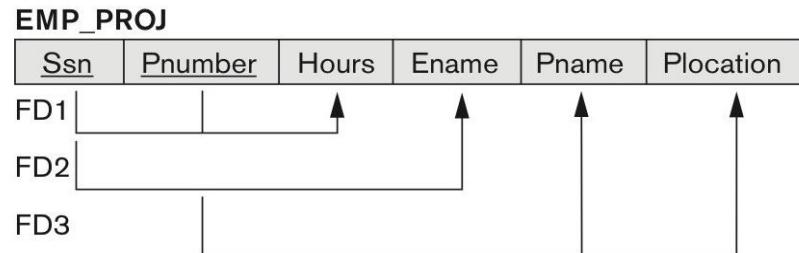
$$\text{SSN} \rightarrow \text{ENAME}$$

Project number determines project name and location

$$\text{PNUMBER} \rightarrow \{\text{PNAME}, \text{PLOCATION}\}$$

Employee ssn and project number determines the hours per week that the employee works on the project

$$\{\text{SSN}, \text{PNUMBER}\} \rightarrow \text{HOURS}$$



Examples of FD constraints (2)

An FD is a property of the attributes in the schema R

The constraint must hold on *every* relation instance $r(R)$

If K is a key of R, then K functionally determines all attributes in R

Defining FDs from instances

Note that in order to define the FDs, we need to understand the meaning of the attributes involved and the relationship between them.

Given the instance (population) of a relation, all we can conclude is that an FD *may exist* between certain attributes.

What we can definitely conclude is – that certain FDs *do not exist* because there are tuples that show a violation of those dependencies.

What can we say about FDs?

TEACH

| Teacher | Course | Text |
|---------|-----------------|----------|
| Smith | Data Structures | Bartram |
| Smith | Data Management | Martin |
| Hall | Compilers | Hoffman |
| Brown | Data Structures | Horowitz |

Ruling Out FDs

Note that given the state of the TEACH relation, we can say that the FD: Text → Course may exist. However, the FDs Teacher → Course, Teacher → Text and Course → Text are ruled out.

TEACH

| Teacher | Course | Text |
|---------|-----------------|----------|
| Smith | Data Structures | Bartram |
| Smith | Data Management | Martin |
| Hall | Compilers | Hoffman |
| Brown | Data Structures | Horowitz |

What FDs may exist?

A relation $R(A, B, C, D)$, which FDs may exist in this relation?

| A | B | C | D |
|----|----|----|----|
| a1 | b1 | c1 | d1 |
| a1 | b2 | c2 | d2 |
| a2 | b2 | c2 | d3 |
| a3 | b3 | c4 | d3 |

What FDs may exist?

A relation $R(A, B, C, D)$, which FDs may exist in this relation?

| A | B | C | D |
|----|----|----|----|
| a1 | b1 | c1 | d1 |
| a1 | b2 | c2 | d2 |
| a2 | b2 | c2 | d3 |
| a3 | b3 | c4 | d3 |

$B \rightarrow C; C \rightarrow B; \{A,B\} \rightarrow C; \{A,B\} \rightarrow D; \{C,D\} \rightarrow B$

How about $A \rightarrow B$? $B \rightarrow A$? $D \rightarrow C$?

Normal Forms Based on Primary Keys

Normalization of Relations

Practical Use of Normal Forms

Definitions of Keys and Attributes Participating in Keys

First Normal Form

Second Normal Form

Third Normal Form

3.1 Normalization of Relations (1)

Normalization:

The process of decomposing unsatisfactory "bad" relations by breaking up their attributes into smaller relations

Normal form:

Condition using keys and FDs of a relation to certify whether a relation schema is in a particular normal form

Normalization of Relations (2)

2NF, 3NF, BCNF

based on keys and FDs of a relation schema

4NF

based on keys, multi-valued dependencies: MVDs;

5NF

based on keys, join dependencies: JDs

Additional properties may be needed to ensure a good relational design (lossless join, dependency preservation; see Chapter 15)

3.2 Practical Use of Normal Forms

Normalization is carried out in practice so that the resulting designs are of high quality and meet the desirable properties

The practical utility of these normal forms becomes questionable when the constraints on which they are based are *hard to understand* or to *detect*

The database designers *need not* normalize to the highest possible normal form (usually up to 3NF and BCNF. 4NF rarely used in practice.)

Denormalization:

The process of storing the join of higher normal form relations as a base relation— which is in a lower normal form

Activity: 10th Nov

Infinium, Develop 2 examples relation where mixing of attributes exist?

Give an example of INSERT, DELETE, UPDATE anomaly in Infinium DB

Write at least 2 FDs that may exist in Infinium DB

Write at least 2 FDs that may not exist in Infinium DB

Write 2 prime attributes in Infinium DB

Write 2 non-prime attributes in Infinium DB

Administrivia

Quiz 3 on 17th Nov

How many will take the Bonus quiz? We will do best of 4; Bonus will be all topics covered and expected to be hard

Bibliography / Acknowledgements

Instructor materials from Elmasri & Navathe 7e



pk.profgiri



Ponnurangam.kumaraguru



/in/ponguru



ponguru



pk.guru@iiit.ac.in

Thank you
for attending
the class!!!