

CS4.301 Data & Applications

Ponnurangam Kumaraguru ("PK")
#ProfGiri @ IIIT Hyderabad



pk.profgiri



/in/ponguru



@ponguru



Ponnurangam.kumaraguru

1.1 Semantics of the Relational Attributes must be clear

GUIDELINE 1: Informally, each tuple in a relation should represent one entity or relationship instance. (Applies to individual relations and their attributes).

Attributes of different entities (EMPLOYEEs, DEPARTMENTs, PROJECTs) should not be mixed in the same relation

Only foreign keys should be used to refer to other entities

Bottom Line: *Design a schema that can be explained easily relation by relation. The semantics of attributes should be easy to interpret.*

(a)

EMP_DEPT

Ename	Ssn	Bdate	Address	Dnumber	Dname	Dmgr_ssn

```
graph TD; Ename[ ] --> Ename; Ssn[ ] --> Ssn; Bdate[ ] --> Bdate; Address[ ] --> Address; Dname[ ] --> Dname; Dmgr_ssn[ ] --> Dmgr_ssn;
```

Any concerns here?

(b)

EMP_PROJ

Ssn	Pnumber	Hours	Ename	Pname	Plocation
FD1					
FD2					
FD3					

```
graph TD; Ssn[FD1] --> Ssn; Pnumber[FD1] --> Pnumber; Hours[FD1] --> Hours; Ename[FD2] --> Ename; Pname[FD2] --> Pname; Plocation[FD2] --> Plocation; FD3[FD3] --> Ssn; FD3[FD3] --> Pnumber; FD3[FD3] --> Hours; FD3[FD3] --> Ename; FD3[FD3] --> Pname; FD3[FD3] --> Plocation;
```

EMP_DEPT: mixing attributes of employees & departments

EMP_PROJ: mixing attributes of employees, projects & works_on

1.2 Redundant Information in Tuples and Update Anomalies

Information is stored redundantly

- Wastes storage

- Causes problems with update anomalies

 - Insertion anomalies

 - Deletion anomalies

 - Modification anomalies

EXAMPLE OF AN INSERT ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Insert Anomaly:

Cannot insert a project unless an employee is assigned to it

Conversely

Cannot insert an employee unless an he/she is assigned to a project

EXAMPLE OF A DELETE ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Delete Anomaly:

When a project is deleted, it will result in deleting all the employees who work on that project.

Alternately, if an employee is the sole employee on a project, deleting that employee would result in deleting the corresponding project.

EXAMPLE OF AN UPDATE ANOMALY

Consider the relation:

EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)

Update Anomaly:

Changing the name of project number P1 from “Billing” to “Customer-Accounting” may cause this update to be made for all 100 employees working on project P1.

Guideline for Redundant Information in Tuples and Update Anomalies

GUIDELINE 2:

Design a schema that does not suffer from the insertion, deletion and update anomalies

If there are any anomalies present, then note them so that applications can be made to take them into account

1.3 Null Values in Tuples

GUIDELINE 3:

Relations should be designed such that their tuples will have as few NULL values as possible

Attributes that are NULL frequently could be placed in separate relations (with the primary key)

Reasons for nulls; different meanings for null:

Attribute not applicable or invalid [visa status to US students]

Attribute value unknown [DOB of an employee]

Value is known but absent; it has not been recorded yet [phone # of employee]

1.3 Null Values in Tuples

Poor design:

Employee(EmpID, Name, Department, Salary, Commission, Bonus)

Why? Better design?

Employee(EmpID, Name, Department, Salary)

SalesCommission(EmpID, Commission)

BonusPayment(EmpID, Bonus)

1.4 Generation of Spurious Tuples – avoid at any cost

Bad designs for a relational database may result in erroneous results for certain JOIN operations

GUIDELINE 4:

No spurious tuples should be generated by doing a natural-join of any relations.

Ssn	Pnumber	Hours	Pname	Plocation	Ename
123456789	1	32.5	ProductX	Bellaire	Smith, John B.
*	123456789	1	32.5	ProductX	English, Joyce A.
123456789	2	7.5	ProductY	Sugarland	Smith, John B.
*	123456789	2	7.5	ProductY	English, Joyce A.
*	123456789	2	7.5	ProductY	Wong, Franklin T.
666884444	3	40.0	ProductZ	Houston	Narayan, Ramesh K.
*	666884444	3	40.0	ProductZ	Wong, Franklin T.
*	453453453	1	20.0	ProductX	Bellaire
453453453	1	20.0	ProductX	Bellaire	Smith, John B.
*	453453453	2	20.0	ProductY	English, Joyce A.
453453453	2	20.0	ProductY	Sugarland	Smith, John B.
*	453453453	2	20.0	ProductY	English, Joyce A.
453453453	2	20.0	ProductY	Sugarland	Wong, Franklin T.
*	333445555	2	10.0	ProductY	Smith, John B.
*	333445555	2	10.0	ProductY	Sugarland
333445555	2	10.0	ProductY	English, Joyce A.	
*	333445555	2	10.0	ProductY	Wong, Franklin T.
*	333445555	3	10.0	ProductZ	Houston
333445555	3	10.0	ProductZ	Narayan, Ramesh K.	
*	333445555	3	10.0	ProductZ	Wong, Franklin T.
333445555	10	10.0	Computerization	Stafford	Wong, Franklin T.
*	333445555	20	10.0	Reorganization	Houston
333445555	20	10.0	Reorganization	Narayan, Ramesh K.	
*	333445555	20	10.0	Reorganization	Wong, Franklin T.

*

*

*

Additional tuples that were not there in Emp_proj is here, they are called spurious tuples

2. Functional Dependencies

Functional dependencies (FDs)

Are used to specify *formal measures* of the "goodness" of relational designs

And keys are used to define **normal forms** for relations

Are **constraints** that are derived from the *meaning* and *interrelationships* of the data attributes

A set of attributes X *functionally determines* a set of attributes Y if the value of X determines a unique value for Y

2.1 Defining Functional Dependencies

$X \rightarrow Y$ holds if whenever two tuples have the same value for X , they *must have the same value for Y*

For any two tuples t_1 and t_2 in any relation instance $r(R)$: If $t_1[X]=t_2[X]$, then $t_1[Y]=t_2[Y]$

$X \rightarrow Y$ in R specifies a *constraint* on all relation instances $r(R)$

Written as $X \rightarrow Y$; can be displayed graphically on a relation schema as in Figures; denoted by the arrow \rightarrow

FDs are derived from the real-world constraints on the attributes

Examples of FD constraints (1)

Social security number determines employee name

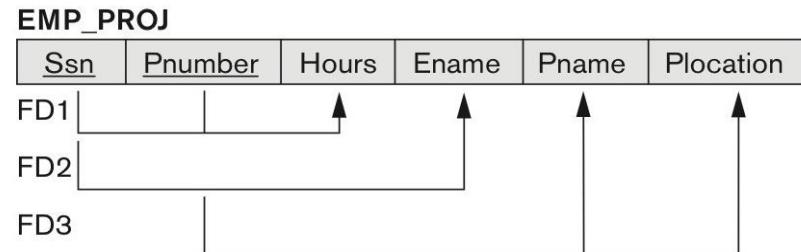
$$\text{SSN} \rightarrow \text{ENAME}$$

Project number determines project name and location

$$\text{PNUMBER} \rightarrow \{\text{PNAME}, \text{PLOCATION}\}$$

Employee ssn and project number determines the hours per week that the employee works on the project

$$\{\text{SSN}, \text{PNUMBER}\} \rightarrow \text{HOURS}$$



Defining FDs from instances

Note that in order to define the FDs, we need to understand the meaning of the attributes involved and the relationship between them.

Given the instance (population) of a relation, all we can conclude is that an FD *may exist* between certain attributes.

What we can definitely conclude is – that certain FDs *do not exist* because there are tuples that show a violation of those dependencies.

What FDs may exist?

A relation $R(A, B, C, D)$, which FDs may exist in this relation?

A	B	C	D
a1	b1	c1	d1
a1	b2	c2	d2
a2	b2	c2	d3
a3	b3	c4	d3

$B \rightarrow C; C \rightarrow B; \{A,B\} \rightarrow C; \{A,B\} \rightarrow D; \{C,D\} \rightarrow B$

How about $A \rightarrow B$? $B \rightarrow A$? $D \rightarrow C$?

3.1 Normalization of Relations (1)

Normalization:

The process of decomposing unsatisfactory "bad" relations by breaking up their attributes into smaller relations

Normal form:

Condition using keys and FDs of a relation to certify whether a relation schema is in a particular normal form

Activity: 10th Nov

Infinium, Develop 2 examples relation where mixing of attributes exist?

Give an example of INSERT, DELETE, UPDATE anomaly in Infinium DB

Write at least 2 FDs that may exist in Infinium DB

Write at least 2 FDs that may not exist in Infinium DB

This Lecture

Definitions of Keys and Attributes Participating in Keys (2)

A **superkey** of a relation schema $R = \{A_1, A_2, \dots, A_n\}$ is a set of attributes S *subset-of* R with the property that no two tuples t_1 and t_2 in any legal relation state r of R will have $t_1[S] = t_2[S]$

If a relation schema has more than one key, each is called a **candidate key**.

One of the candidate keys is *arbitrarily* designated to be the **primary key**, and the others are called **secondary keys**.

A **Prime attribute** must be a member of *some* candidate key

A **Nonprime attribute** is not a prime attribute—that is, it is not a member of any candidate key.

Prime attributes

Consider the relation:

$R(A, B, C, D)$

and suppose the **candidate keys** are:

- $\{A, B\}$
- $\{C, D\}$

Then:

- **Prime attributes:** A, B, C, D (since all are part of some candidate key).
- **Non-prime attributes:** None (since every attribute participates in at least one key).

Prime attributes

Consider the relation:

$R(A, B, C, D)$

and suppose the **candidate keys** are:

- $\{A, B\}$
- $\{C, D\}$

Then:

- **Prime attributes:** A, B, C, D (since all are part of some candidate key)
- **Non-prime attributes:** None (since every attribute participates in a candidate key)

Another example:

$R(A, B, C, D)$ with **candidate key** $\{A, B\}$ only.

Then:

- **Prime attributes:** A, B
- **Non-prime attributes:** C, D

3.4 First Normal Form

Disallow

- composite attributes

- multivalued attributes

- nested relations**; attributes whose values for an *individual tuple* are non-atomic

Considered to be part of the definition of a relation

Most RDBMSs allow only those relations to be defined that are in First Normal Form

Normalization into 1NF

(a)

DEPARTMENT

Dname	Dnumber	Dmgr_ssn	Dlocations



(b)

DEPARTMENT

Dname	Dnumber	Dmgr_ssn	Dlocations
Research	5	333445555	{Bellaire, Sugarland, Houston}
Administration	4	987654321	{Stafford}
Headquarters	1	888665555	{Houston}

Figure 14.9
Normalization into 1NF. (a)
A relation schema that is
not in 1NF. (b) Sample
state of relation
DEPARTMENT

Ways to make it make it 1NF?

1NF

2 = Redundancy,
Dnumber & Dlocation
primary key

DEPARTMENT		
Dname	Dnumber	Dmgr_ssn

F.K.

P.K.

DEPT_LOCATIONS	
Dnumber	Dlocation

F.K.

DEPARTMENT			
Dname	Dnumber	Dmgr_ssn	Dlocation
Research	5	333445555	Bellaire
Research	5	333445555	Sugarland
Research	5	333445555	Houston
Administration	4	987654321	Stafford
Headquarters	1	888665555	Houston

1 = Two 1NF relations

3 = If the maximum number of values (n) for location is known, replace it with n attributes
e.g. Only 3 locations for the company – Dlocation1, Dlocation2, Dlocation3

Introducing NULL if most departments have fewer than 3 locations

Hard to query, e.g. List the departments that have ‘Bellaire’ as one of the locations

1st option is commonly used one

Normalizing nested relations into 1NF

(a)

EMP_PROJ		Projs	
Ssn	Ename	Pnumber	Hours

(b)

EMP_PROJ			
Ssn	Ename	Pnumber	Hours
123456789	Smith, John B.	1	32.5
		2	7.5
666884444	Narayan, Ramesh K.	3	40.0
		1	20.0
453453453	English, Joyce A.	2	20.0
		10	10.0
333445555	Wong, Franklin T.	2	10.0
		3	10.0
		10	10.0
		20	10.0
		30	30.0
999887777	Zelaya, Alicia J.	10	10.0
		30	35.0
987987987	Jabbar, Ahmad V.	30	5.0
		30	20.0
987654321	Wallace, Jennifer S.	20	15.0
		20	NULL
888665555	Borg, James E.	20	

(c)

EMP_PROJ1	
Ssn	Ename

EMP_PROJ2		
Ssn	Pnumber	Hours

Ssn is the primary key, Pnumber is the partial key

Remove the nested relation attributes into a new relation and propagate primary key

This idea can be applied recursively to a relation with multiple-level nesting to unnest

BLOB, CLOB – atomic, single-valued so 1NF

Figure 14.10

Normalizing nested relations into 1NF. (a) Schema of the EMP_PROJ relation with a nested relation attribute PROJS. (b) Sample extension of the EMP_PROJ relation showing nested relations within each tuple. (c) Decomposition of EMP_PROJ into relations EMP_PROJ1 and EMP_PROJ2 by propagating the primary key.

3.5 Second Normal Form (1)

Uses the concepts of **FDs, primary key**

Definitions

Prime attribute: An attribute that is member of some candidate key K

Full functional dependency: a FD $Y \rightarrow Z$ where removal of any attribute from Y means the FD does not hold any more

Examples:

$\{SSN, PNUMBER\} \rightarrow HOURS$ is a full FD since neither $SSN \rightarrow HOURS$ nor $PNUMBER \rightarrow HOURS$ hold

$\{SSN, PNUMBER\} \rightarrow ENAME$ is not a full FD (it is called a partial dependency) since $SSN \rightarrow ENAME$ also holds

Fully functional dependency

If X and Y are an attribute set of a relation, Y is fully functional dependent on X, if Y is functionally dependent on X but not on any proper subset of X.

e.g. In the relation ABC->D, attribute D is fully functionally dependent on ABC and not on any proper subset of ABC. That means that subsets of ABC like AB, BC, A, B, etc cannot determine D.

supplier_id	item_id	price
1	1	540
2	1	545
1	2	200
2	2	201
1	1	540
2	2	201
3	1	542

{ supplier_id , item_id } -> price

Partial dependency

A functional dependency $X \rightarrow Y$ is a partial dependency if Y is functionally dependent on X and Y can be determined by any proper subset of X .

name	roll_no	course
Ravi	2	DBMS
Tim	3	OS
John	5	Java

$\{name\} \rightarrow course$
 $\{roll_no\} \rightarrow course$

Second Normal Form (2)

A relation schema R is in **second normal form (2NF)** if every non-prime attribute A in R is fully functionally dependent on the primary key

A **Prime attribute** must be a member of *some* candidate key

A **Nonprime attribute** is not a prime attribute—that is, it is not a member of any candidate key.

R can be decomposed into 2NF relations via the process of 2NF normalization or “second normalization”

2NF

EMP_PROJ

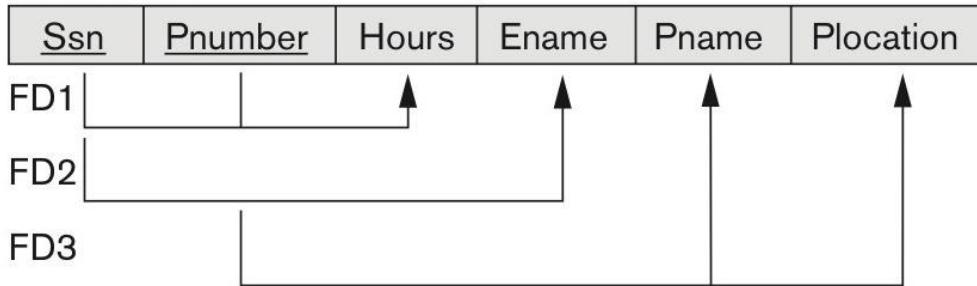
Ssn	Pnumber	Hours	Ename	Pname	Plocation
FD1					
FD2					
FD3					

The diagram illustrates the functional dependencies for the EMP_PROJ relation. FD1 is shown as covering all six columns. FD2 is shown as covering the last five columns. FD3 is shown as covering the last four columns.

Any violation?

2NF

EMP_PROJ



Any violation?

FD2, FD3 violates 2NF

Partial dependency: $\{ssn\} \rightarrow \{Ename\}$, $\{pnumber\} \rightarrow \{pname, plocation\}$ & $\{ssn, pnumber\}$ are primary keys; not fully functionally dependent

Normalizing into 2NF

(a)

EMP_PROJ

Ssn	Pnumber	Hours	Ename	Pname	Plocation
-----	---------	-------	-------	-------	-----------



FD3

2NF Normalization

EP1

Ssn	Pnumber	Hours
-----	---------	-------



EP2

Ssn	Ename
-----	-------



EP3

Pnumber	Pname	Plocation
---------	-------	-----------



Figure 14.11
Normalizing into 2NF and 3NF.
(a) Normalizing EMP_PROJ into
2NF relations.

EP1, EP2, EP3 are fully functionally dependent

Summary

Property	Explanation
Partial Dependencies	Yes — Ssn → Ename and Pnumber → Pname, Plocation
Primary Key	(Ssn, Pnumber)
Why Not 2NF	Non-prime attributes depend on <i>part</i> of the composite key
Fix	Decompose into EMPLOYEE, PROJECT, and EMP_PROJ

Third normal form

Transitive Dependency

$X \rightarrow Y$ in R is transitive dependency, if there exists a set of attributes Z in R that is neither a candidate key nor a subset of any key of R and both $X \rightarrow Z$ & $Z \rightarrow Y$ hold.

Third normal form

Transitive Dependency

$X \rightarrow Y$ in R is transitive dependency, if there exists a set of attributes Z in R that is neither a candidate key nor a subset of any key of R and both $X \rightarrow Z$ & $Z \rightarrow Y$ hold.

EMP_DEPT

Ename	Ssn	Bdate	Address	Dnumber	Dname	Dmgr_ssn

```
graph TD; Ename --> Dnumber; Ssn --> Dnumber; Bdate --> Dnumber; Address --> Dnumber; Dmgr_ssn --> Dnumber; Dnumber --> Dname;
```

Any transitivity?

Third normal form

Transitive Dependency

$X \rightarrow Y$ in R is transitive dependency, if there exists a set of attributes Z in R that is neither a candidate key nor a subset of any key of R and both $X \rightarrow Z$ & $Z \rightarrow Y$ hold.

EMP_DEPT

Ename	Ssn	Bdate	Address	Dnumber	Dname	Dmgr_ssn

```
graph TD; Ssn --> Dnumber; Ssn --> Dmgr_ssn; Dnumber --> Dmgr_ssn;
```

Any transitivity?

$Ssn \rightarrow dmgr_ssn$ is transitive through dnumber

Both $ssn \rightarrow dnumber$ & $dnumber \rightarrow dmgr_ssn$ hold & dnumber is neither a key nor a subset of a key

Third normal form

R is in 3NF if it satisfies 2NF and no nonprime attribute of R is transitively dependent on the primary key

EMP_DEPT

Ename	<u>Ssn</u>	Bdate	Address	Dnumber	Dname	Dmgr_ssn

```
graph TD; Ssn --> Dname; Bdate --> Dname; Address --> Dname; Dnumber --> Dname; Dname --> Dmgr_ssn;
```

Normalizing EMP_DEPT into 3NF relations.

3NF Normalization

ED1

Ename	<u>Ssn</u>	Bdate	Address	Dnumber

```
graph TD; Ssn --> Dname; Bdate --> Dname; Address --> Dname; Dnumber --> Dname;
```

ED2

Dnumber	Dname	Dmgr_ssn

```
graph TD; Dname --> Dmgr_ssn;
```

ED1 & ED2 represent independent facts about employees & departments

Another example

STUDENT (RollNo, Sname, DeptNo, DeptName, DeptLocation)

Functional Dependencies:

1. RollNo → Sname, DeptNo
2. DeptNo → DeptName, DeptLocation

Step 1: Identify the key

- RollNo is the **primary key** (uniquely identifies a student).

Step 2: Find indirect dependencies

- RollNo → DeptNo (direct)
- DeptNo → DeptName, DeptLocation
→ So, RollNo → DeptName, DeptLocation (transitively)

Another example

◆ Why It's a Problem

Attributes like DeptName and DeptLocation depend **indirectly** on RollNo — through DeptNo.

That means the department info is **redundantly stored** for every student in that department.

If a department moves to a new location, we must update many rows — this leads to **update anomalies**.

◆ Solution

To eliminate transitive dependencies, we move to **Third Normal Form (3NF)**:

STUDENT (RollNo, Sname, DeptNo)

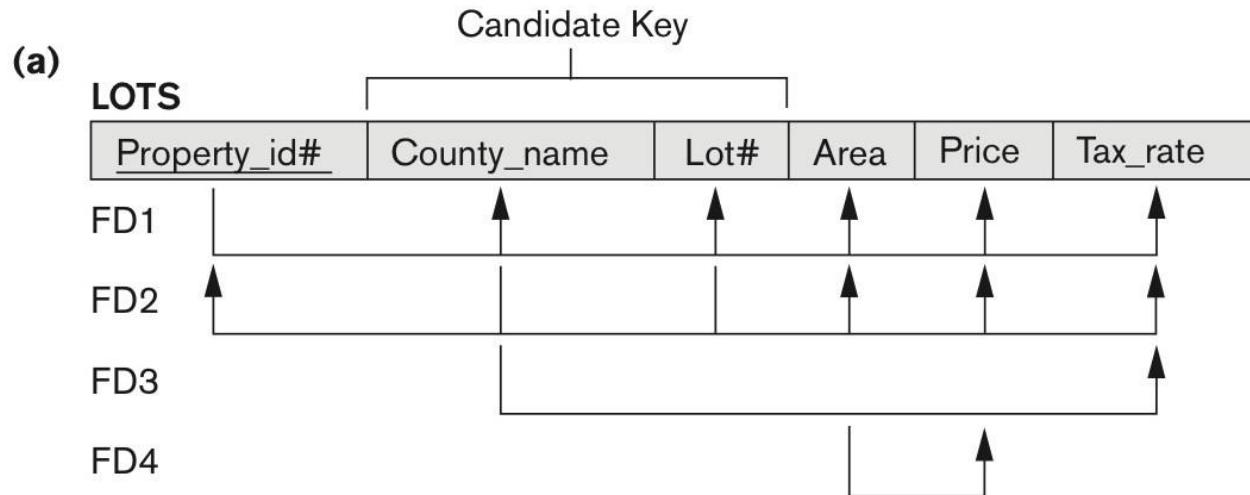
DEPARTMENT (DeptNo, DeptName, DeptLocation)

RollNo determines Sname, DeptNo — direct dependency.

DeptNo determines DeptName, DeptLocation — separate table.

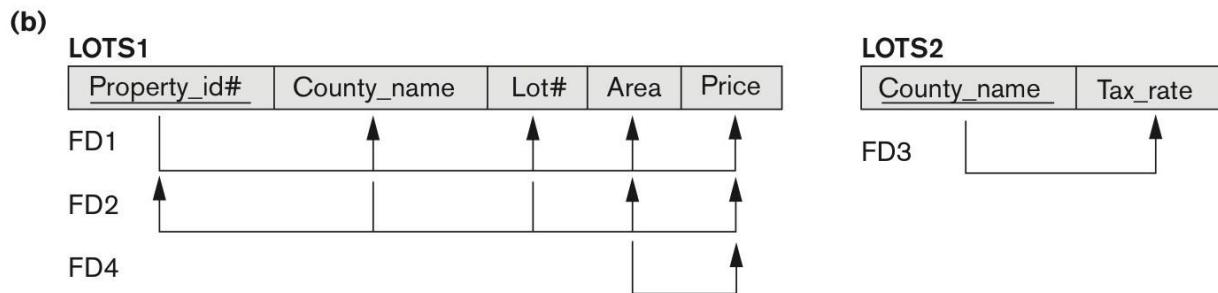
No transitive dependencies.

Figure 14.12a Normalization into 2NF and 3NF. The LOTS relation with its functional dependencies FD1 through FD4.



continued on next slide

Figure 14.12b Normalization into 2NF and 3NF.
Decomposing into the 2NF relations LOTS1 and LOTS2.



continued on next slide

Figure 14.12c Normalization into 2NF and 3NF.
Decomposing LOTS1 into the 3NF relations LOTS1A and
LOTS1B.

(c)
LOTS1A

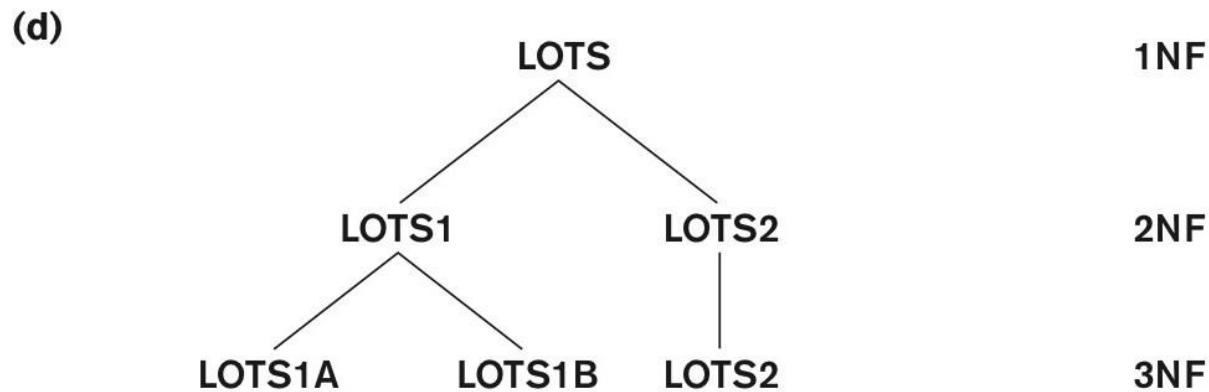
<u>Property_id#</u>	County_name	Lot#	Area
FD1			
FD2			

LOTS1B

<u>Area</u>	Price
FD4	

continued on next slide

Figure 14.12d Normalization into 2NF and 3NF.
Progressive normalization of LOTS into a 3NF design.



5. BCNF (Boyce-Codd Normal Form)

A relation schema R is in **Boyce-Codd Normal Form (BCNF)** if whenever an **FD** $X \rightarrow A$ holds in R, then **X is a superkey** of R

Refresher: Roll number: Candidate key; Roll number + Name: Super key

Each normal form is strictly stronger than the previous one

Every 2NF relation is in 1NF

Every 3NF relation is in 2NF

Every BCNF relation is in 3NF

There exist relations that are in 3NF but not in BCNF

Hence BCNF is considered a **stronger form of 3NF**

The goal is to have each relation in BCNF (or 3NF)

Figure 14.13 Boyce-Codd normal form

LOTS1A

<u>Property_id#</u>	County_name	Lot#	Area
FD1			
FD2			

Is this in BCNF?

LOTS1A

<u>Property_id#</u>	County_name	Lot#	Area
FD1			
FD2			
FD5			

Figure 14.13 Boyce-Codd normal form

LOTS1A

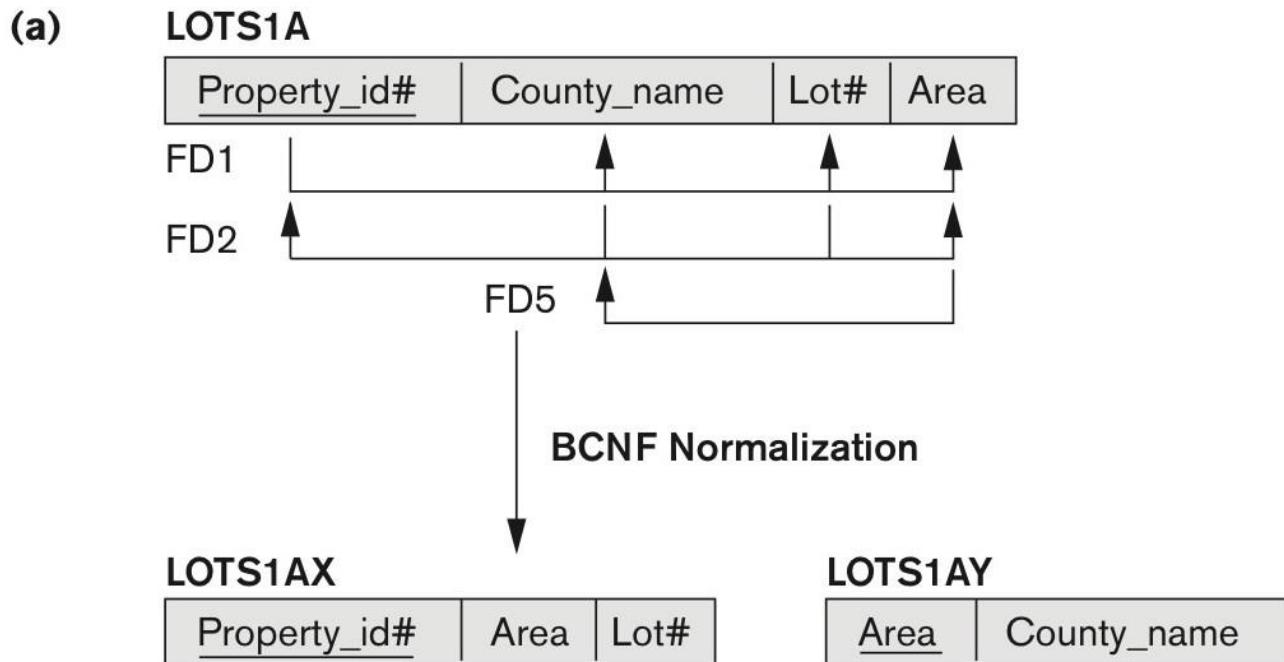
Property_id#	County_name	Lot#	Area
FD1			
FD2			

Is this in BCNF?
Area is not a superkey of LOTS1A

LOTS1A

Property_id#	County_name	Lot#	Area
FD1			
FD2			
FD5			

Figure 14.13 Boyce-Codd normal form

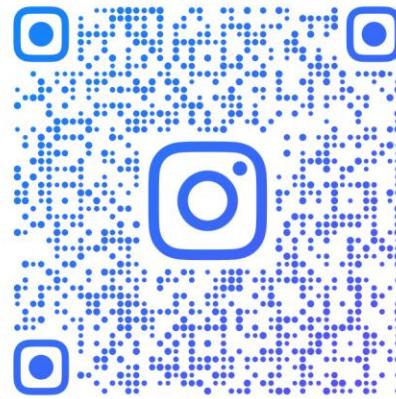


Normalization

<https://www.youtube.com/watch?v=ABwD8IYByfk>

Search for: Ponnurangam Kumaraguru

<https://www.linkedin.com/in/ponguru/>



@PK.PROFGIRI



<https://twitter.com/ponguru>

Interested in working with Precog?

Look for an email in Jan 1st week after start of classes

Areas / Questions: NLP & Responsible AI; Computer Vision; Graphs; Math / Quant

Process: Apply (SOP, CV, etc.) – Task – Technical Interview

What to prepare? Pick an area from above & try out topics

All process done by mid sem

Filter by

- <> Code 24.7k
- Repositories 124**
- Issues 29
- Pull requests 32
- Discussions 2
- Users 0
- More

Languages

- Jupyter Notebook
- Python

124 results (189 ms)

Sort by: Best match

Save ...

anant14014/TwitterHinglishTranslation
Submission for the Programming Task for the Precog Recruitment Process (II)
Python · ⭐ 14 · Updated on Jul 29, 2016

shirebox/Precog-Task1
Recruitment Task for Precog
HTML · ⭐ 3 · Updated on Dec 19, 2017

JJumSSu/PRECOG
Prediction of Results from Coarse Task Descriptions Gathered from the Literature
Jupyter Notebook · ⭐ 2 · Updated on Oct 5

Administrivia

Quiz 3 on 17th Nov

Make up quiz on the day of end-sem exams; if you take the make up quiz it will be counted one of the 3

Same day as end sem at 8 – 9PM

Activity: 13th Nov

- (0) Write 2 prime & 2 full dependency attributes in Infinium DB
- (1) Write 2 non-prime & 2 partial dependency attributes in Infinium DB

Relation: EMP_PROJECT

| EmpNo | EmpName | ProjNo | ProjName | DeptNo | DeptName | HoursWorked | DeptLocation |

Given information:

1. Each employee works in one department, but a department can have many employees.
2. Each project is controlled by one department.
3. An employee can work on multiple projects, and a project can have multiple employees.
4. HoursWorked is the number of hours that an employee works on a particular project.

- (2a) Identify all functional dependencies (FDs).
- (2b) Find the candidate keys.
- (2c) Convert the relation into 1NF, 2NF, 3NF, and BCNF.
- (2d) Draw the resulting schemas at each stage.

Bibliography / Acknowledgements

Instructor materials from Elmasri & Navathe 7e



pk.profgiri



Ponnurangam.kumaraguru



/in/ponguru



ponguru



pk.guru@iiit.ac.in

Thank you
for attending
the class!!!