

# CS3.301 Operating Systems and Networks

## Concurrency - Locks

Karthik Vaidhyanathan

<https://karthikvaidhyanathan.com>

1



# Acknowledgement

The materials used in this presentation have been gathered/adapted/generate from various sources as well as based on my own experiences and knowledge -- Karthik Vaidhyanathan

## Sources:

- Operating Systems in three easy pieces by Remzi et al.



# Shared Data - More Tricky

```
void *worker_thread(void *arg)
{
    int index;
    for (index =0; index<max_index; index++)
    {
        counter++;
    }
}
```

Initial value of the counter 0  
Final value of the counter 4000

Initial value of the counter 0  
Final value of the counter 3790

- Max size: 2000, assume global variable counter initialised to 0
- Desired final result: **4000!**, even on a single processor system there is no guarantee

• Why does this happen?

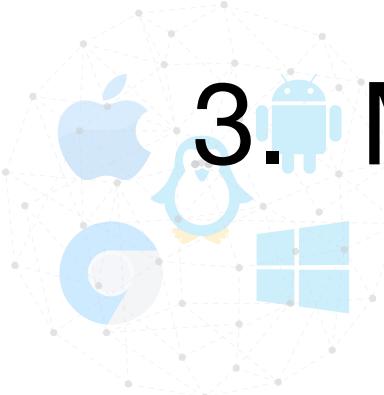


# Lets break the code down in assembly

**counter = counter + 1**

```
mov 0x8049a1c, %eax  
add &0x1, %eax  
mov %eax, 0x8049a1
```

1. Load memory value to register eax
2. Increment the value in the register by 1
3. Move the value from register back to memory



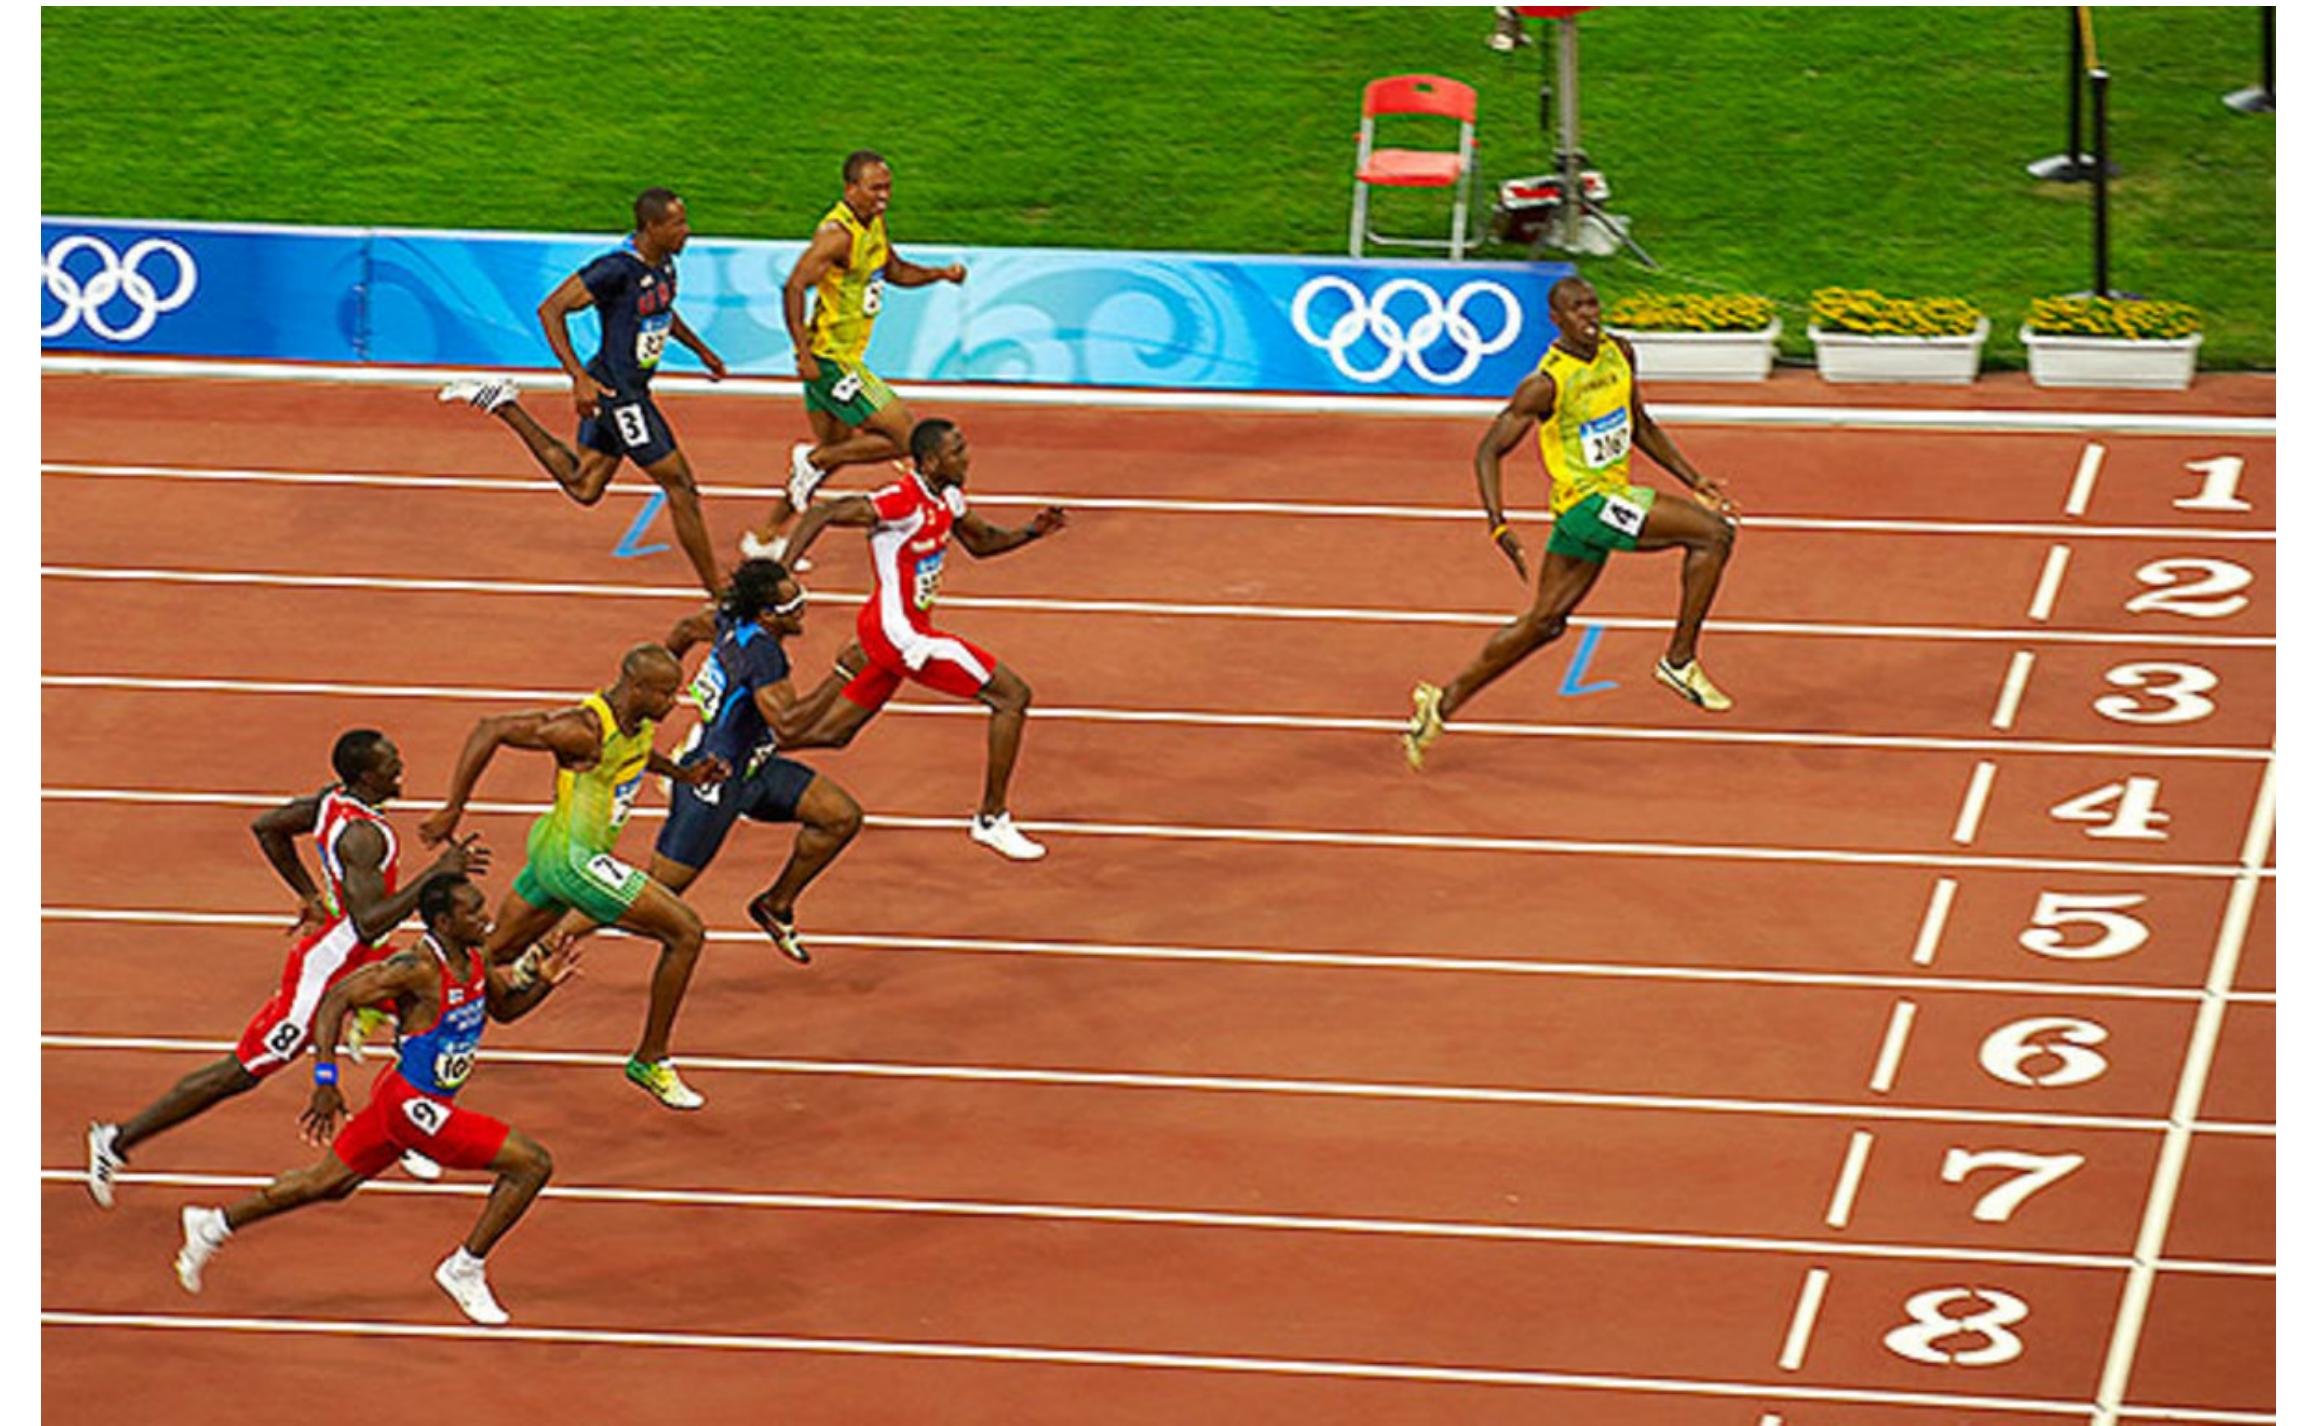
# What can happen?

OS	Thread 1	Thread 2	Counter and Register (Initial value of counter = 50)
	mov 0x8049a1c, %eax add \$0x1, %eax		eax = 51 <b>counter = 50</b>
<b>interrupt</b> Save T1' state Restore T2's state			
		mov 0x8049a1c, %eax add \$0x1, %eax Mov %eax, 0x8049a1c	eax = 51 <b>counter = 51</b>
<b>interrupt</b> Save T2' state Restore T1's state			



# Race Condition and Critical Section

- **Race Condition:** Condition where
  - Multiple threads executing concurrently and
  - results depend on order of execution (time)
  - Scheduler can swap threads, also interrupts
  - Non-deterministic results
- **Critical Section:**
  - The section of code that is shared between the threads (leads to race conditions)
  - Shared variables or data



Source: <https://www.si.com/olympics/2016/07/14/usain-bolt-2016-rio-olympics>



# Concurrency is tricky!

## Race conditions can result in fatal issues



**Therac 25**

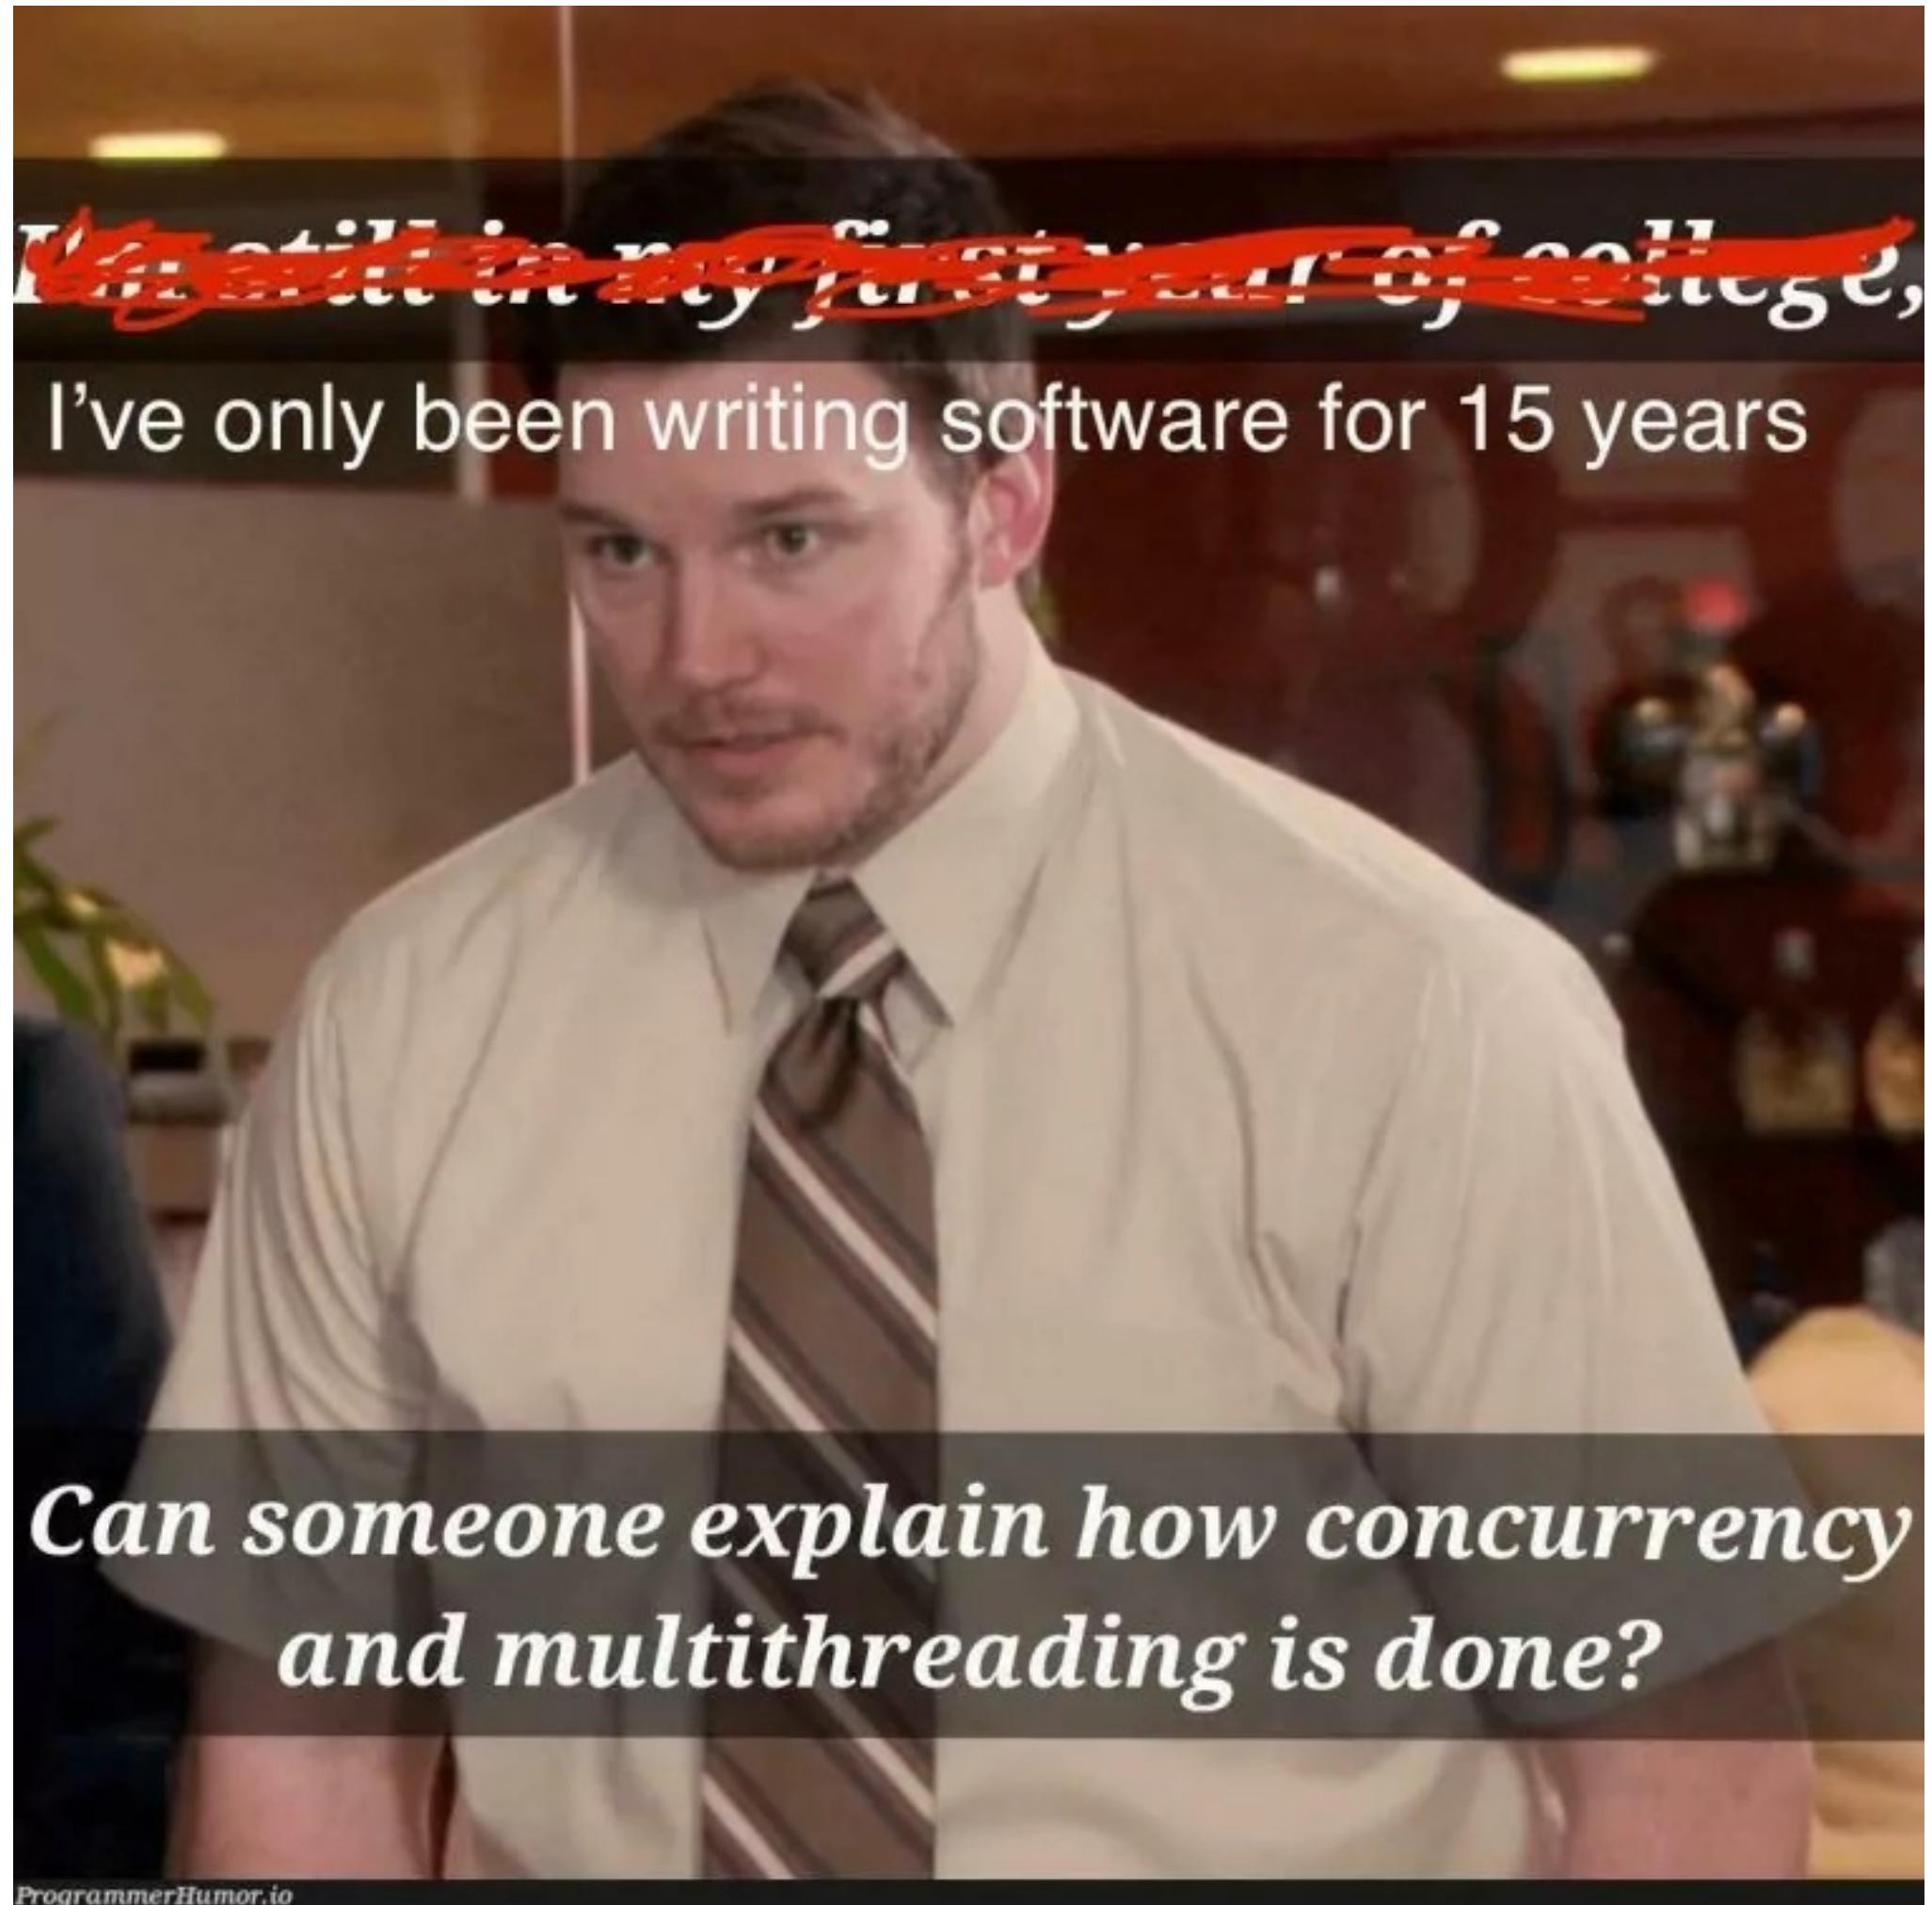


**Northeast Blackout of 2003, US**

<https://www.everydayshouldbesaturday.com/2018/8/14/17687734/flashback-the-blackout-of-2003>



# Concurrency can be tricky!



Source: [programmerhumor.io](http://programmerhumor.io)

# What can be done?

## Bring in Atomicity

- What we want here is mutual exclusion!
  - When one thread is accessing critical section, others should wait
  - No two threads should access critical section at the same time
- In other words, **atomicity** needs to be provided
  - What if there was one instruction in assembly:
    - memory-add 0x8049a1c, &0x1 - Reality is not this!!

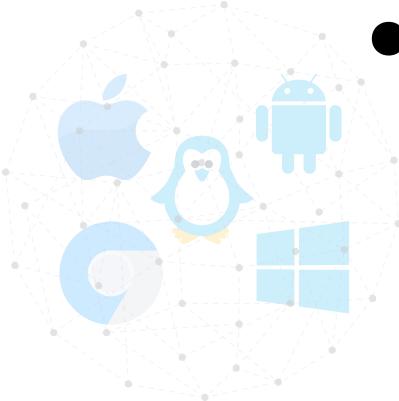
```
mov 0x8049a1c, %eax  
add &0x1, %eax  
mov %eax, 0x8049a1
```

**This should execute atomically!!**



# What we need?

- Need to build synchronization primitives
  - Hardware + software support
  - Ensure that critical section is accessed in synchronised and controlled manner
- **One part:** Build some primitives for synchronisation
  - This will ensure atomicity (avoid race conditions)
- **Second part:** Ensure every thread gets access!
  - No one should starve



# Some Issues needs to be addressed

- What support do we need from hardware?
- What support is needed from software?
- How to build primitives correctly and effectively?
- How can programs use these primitives?



# Locks: A Basic Idea

- Lets go back to the shared variable code in **Critical Section (CS)**:

*counter = counter + 1*

- What if we can have a lock surrounding the statement

```
● ● ●  
lock_t mutex; //a global lock  
....  
lock (&mutex); // lock the critical section  
counter = counter + 1;  
unlock (&mutex); // release the lock
```



# What are locks?

- Lock here is just a variable
- The lock variable holds the state of lock at any instant of time
- It is either available (free) or acquired:
  - **Available:** No threads hold the lock
  - **Acquired:** Lock not available, one thread is holding it and in CS
- We can also enrich with more information - which thread holds the lock, create a queue for threads to get locks, etc.



# Lock and Unlock

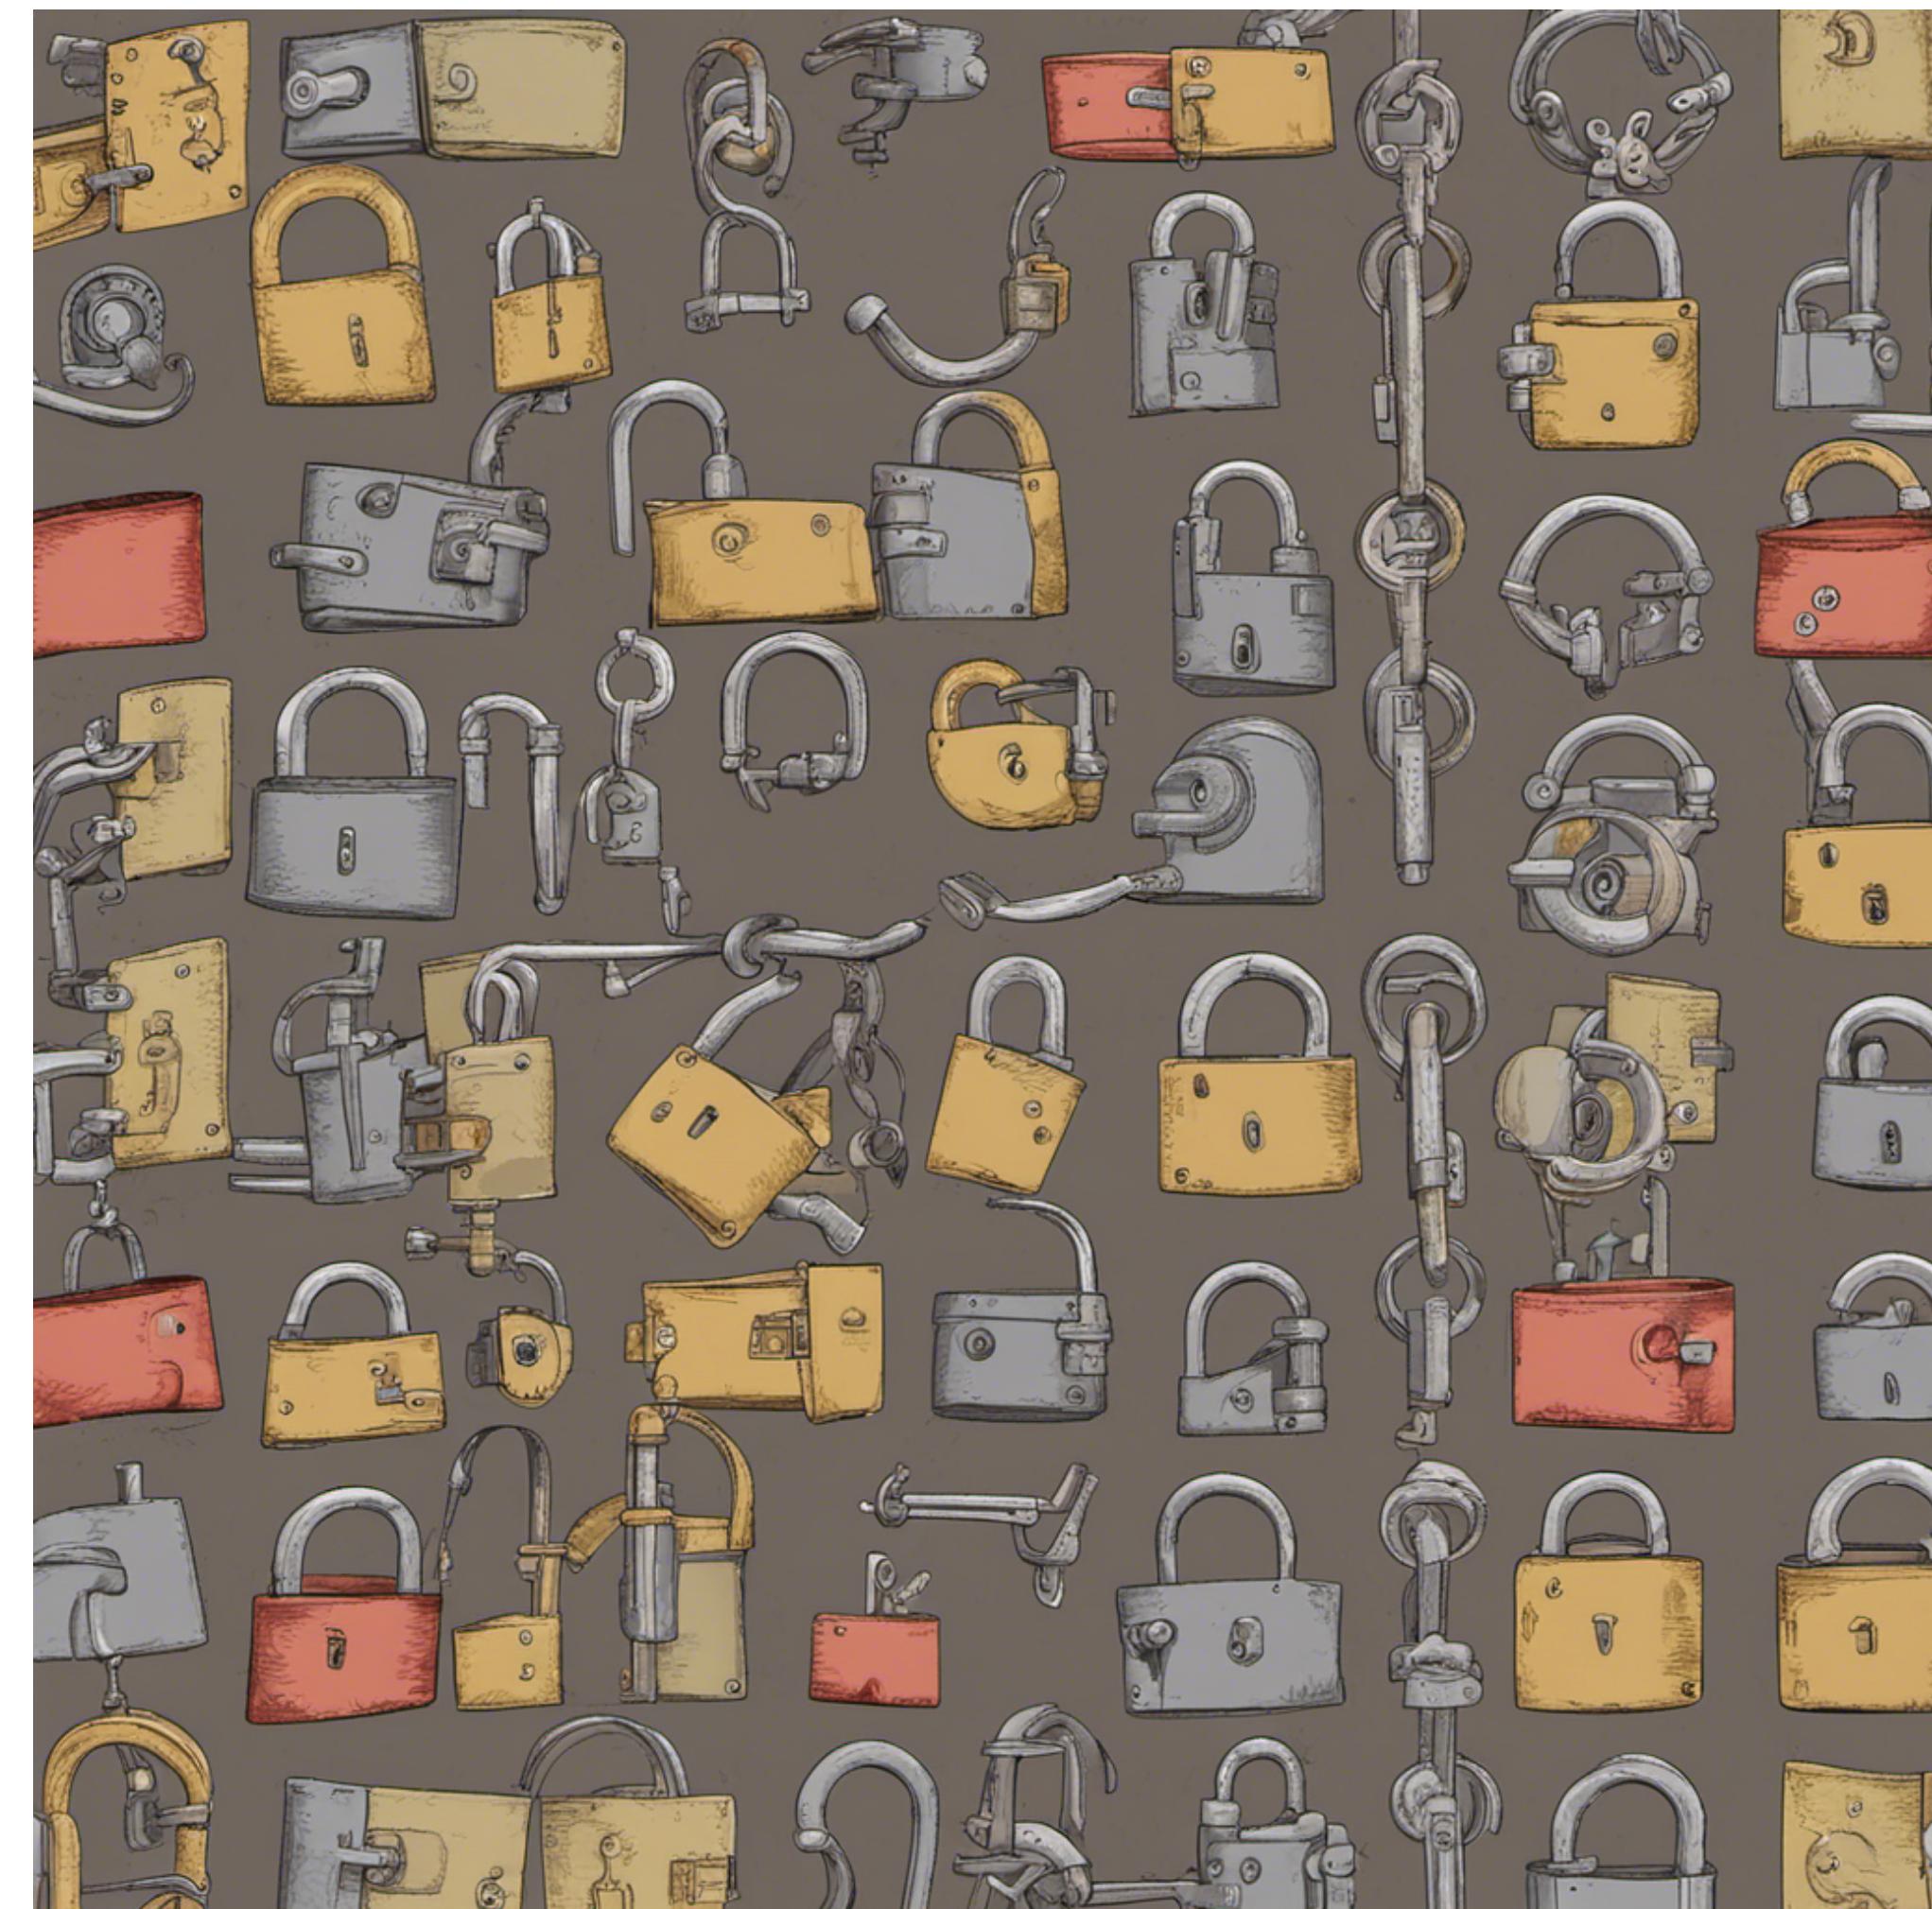
- The thread that holds the lock - **Owner**
- Owner needs to call unlock to free the lock
  - The lock becomes free
  - There may be threads waiting for the lock, one of them will acquire it
  - The next thread with lock enters CS
  - When no thread is waiting for the lock, the lock stays as free

• **How to go about building a lock?**



# How to go about building a lock?

- Think about classrooms and locks
- In the physical locks itself there are many options
- Many hardware primitives have been added to instruction set architecture to support locks
- The way in which the primitives are used + OS support forms the key



# Criteria to evaluate

- **Mutual Exclusion**
  - Does the lock prevent multiple threads from entering CS at same time?
- **Fairness**
  - Does each thread get a fair chance to enter into the CS? - Avoid starvation!!
- **Performance**
  - When there is only one thread what is the overhead?
  - Multiple threads on single CPU - What about overhead?
  - Multiple threads on multiple CPU - What's are overhead?



# Why threading is challenging?

- Two threads running at the same time
  - Interrupts - Can we disable interrupt inside a lock?

```
● ● ●

void lock()
{
    DisableInterrupts();
    //disables interrupts
}

void unlock()
{
    EnableInterrupt();
}
```

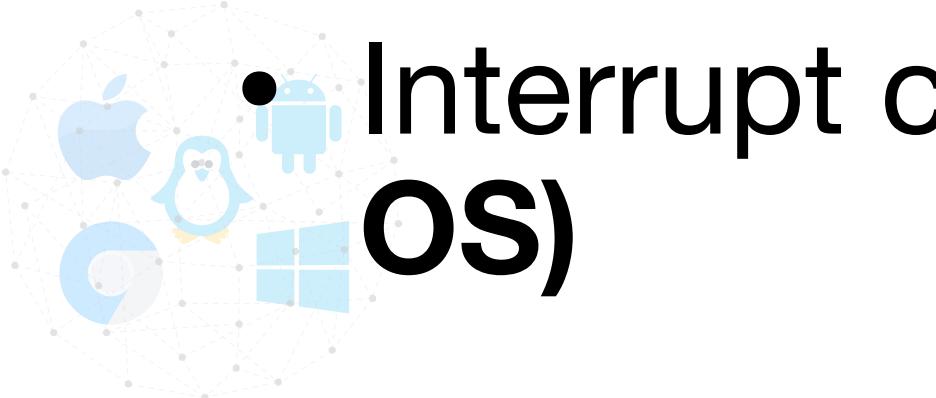
**Main positive approach is simplicity**

**Are there any issues?**



# There are Negatives

- Thread gets a **very high privilege**
  - Thread can switch on and off the interrupts
  - Arbitrary thread can monopolize the processor
  - Errant program could call lock() and get into endless loop
- The approach does not work on multiple processor systems
  - Even if interrupts are disabled, other threads can run on different processor
- **Inefficiency:** Code that masks or unmasks interrupts are executed slowly
- Interrupt control is used in limited context as mutual exclusion primitive (**inside OS**)



# Can we try with a software lock?



Software based lock

```
typedef struct __lock_t
{
    int flag;
} lock_t;

void init (lock_t *mutex)
{
    mutex -> flag = 0;
}

void lock(lock_t *mutex)
{
    while(mutex->flag==1)
        //do nothing
    mutex -> flag = 1;
}

void unlock (lock_t *mutex)
{
    mutex->flag = 0;
}
```

- Use software based locking mechanism
- Create a lock which has a flag value
- Every time a thread wants to enter CS
  - Invoke lock function, if flag = 1, wait for the lock to be available
  - Else acquire the lock and enter CS
- **Do you foresee some problem here?**



# Simple Trace

Thread 1	Thread 2
<b>Call lock()</b> <b>while (flag== 1)</b> <b>Interrupt to Thread 2</b>	
	<b>Call lock()</b> <b>While (flag==1)</b> <b>flag = 1;</b> <b>Interrupt to thread 1</b>
<b>flag = 1</b>	

- No mutual exclusion - Both threads have flag set to 1
- Performance overhead due to spin-waiting



# Working Spin Lock with Test-And-Set

## Test and Set Instruction



Test and Set Lock

```
int TestAndSet(int *ptr, int new)
{
    int old = *ptr;
    *ptr = new;
    return old;
}
```

C Pseudocode

- Simplest hardware primitive - test-and-set or atomic exchange instruction
- Sequence of instructions executed atomically
- Enables testing of old values while setting the main to a new value
- Think about implementing a CS code with this lock!



# Implementing using Test and Set Lock



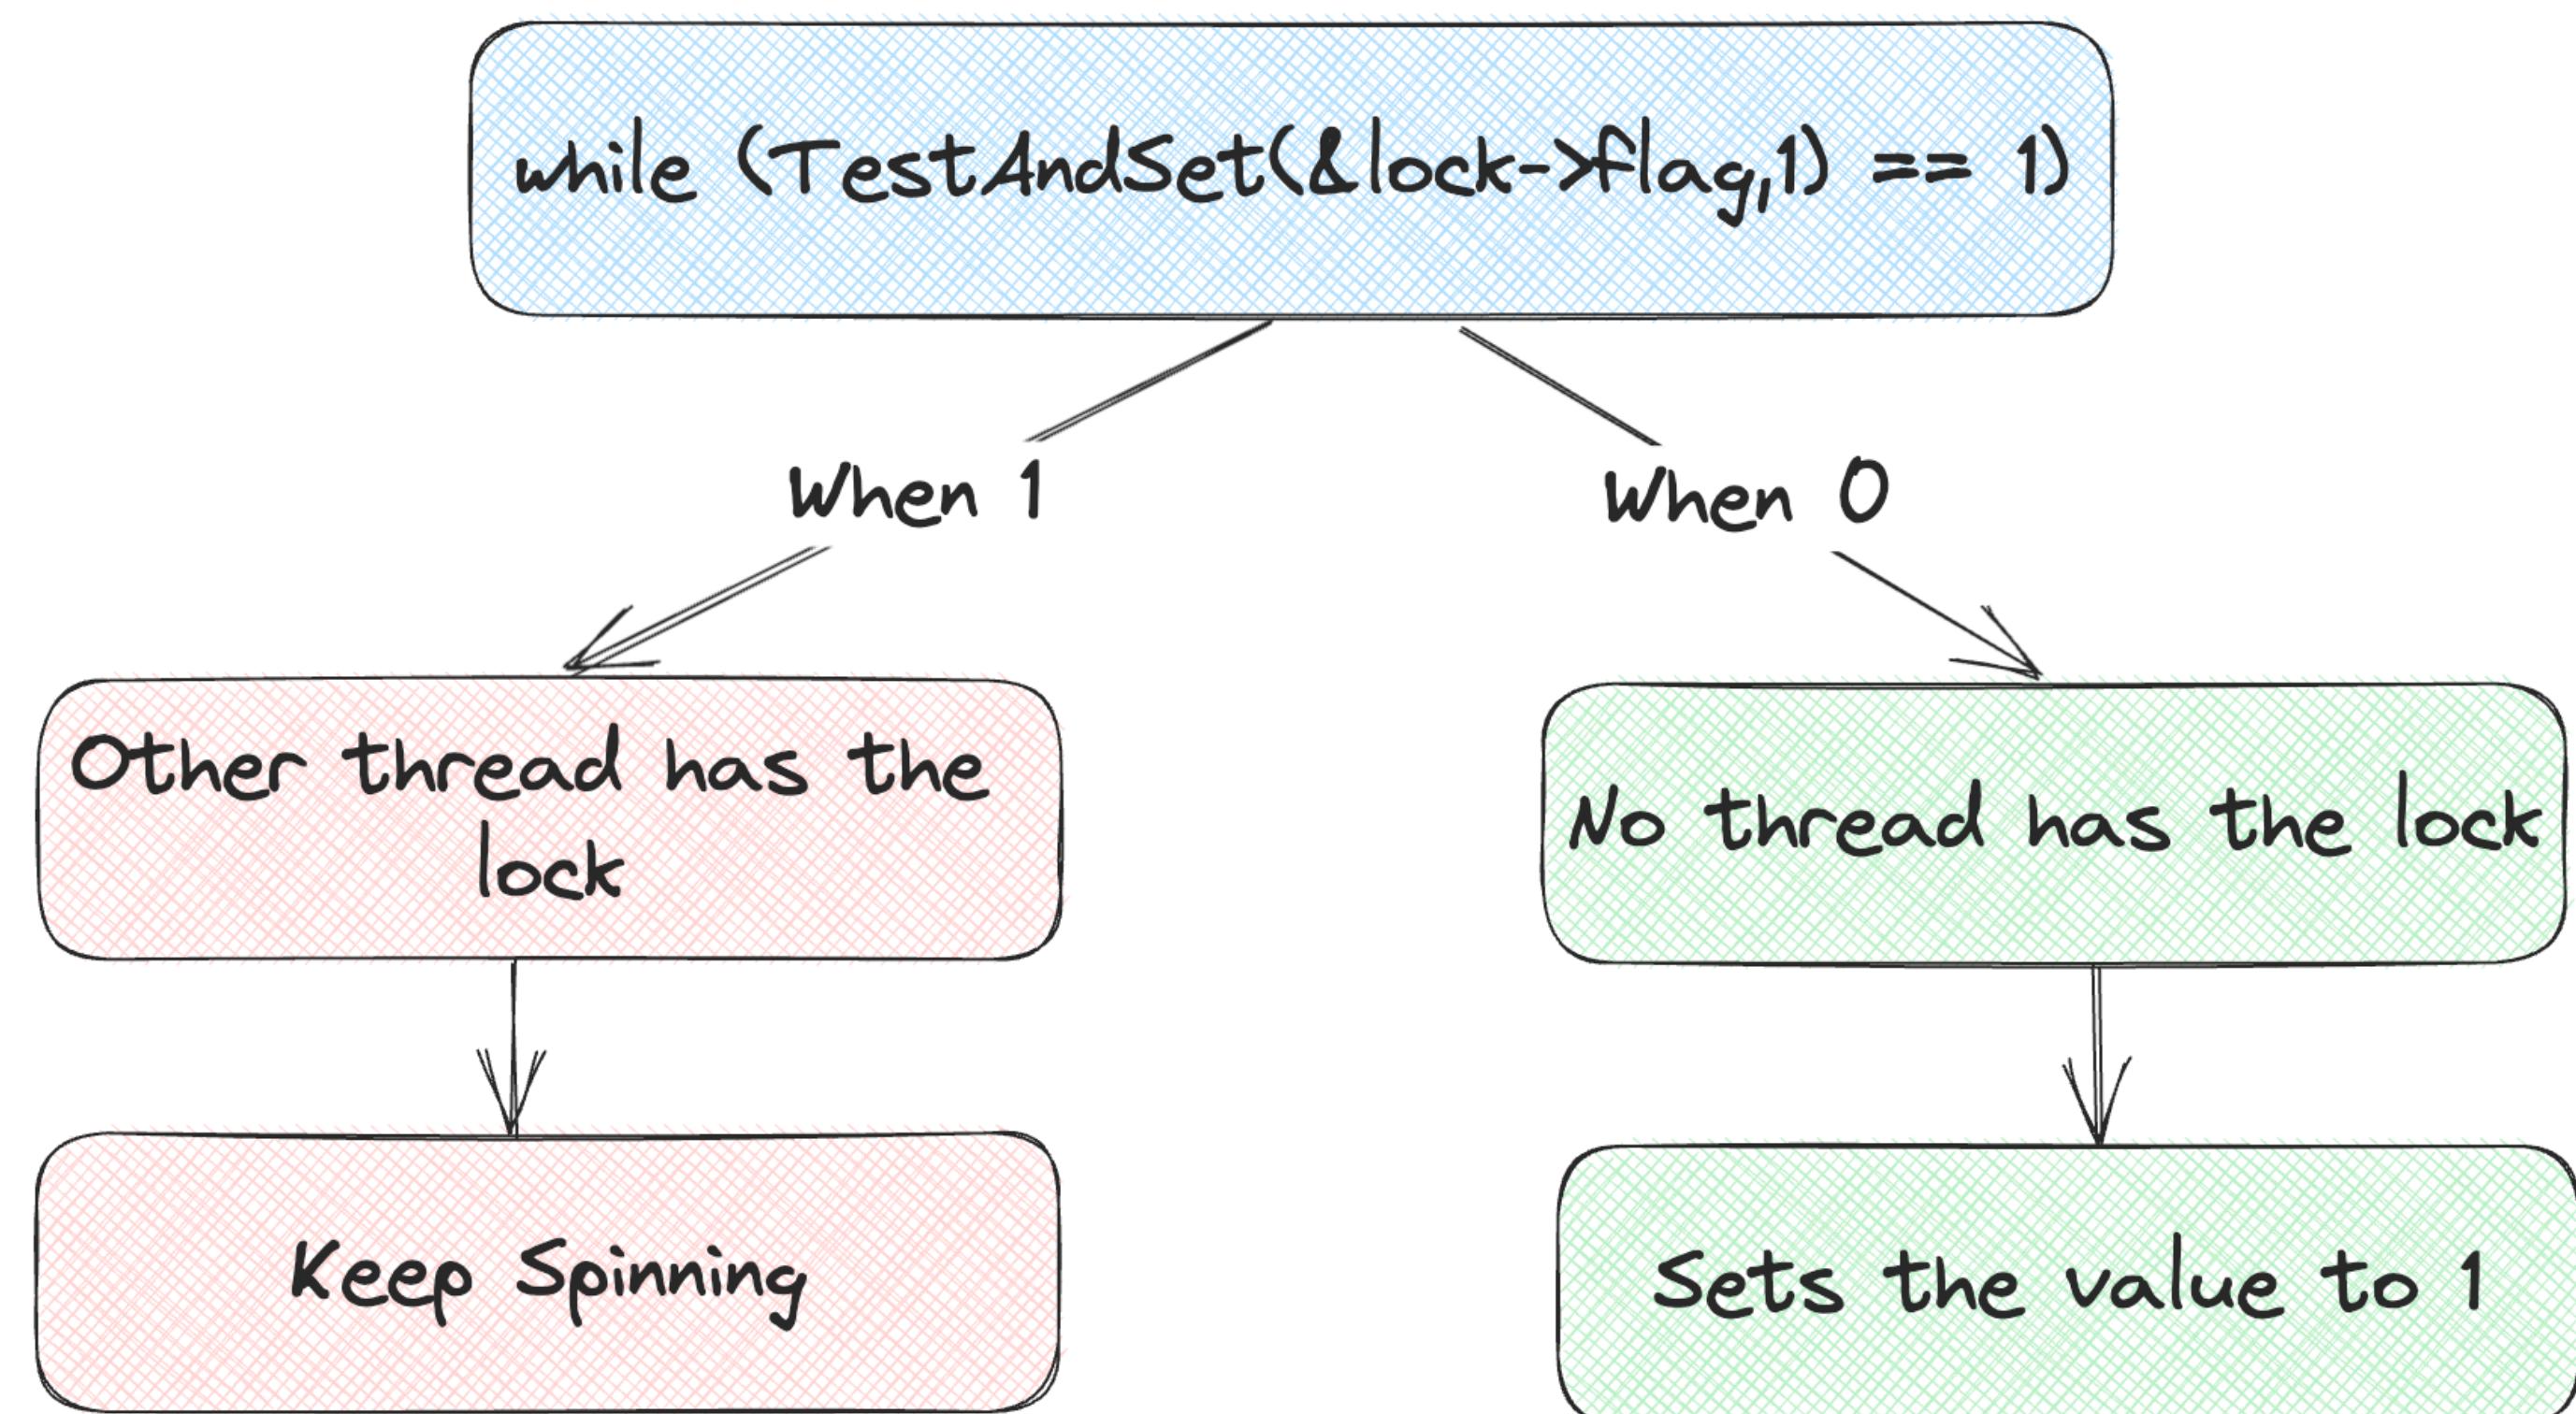
Using TestAndSet Lock

```
typedef struct __lock_t
{
    int flag
} lock_t;

void init (lock_t *lock)
{
    lock -> flag = 0;
}

void lock (lock_t *lock)
{
    while (TestAndSet(&lock->flag, 1)==1)
        // Keep spinning
}

void unlock (lock_t *lock)
{
    lock -> flag = 0;
}
```



Only one thread can acquire the lock



# Evaluating Spin Locks

- **Correctness perspective** Provides mutual exclusion, so correct!
- **Fairness perspective**
  - They don't provide fairness guarantee
  - Threads can keep spinning forever leading to starvation
- **Performance perspective**
  - In single CPU - significant overhead, What if thread holding the lock is preempted in critical section? - Run all N-1 threads to waste CPU
  - On multiple CPUs these locks work reasonably well (esp if n. Threads ~ n. CPUs)



# Compare-And-Swap

## Another Hardware Primitive

### C Pseudocode

● ● ●  
Compare-And-Swap

```
int CompareAndSwap(int *ptr, int expected, int new)
{
    int actual = *ptr;
    if (actual == expected)
    {
        *ptr = new;
    }
    return actual;
}
```

Inside the lock function the call will be:

```
while(CompareAndSwap(&lock->flag, 0,1) ==1)
```

- **Basic idea:** Test whether the value at address specified by ptr is equal to expected
  - If yes, update the memory location pointed to ptr by new value
  - If not, do nothing!
- In either case, return actual value at the memory location



# Load-Linked and Store-Conditional (LL SC)

- In MIPS architecture, they can be used in tandem to build locks and concurrent structures



Load Linked

```
int LoadLinked (int *ptr)
{
    return *ptr;
}
```

C Pseudocode



Store Conditional

```
int StoreConditional(int *ptr, int value)
{
    if (no one has updated ptr since load linked)
    {
        *ptr = value;
        return 1;
    }
    else
    {
        return 0;
    }
}
```



# LL/SC for building locks

- Load linked is like a typical load operation
  - Simply fetches a value from memory and places it in a register
- Store conditional succeeds if no intermittent store to address has taken place
  - In case of success, it updates ptr to value and returns 1 else returns 0

When does failure condition of store conditional arise?



LL/SC for building locks

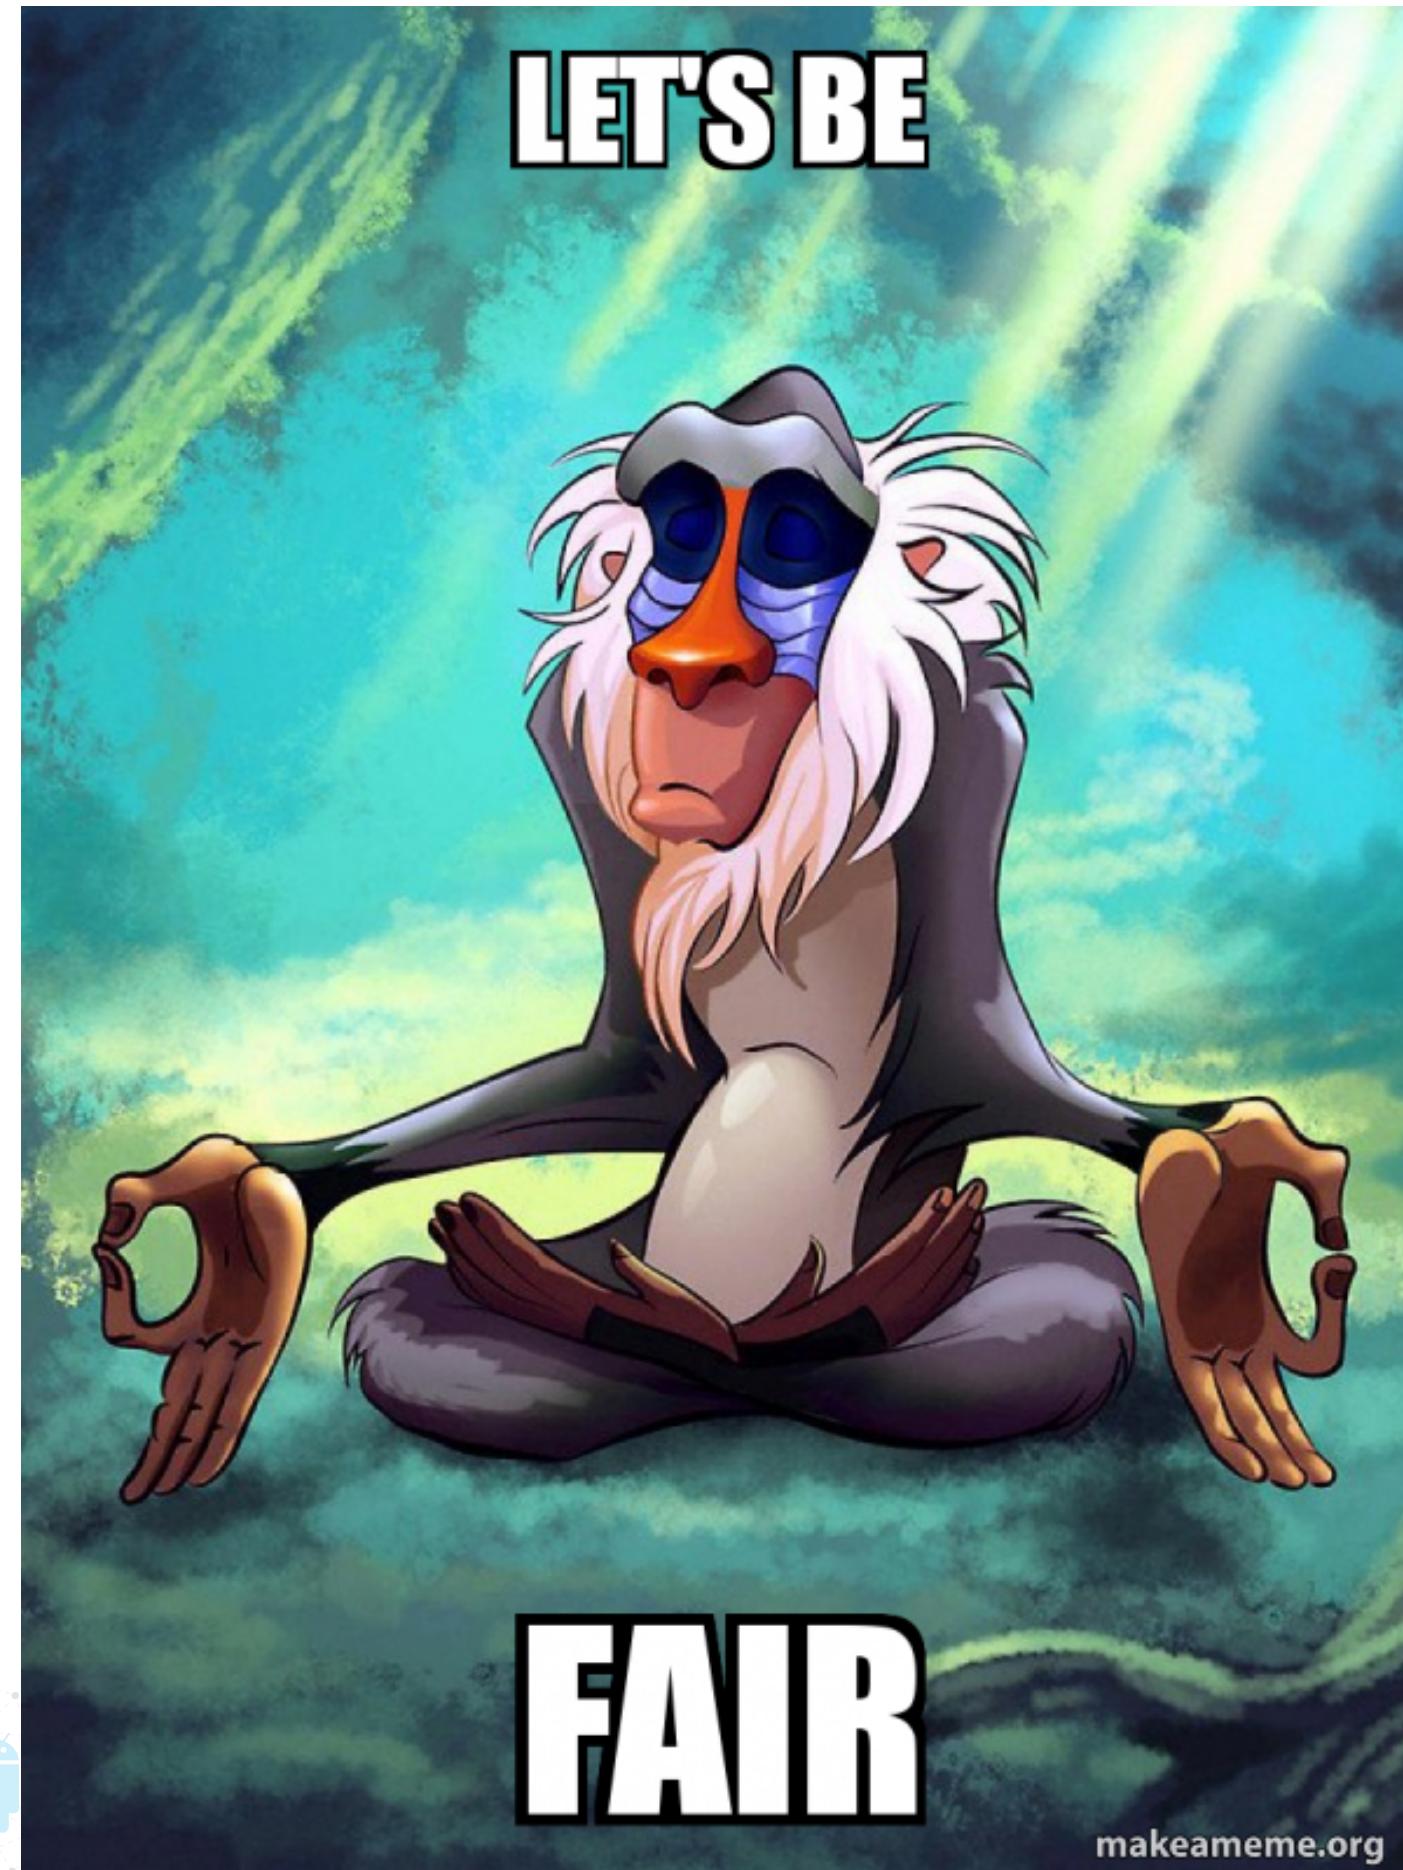
```
void lock(lock_t *lock)
{
    while (1)
    {
        while (LoadLinked(&lock->flag)==1)
            ; //Keep spinning

        if (StoreConditional(&lock->flag,1)==1)
        {
            //store is successful, retrun else repeat all again
            return;
        }
    }
}

void unlock(lock_t *lock)
{
    lock->flag = 0;
}
```

# What about Fairness?

Can we incorporate fairness through locks?



- Lock only has a flag variable
- Every time a thread acquires, it checks for flag
- However, which threads are checking is not recorded
- The threads that are looking for locks may have to be stored somewhere
- Can we store more information within each lock?

# What if we leverage Turns and Tickets

Think about going to some crowded office space



**Current Number and Window**



# Fetch and Add

Yet another hardware primitive but very powerful



## Fetch And Add

```
int FetchAndAdd( int *ptr )
{
    int old = *ptr;
    *ptr = old + 1;
    return old;    //returns old value
}
```

- Atomically increment a value while returning the old value at a particular address
- Used to build interesting type of lock - **The ticket lock**
- Instead of single variable a combination of ticket and turn variable is used
- Not just flag: ticket and turn



# Ticket Lock



Ticket Lock

```
typedef struct __lock_t
{
    int ticket;
    int turn;
} lock_t;

void lock_init (lock_t *lock)
{
    lock->ticket = 0;
    lock->turn = 0;
}
```



Ticket Lock

```
void lock (lock_t *lock)
{
    int myturn = FetchAndAdd(&lock->ticket);
    //when ticket value = my turn, thread goes into CS
    while (lock->turn != myturn)
    {
        // keep spinning
    }
}

void unlock (lock_t *lock)
{
    // next waiting thread can enter cs
    FetchAndAdd(&lock->turn);
}
```



# An Illustration of Ticket Lock

Four Processor Ticket Lock Example

Row	Action	next_ticket	now_serving	P1 my_ticket	P2 my_ticket	P3 my_ticket	P4 my_ticket
1	Initialized to 0	0	0	-	-	-	-
2	P1 tries to acquire lock (succeed)	1	0	0	-	-	-
3	P3 tries to acquire lock (fail + wait)	2	0	0	-	1	-
4	P2 tries to acquire lock (fail + wait)	3	0	0	2	1	-
5	P1 releases lock, P3 acquires lock	3	1	0	2	1	-
6	P3 releases lock, P2 acquires lock	3	2	0	2	1	-
7	P4 tries to acquire lock (fail + wait)	4	2	0	2	1	3
8	P2 releases lock, P4 acquires lock	4	3	0	2	1	3
9	P4 releases lock	4	4	0	2	1	3
10	...	4	4	0	2	1	3



Source: wikipedia article on ticket lock

# How good is the spin based locks?

- Simple hardware based locks are simple to implement and powerful
- They are also quite inefficient especially when it comes to performance
  - Consider that there are two threads and one thread has the lock
  - When thread has lock, it may get interrupted, the other thread spins for a time slice, waste CPU cycle
  - Think about N threads, N-1 threads might waste CPU cycles in spinning (especially if round robin)
- **Can we come up with something better instead of wasting cycles with spinning?**





**Thank you**

**Course site:** [karthikv1392.github.io/cs3301\\_osn](https://karthikv1392.github.io/cs3301_osn)

**Email:** [karthik.vaidhyanathan@iiit.ac.in](mailto:karthik.vaidhyanathan@iiit.ac.in)

**Twitter:** [@karthyishere](https://twitter.com/karthyishere)

