

CS3.301 Operating Systems and Networks

Persistence: File System Implementation

Karthik Vaidhyanathan

<https://karthikvaidhyanathan.com>

1



Acknowledgement

The materials used in this presentation have been gathered/adapted/generate from various sources as well as based on my own experiences and knowledge -- Karthik Vaidhyanathan

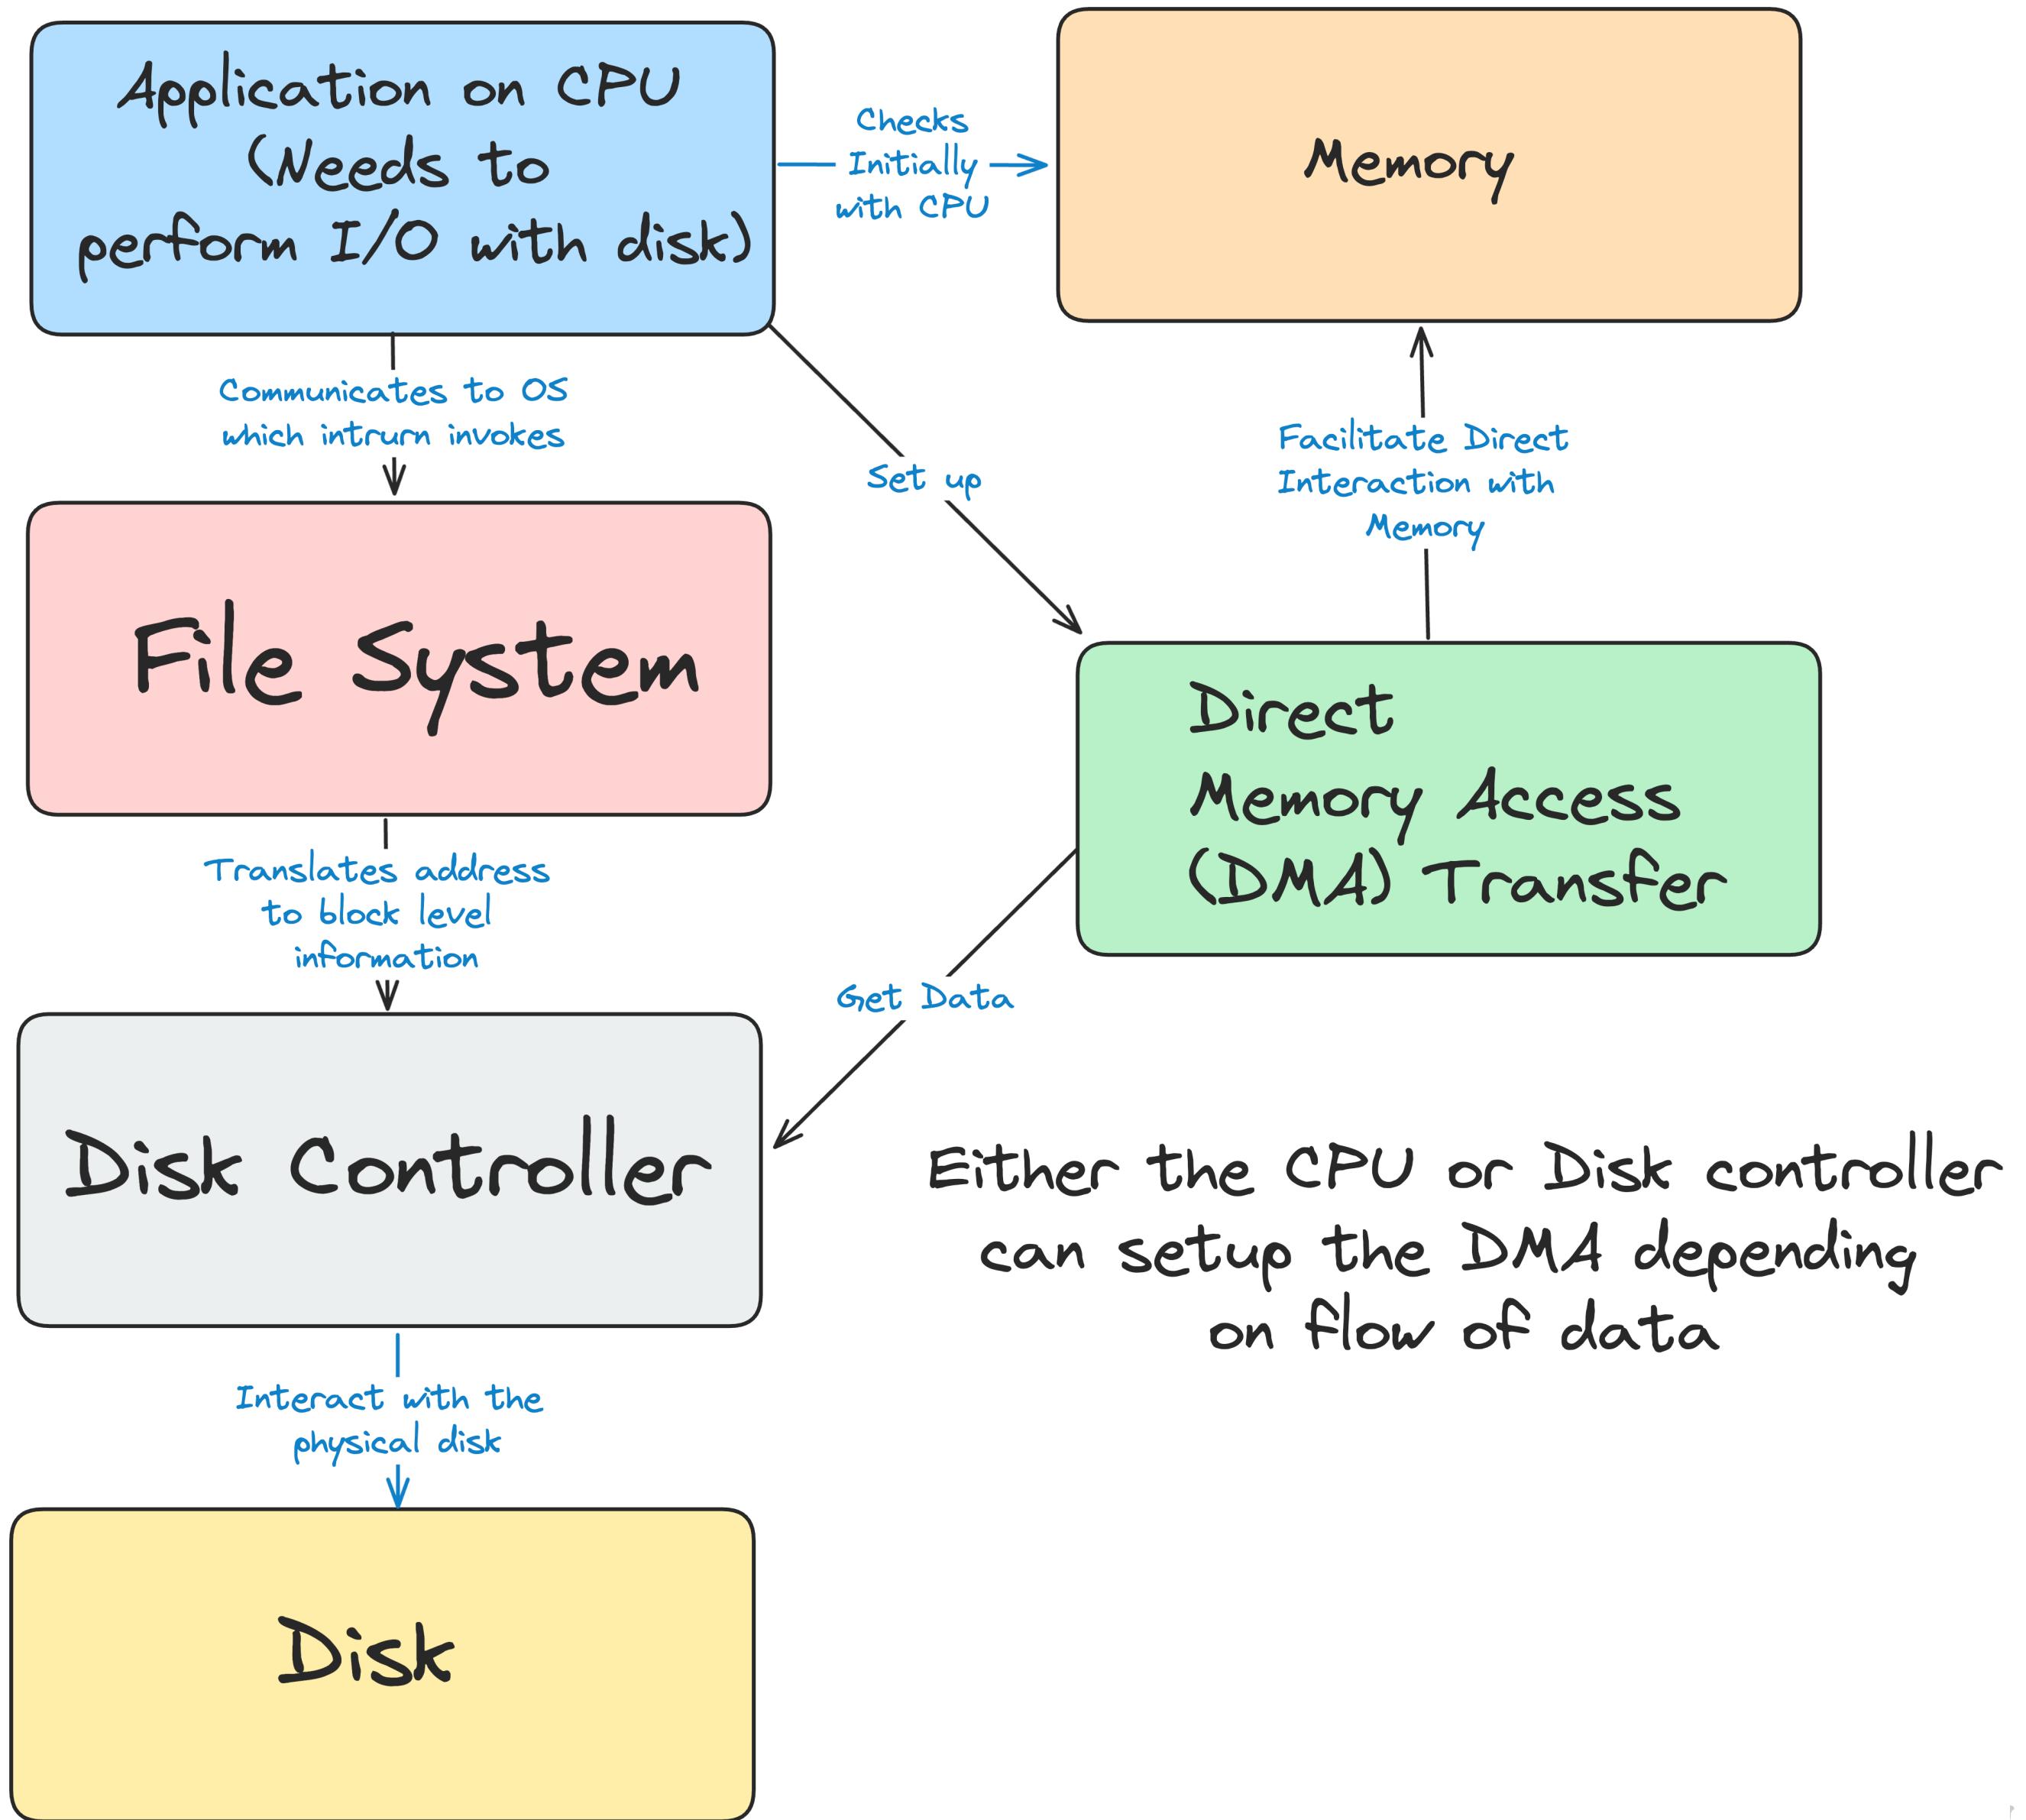
Sources:

- Operating Systems in Three Easy Pieces by Remzi et al.
- File System implementation by Youjip Won, Hanyang University



The flow of access

- Application performs read or write to a file
- CPU communicates to OS which invokes the File System (FS)
- The OS may check in its cache if its already there
- FS prepares block level information to disk controller
- A Direct Memory Access (DMA) is set up
- Disk controller performs the physical read or write based on commands from DMA and file system
- If its read, Disk -> DMA, for writes, DMA -> Disk



How can we build a simple File System?

What structures are needed in disk and how to access?



Breaking down into two main aspects

- Lets try building a simple file system - **Very Simple File System (VSFS)**
- In any FS, two key things make the difference

Data Structures

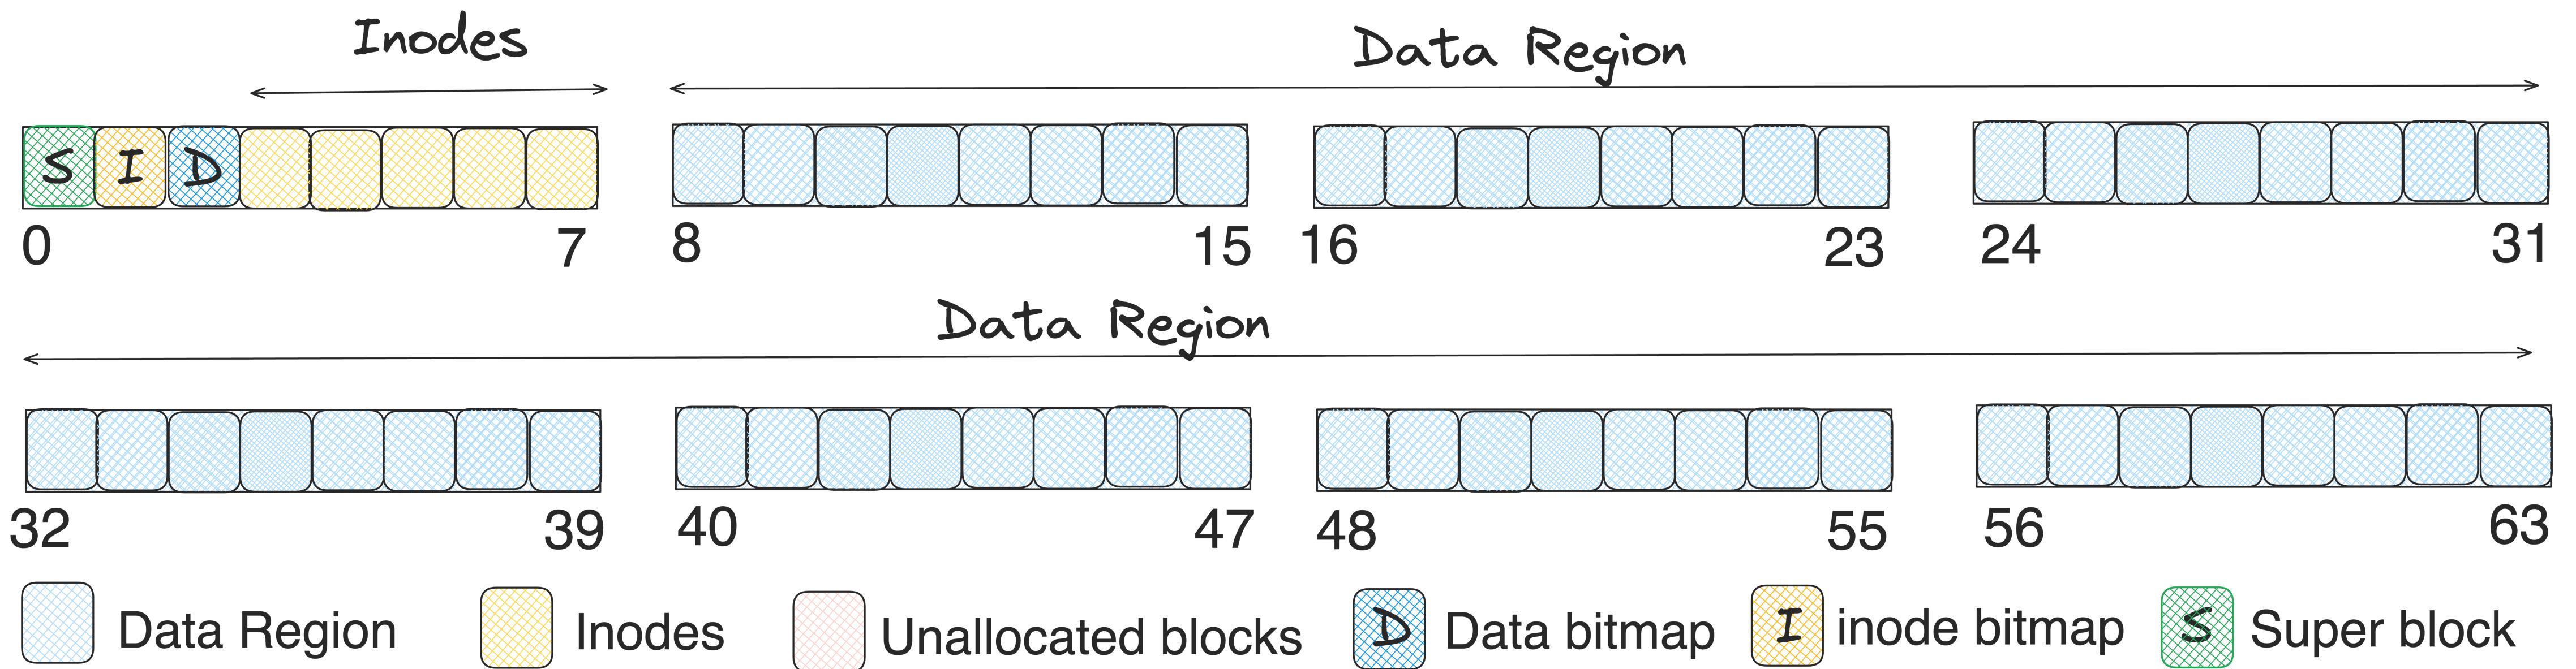
- What types of on-disk data structures are utilized by the file system to organise its data and metadata?
- VSFS can make use of simple structures like array of blocks (complex ones: trees)

Access Methods

- How can the calls like open(), read(), write(), etc made by process be mapped?
 - Which structures are read during the execution of a system call?
- What about the efficiency?



A more complete representation

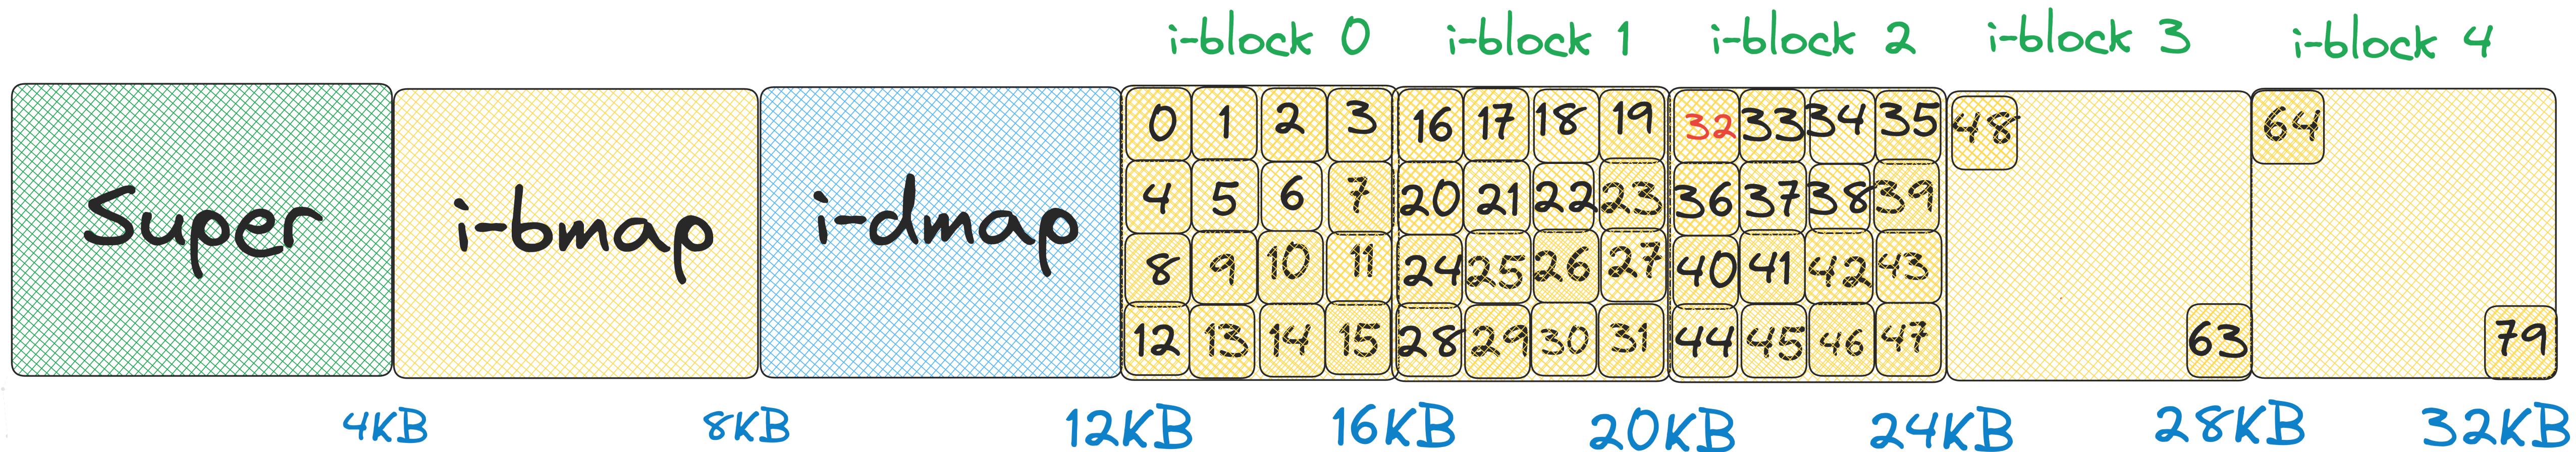


- **Super block** holds the entire organisation of all other blocks
 - Which blocks are inodes, which are data blocks, where does data block start, where Inode begins, type of file system, etc
 - During the mount, OS reads super block to initialise various parameters



File Organization: The inode

- Each inode is referred to by the inode number
 - Using inode number, FS can locate inode, eg: inode number: 32
 - Calculate offset into inode: $32 \times (\text{sizeof(inode)}) = 32 * 256 = 8192 \Rightarrow 8 \text{ KB}$
 - Add offset with start address of inode = 12KB + 8KB = **20KB**



What does inode contain?

- inode contains all the information about a file - The metadata
 - File type (regular file, directory, etc.)
 - Size, number of blocks allocated to it
 - Protection information (who can access, what access, etc.)
 - Time information (modified time, access time, etc)
 - Many more

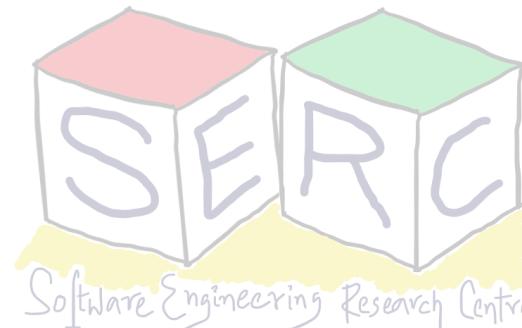


Simplified EXT2 inode

Size	Name	What is this inode field for?
2	mode	can this file be read/written/executed?
2	uid	who owns this file?
4	size	how many bytes are in this file?
4	time	what time was this file last accessed?
4	ctime	what time was this file created?
4	mtime	what time was this file last modified?
4	dtime	what time was this inode deleted?
4	gid	which group does this file belong to?
2	links_count	how many hard links are there to this file?
2	blocks	how many blocks have been allocated to this file?
4	flags	how should ext2 use this inode?
4	osd1	an OS-dependent field
60	block	a set of disk pointers (15 total)
4	generation	file version (used by NFS)
4	file_acl	a new permissions model beyond mode bits
4	dir_acl	called access control lists
4	faddr	an unsupported field
12	i_osd2	another OS-dependent field

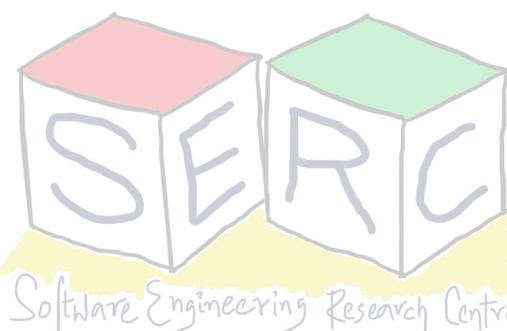
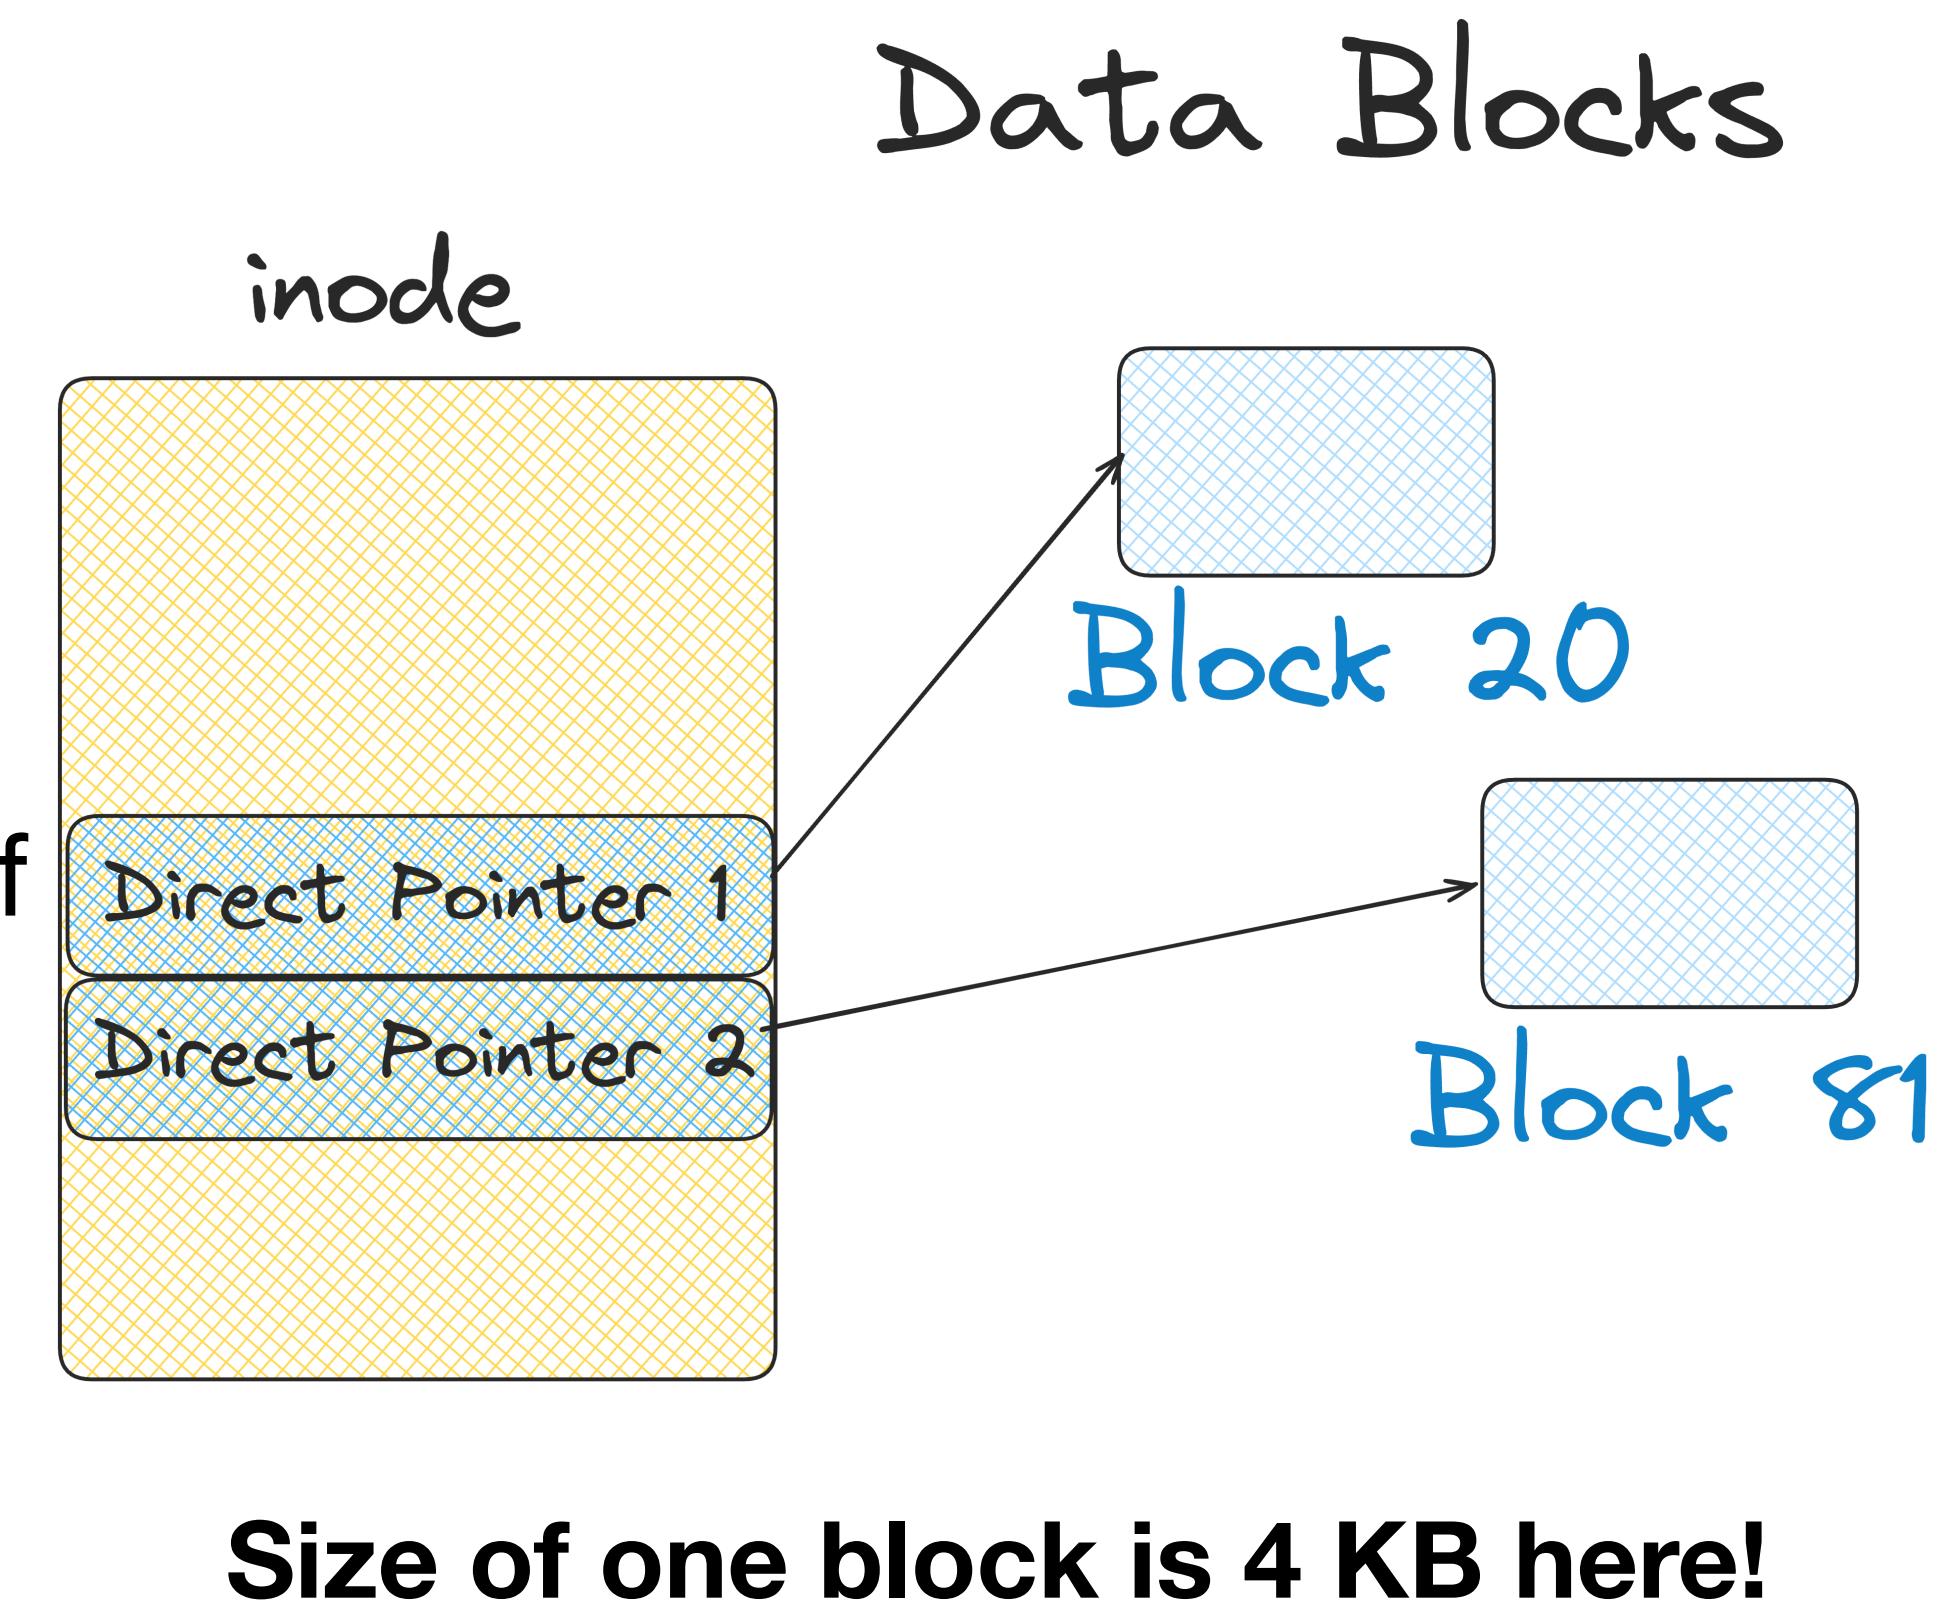
Total 128 bytes

**How can inode get
to data blocks?**



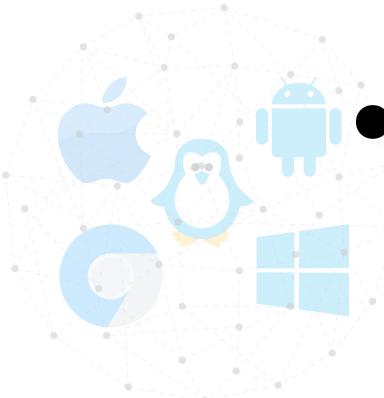
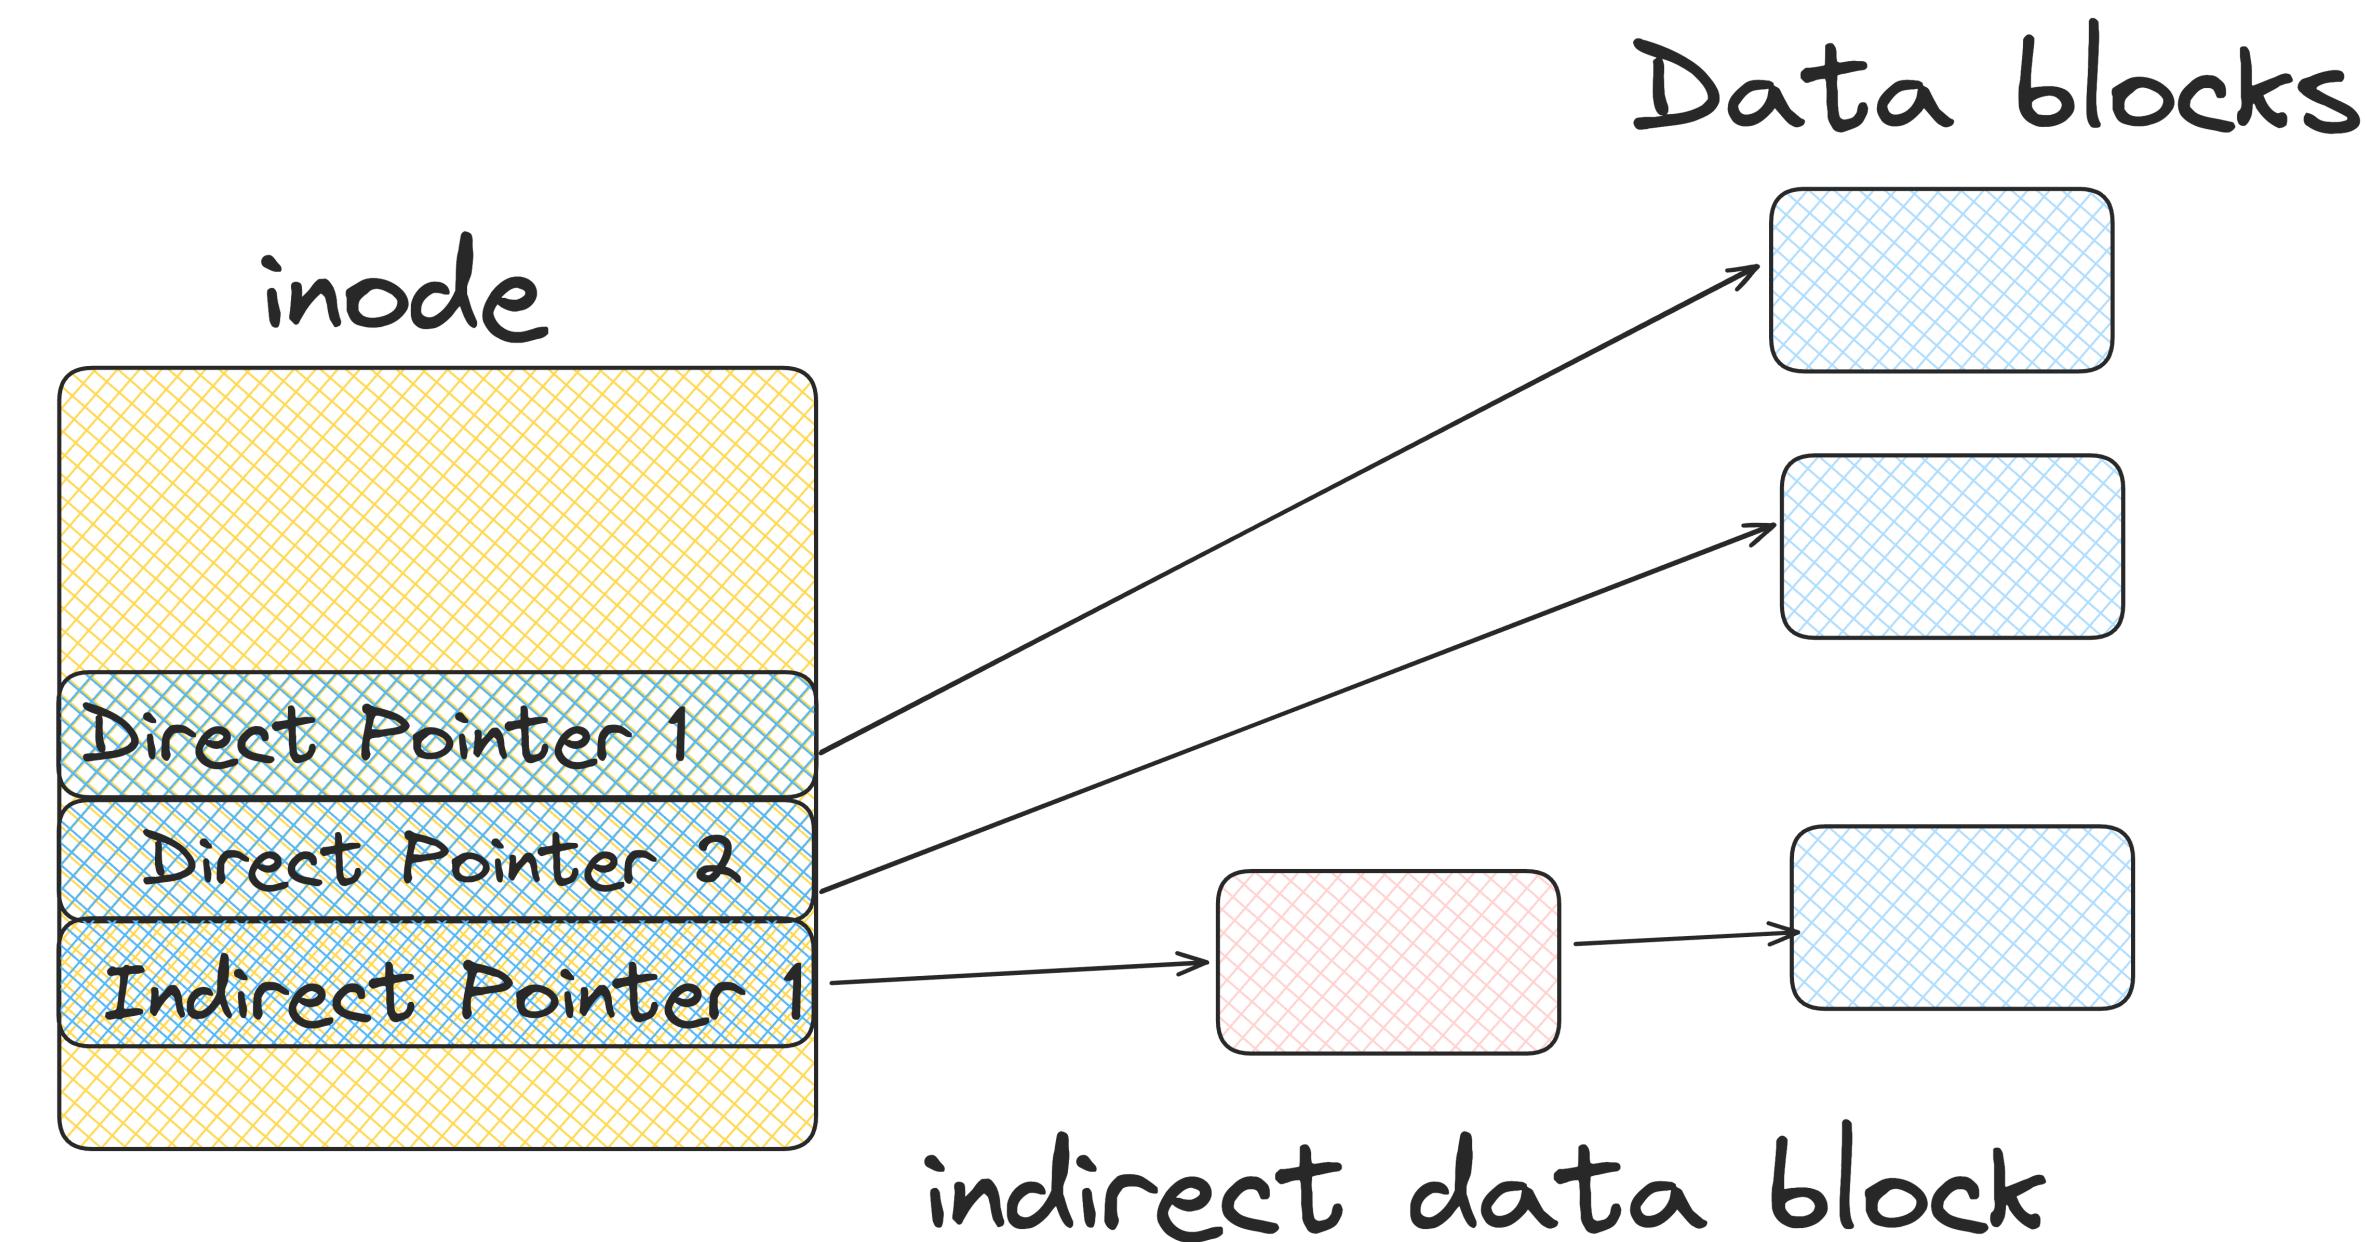
More about inodes

- Each inode needs to track disk block numbers of a file
- File data is not stored contiguously on disk
 - How to track multiple block numbers of a file?
 - Store pointer to the block inside the inode
 - Numbers of first few blocks are stored in inode itself
 - Each pointer can point to the location in the disk block - **direct pointers**
 - **What if the file size is large?** - How many block numbers can i-node store?
 - Need for better mechanism



Indirect Pointers

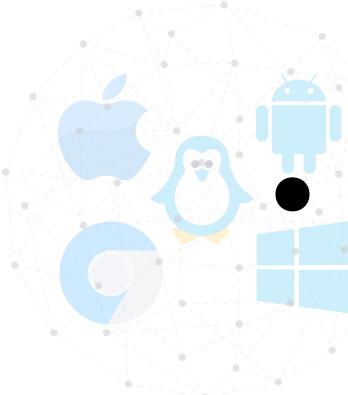
- To support large files, few direct pointers may not suffice!
- Use a special pointer - **indirect pointer**
 - Point to a block that contains more pointers - **indirect data block**
 - Each of the pointer can further point to data blocks
 - The indirect block is allocated from the data region
- Inode array may have 12 direct pointers and one indirect pointer



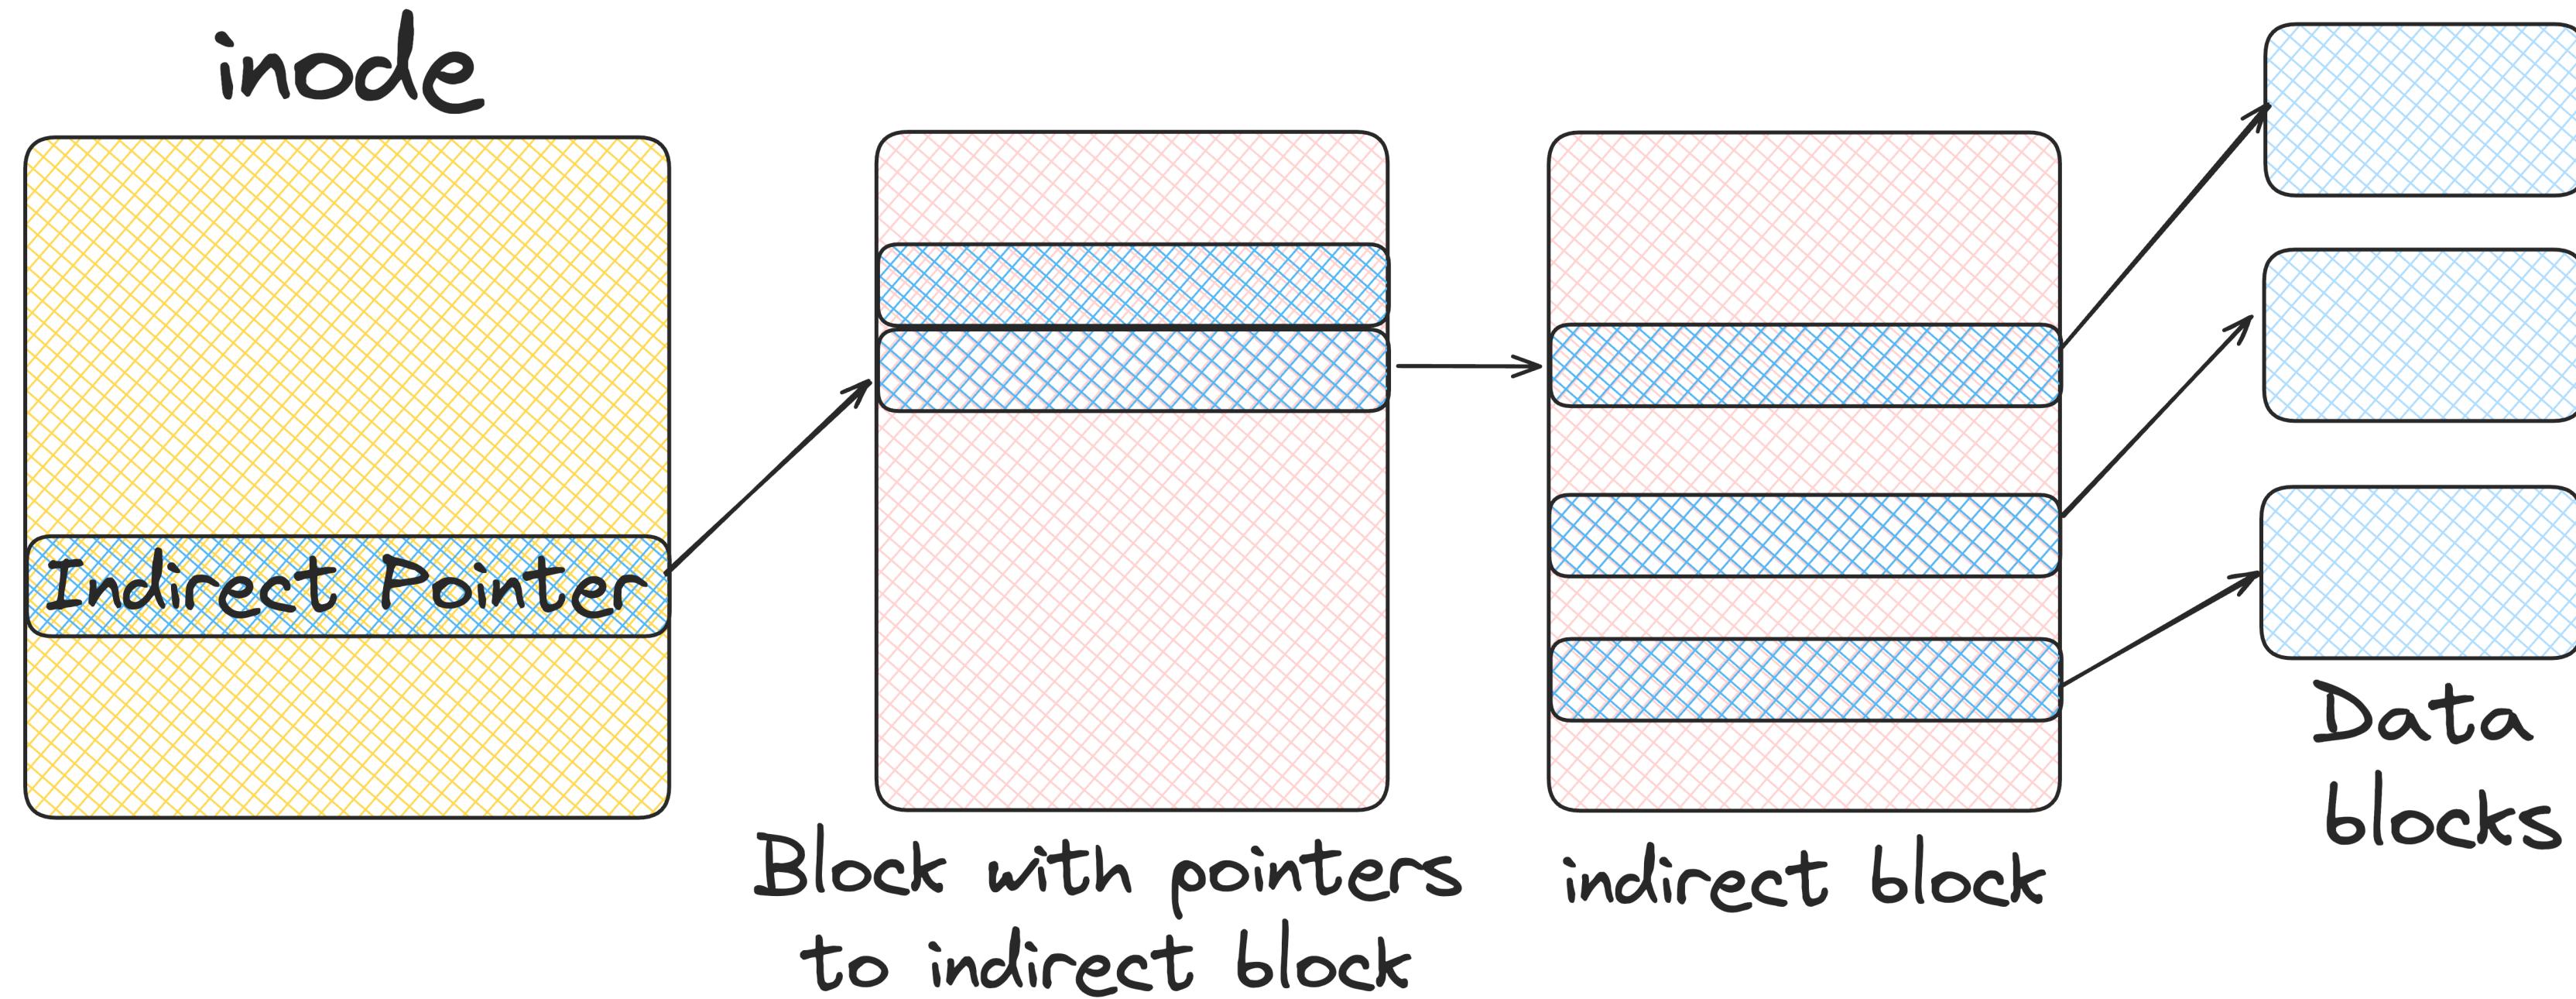
How much files can be supported?

Having one indirect pointer

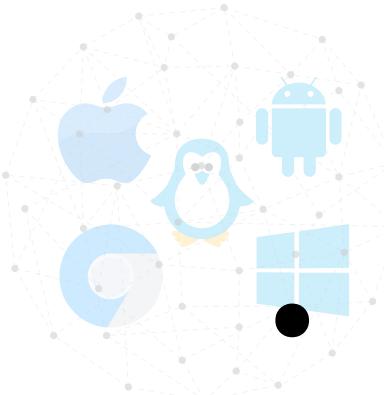
- Each block is 4 KB
- Each inode can contain 12 direct pointers => $12 * 4 = 48$ KB of file can be addressed
- 1 indirect pointer points to a block of size 4 KB
 - Each address takes around 4 bytes
 - Indirect blocks can have around 1024 pointers ($4\text{ KB} / 4$)
 - Total size of file that can be addressed = $(12 + 1024) * 4\text{K} = 4144$ KB
- What if the file is even larger? How can the inode capture all the blocks?



The Multi-Level Index



- **Double indirect pointer:** Points to a block with pointers to indirect block
 - Each of the pointers in indirect block points to data blocks
 - Size now that can be supported is $1024 \times 1024 \times 4 \sim 4\text{GB}$
- For more even **triple indirect pointers** can be sought of



Why this direct and indirect pointers?

- One finding over many years of research: most of files are small
- Thus with small number of direct pointers, inode can point to 48 KB of data
- All that is needed is one or few indirect blocks

Most files are small	~2K is the most common size
Average file size is growing	Almost 200K is the average
Most bytes are stored in large files	A few big files use most of space
File systems contain lots of files	Almost 100K on average
File systems are roughly half full	Even as disks grow, file systems remain ~50% full
Directories are typically small	Many have few entries; most have 20 or fewer



“A Five-Year Study of File-System Metadata” by Nitin Agrawal, William J. Bolosky, John R. Douceur, Jacob R. Lorch. FAST ’07, San Jose, California, February 2007.

What about Directories?

- Directory stores the mapping of file names and their inode numbers
- Each directory has two extra files
 - “.” for current directory and “..” for parent directory
- Assume that a directory “OSN” has three files (I01, I02, lect03)
- Directory is a special type of file and has inode and data blocks (stores file records)

inum	reclen	strlen	name
5	12	2	.
2	12	3	..
12	12	4	I01
13	12	4	I02
24	36	7	lect03

inum - inode number

reclen - total bytes for name

strlen - length of the name

name - actual name



Free Space Management

- FS has to keep track of which inodes and data blocks are free
- Multiple methods can be used and many design choices exist. Eg:
 - Use **bitmaps** for inodes and data blocks, store one bit per block to indicate free or not
 - **Free list:** Super block can store pointer to first free block which can then point to next free block and so on.
- Eg: Linux FS such as ext2 and ext3 checks for sequence of blocks on new file creation
 - Sequence of data blocks are allocated contiguously for performance
 - Pre-allocation policy is commonly used heuristic when allocating data blocks



Access: Reading File From Disks

- FS also needs better ways of managing access to file (apart from data structure)
- Eg: FS has been mounted and read issued to */OSN/I01* - open, read, close
- Assume that file size is 12 KB (3 blocks in size)
 - sys call `open("/OSN/I01", O_RDONLY)`
- Intuitively: FS must traverse the pathname and locate the file
 - What will be the process to achieve this?



Opening Files

- First part of read is always open sys call - Why?
 - Take the inode and load it in the memory for future operations
 - Open returns file descriptor which points to in-memory I-node
 - Reads and writes can access file data from I-node
- Assume a sys call `open("/OSN/lectures/I01.txt", O_RDONLY)`
 - Traverse the path name and then locate desired inode
 - Begin at the root of the FS (/), root inode number is 2 in Unix FS (mostly)
 - FS reads the block that contains inode number 2



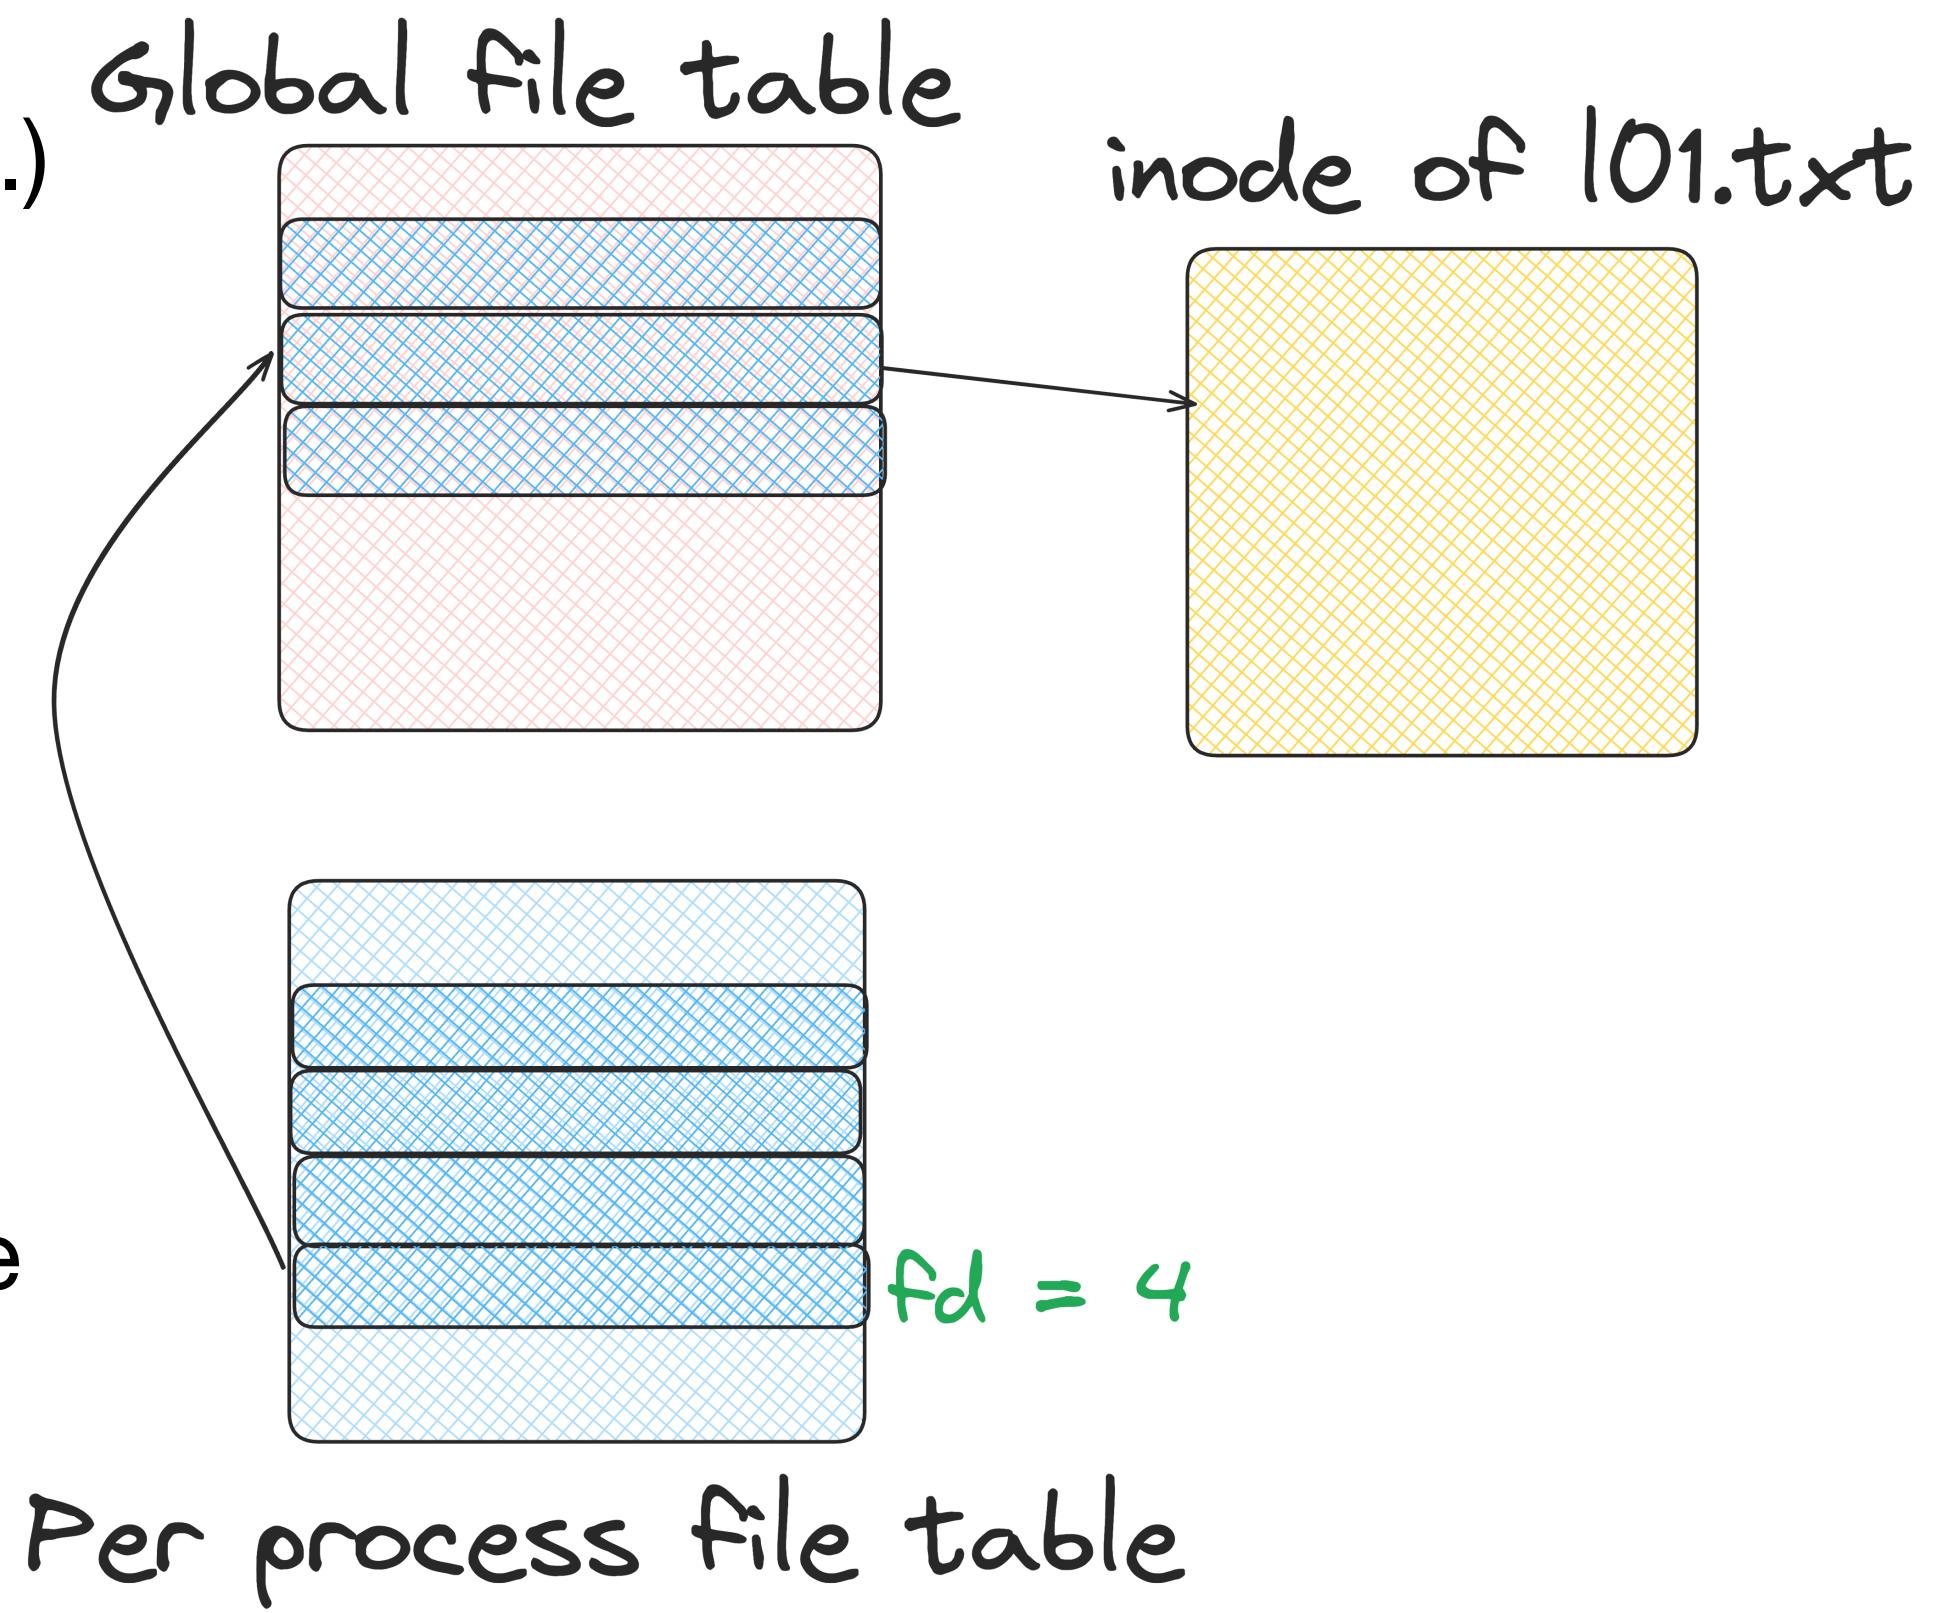
Opening Files

- Recursively: Read the data blocks of root directory, find the name “lectures” and get its inode number
 - Get inode of lectures -> get inode number of “I01.txt” -> get inode
 - Keep repeating the process until the end of the path
- Read inode of “I01.txt” into memory, make final permission check
- Allocate file descriptor for this process and return file descriptor to user
 - Allocation will be done in the in-memory **open file table**. It will be updated for each read - offset
- In the case of new file, new inode and data blocks will be allocated using bitmap and update directory entry



Open File Table

- Kernel uses a set of data structures to track all open files
- **Global open file table**
 - One entry for every open file (stores also sockets, pipes, etc.)
 - Entry points to the in-memory inode of the file (remember opening of file)
- **Per-process open file table**
 - Array of all the files that the process has opened
 - File descriptor is index into the array
 - Per process file entry -> global file table entry -> inode of file
 - Every process has three files (stdin, stdout, err) open by default
- Open system call creates entries in both table and returns file descriptor number

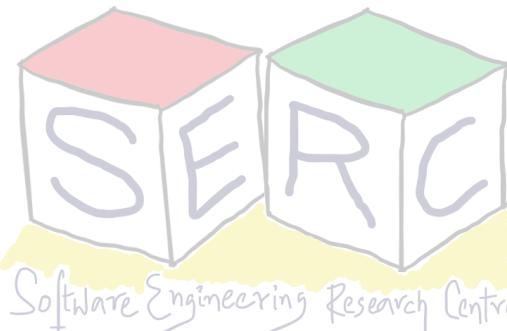
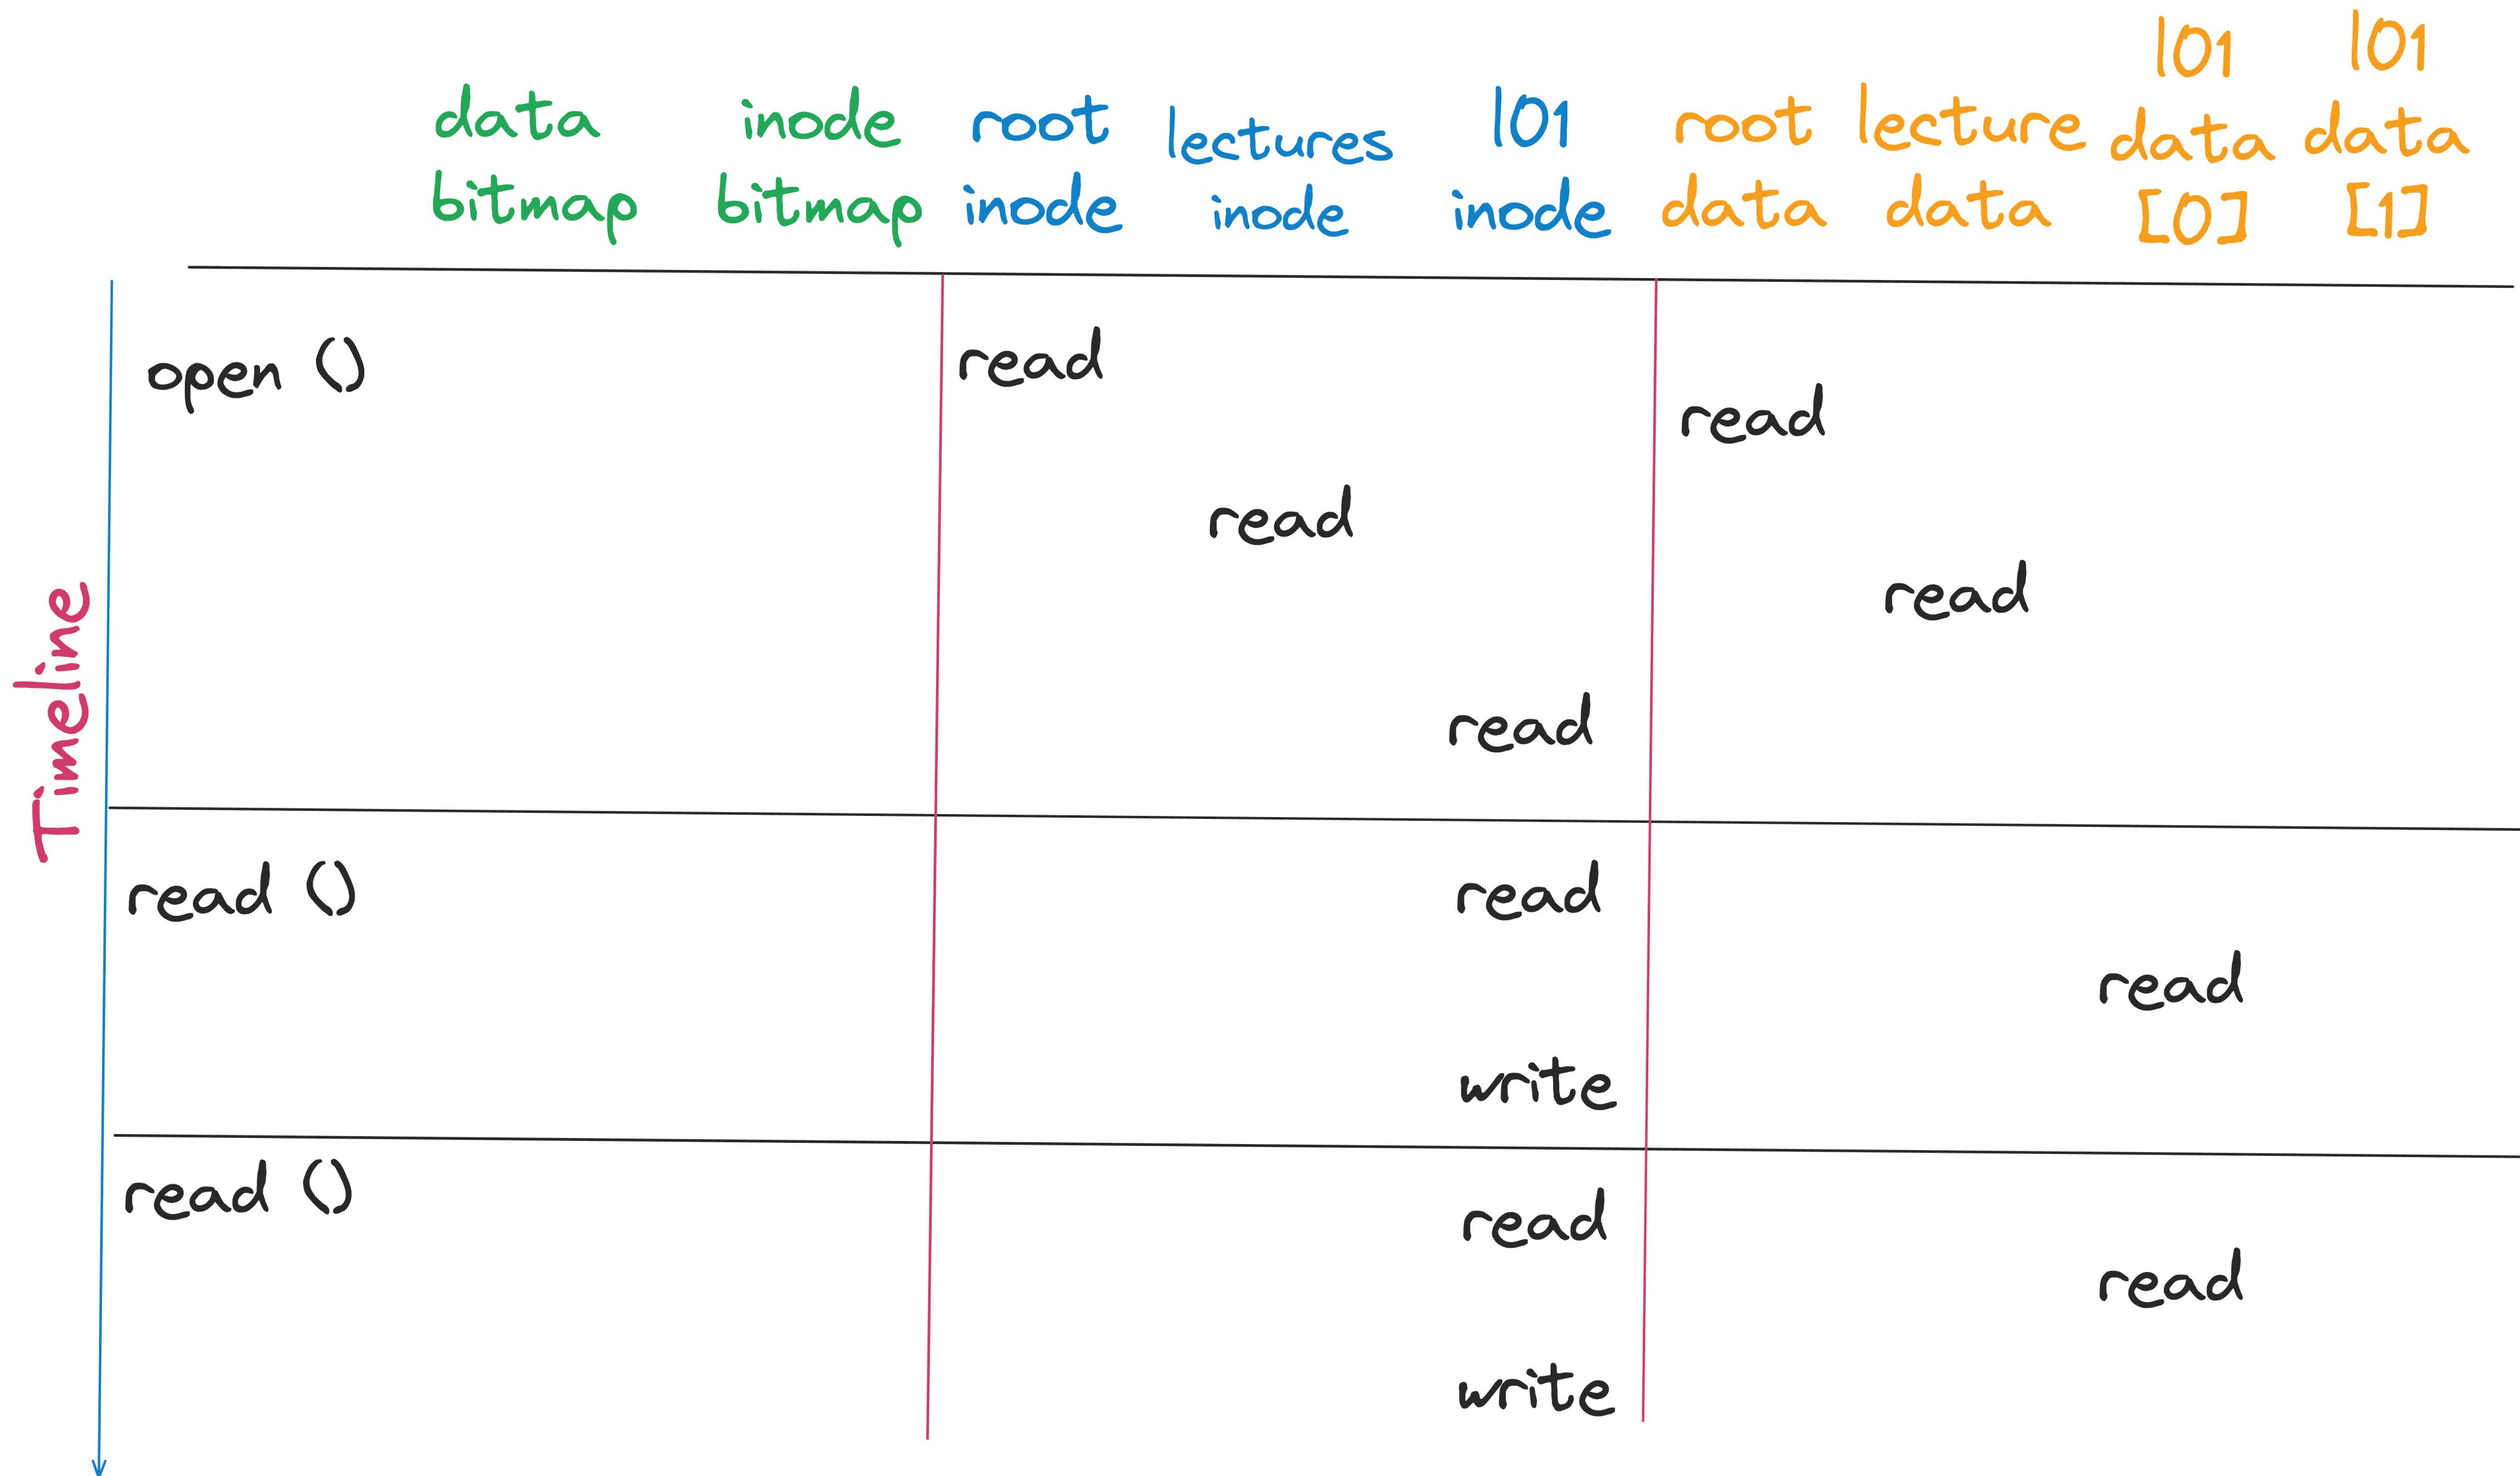


Reading a File

- Make a call read() to read from file
 - Read in the first data block of the file with help of inode
 - Update the inode with last accessed time
 - Update in-memory open file table for file descriptor, file offset
 - Repeat the process for reading each block of data
- Once file is closed
 - Just the file descriptor should be deallocated - No disk I/O



Reading a File From Disk



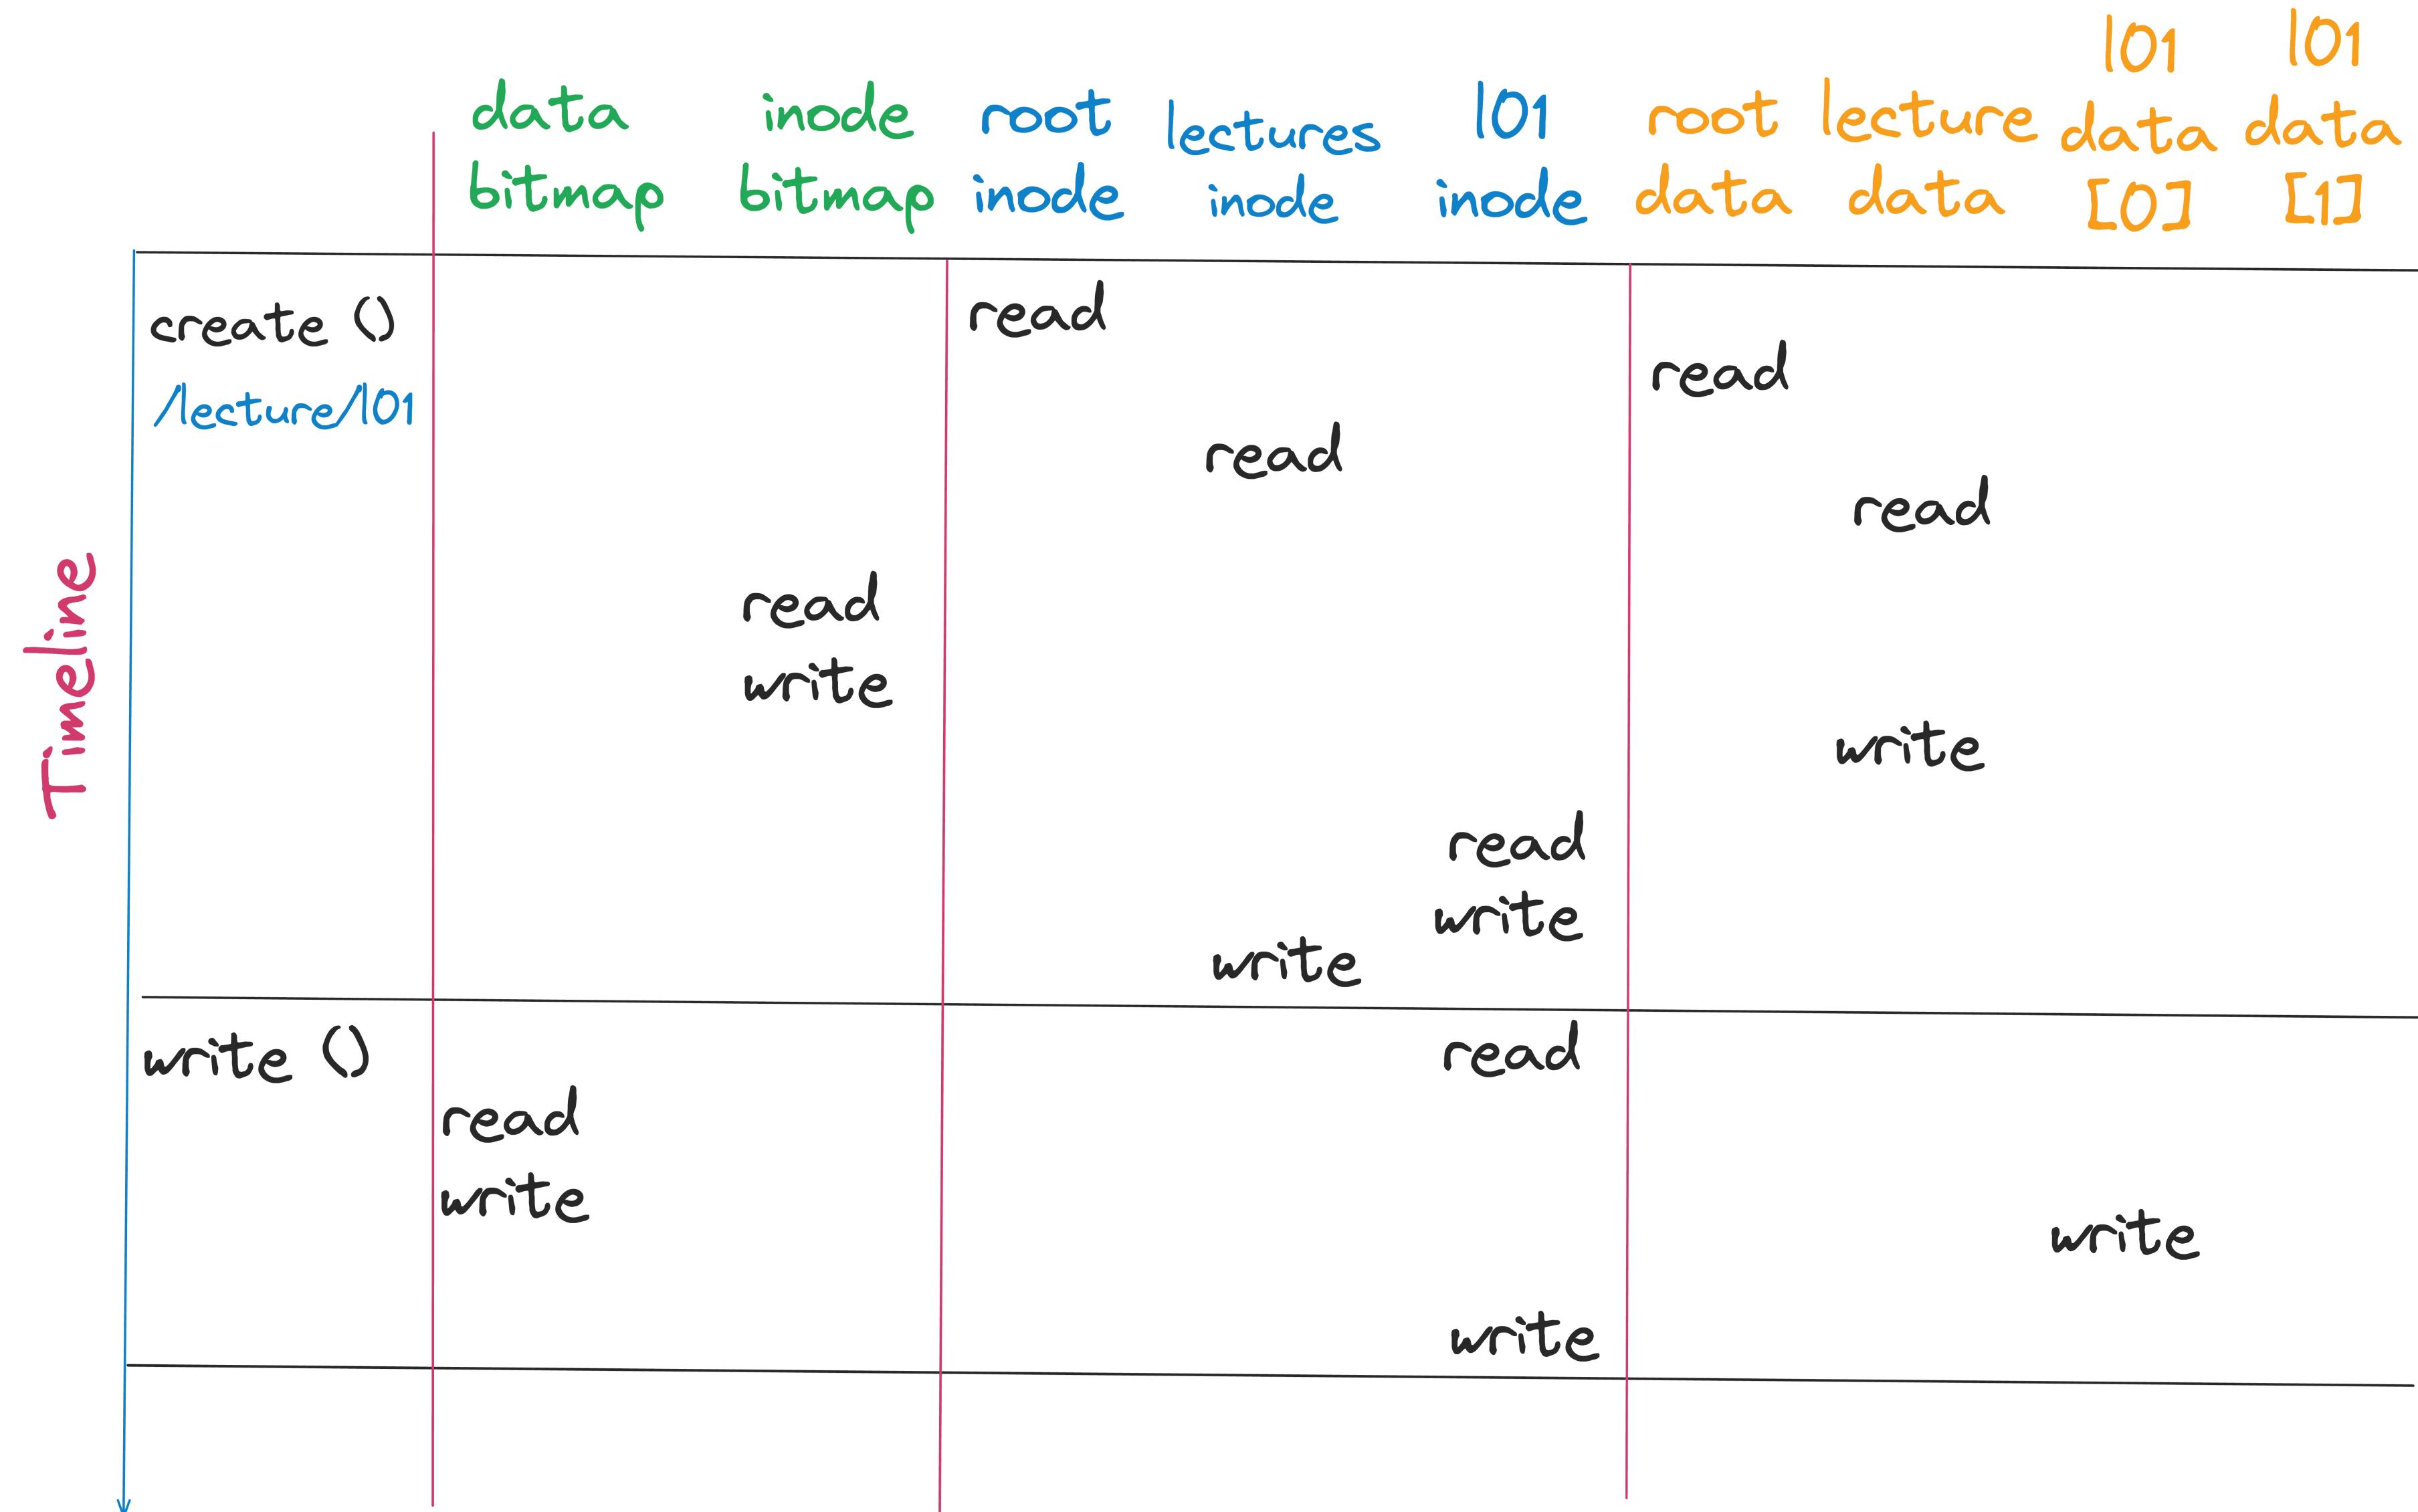
Writes to a File

- Make a call write() to write into the file on the disk
- Data block may have to be allocated (if not overwriting)
 - Need to update data bitmap and data block
 - Total of five I/O:
 - One to read data bitmap
 - Write to data bitmap
 - Two more to read and write the inode
 - Write to the actual block itself

• In case of creation of new file, number of I/Os can go really high!



Writing a File To Disk



Can we do something about performance?

- Reading and writing files are expensive
- Imagine opening and reading a file by providing a long path
 - Each inode needs to be fetched, corresponding data then read of files
 - Can go upto 100s of I/Os
- Use the concept of caching and buffering
 - Use system memory to cache important blocks - **Minimise overheads!**
 - Early FS, used **fixed-size cache** -> store popular blocks (10% at boot time)
 - Use strategies like LRU to evict blocks



Caching and Buffering

- Static partitioning of memory is not always useful - **Wastages!**
- Modern systems employ dynamic partitioning approach
 - Integrate virtual memory pages and FS pages into unified page cache
 - First open may generate lot of I/O but subsequent will be in cache!
- Writes is little tricky as at some point the disk has to be accessed to store
 - **Write buffering** - Delay writes to disk, perform batch I/O
 - Schedule I/Os in a particular order for performance gain
 - Writes can be avoided totally - file is created and deleted in few seconds!
(Don't write)



Caching and Buffering

- Applications like DB avoids caching altogether - direct I/O
 - System calls like fsync() allows writes to be pushed immediately
 - Unexpected data loss may happen since data is in memory
 - Has impact on overall system performance
- At the end its all about trade-off's
 - Durability vs Performance tradeoff
 - Has big dependance on the application
 - Browser vs Transactional database!





Thank you

Course site: karthikv1392.github.io/cs3301_osn
Email: karthik.vaidhyanathan@iiit.ac.in
Twitter: [@karthi_ishere](https://twitter.com/karthyishere)

