

Probability and Statistics

Homework 11 - Solutions

Q1: Let $X_1, X_2, X_3, \dots, X_n$ be a random sample from the following distribution

$$f_X(x) = \begin{cases} \theta \left(x - \frac{1}{2} \right) + 1 & \text{for } 0 \leq x \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

where $\theta \in [-2, 2]$ is an unknown parameter. We define the estimator $\hat{\Theta}_n$ as

$$\hat{\Theta}_n = 12\bar{X} - 6$$

to estimate θ .

- a. Is $\hat{\Theta}_n$ an unbiased estimator of θ ?
- b. Is $\hat{\Theta}_n$ a consistent estimator of θ ?
- c. Find the mean squared error (MSE) of $\hat{\Theta}_n$.

Solution:

To answer all three questions, we first need to find the expected value of X , $\mathbb{E}[X]$.

$$\begin{aligned} \mathbb{E}[X] &= \int_{-\infty}^{\infty} xf_X(x) dx \\ &= \int_0^1 x \left[\theta \left(x - \frac{1}{2} \right) + 1 \right] dx \\ &= \int_0^1 \left(\theta x^2 - \frac{\theta}{2}x + x \right) dx \\ &= \left[\frac{\theta x^3}{3} - \frac{\theta x^2}{4} + \frac{x^2}{2} \right]_0^1 \\ &= \left(\frac{\theta}{3} - \frac{\theta}{4} + \frac{1}{2} \right) - (0) \\ &= \frac{4\theta - 3\theta}{12} + \frac{1}{2} \\ &= \frac{\theta}{12} + \frac{1}{2} = \frac{\theta + 6}{12} \end{aligned}$$

a. Unbiasedness

An estimator $\hat{\Theta}_n$ is unbiased if $\mathbb{E}[\hat{\Theta}_n] = \theta$. We compute the expected value of the estimator.

$$\begin{aligned} \mathbb{E}[\hat{\Theta}_n] &= \mathbb{E}[12\bar{X} - 6] \\ &= 12\mathbb{E}[\bar{X}] - 6 \end{aligned}$$

Since X_1, \dots, X_n are i.i.d., $\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{n} \sum X_i\right] = \frac{1}{n} \sum \mathbb{E}[X_i] = \mathbb{E}[X]$.

$$\begin{aligned}\mathbb{E}[\hat{\Theta}_n] &= 12\mathbb{E}[X] - 6 \\ &= 12\left(\frac{\theta + 6}{12}\right) - 6 \\ &= (\theta + 6) - 6 \\ &= \theta\end{aligned}$$

Since $\mathbb{E}[\hat{\Theta}_n] = \theta$, yes, $\hat{\Theta}_n$ is an unbiased estimator of θ .

b. Consistency

An estimator is consistent if it converges in probability to the true parameter as $n \rightarrow \infty$. That is, $\lim_{n \rightarrow \infty} P(|\hat{\Theta}_n - \theta| \geq \epsilon) = 0$.

By the Weak Law of Large Numbers (WLLN), the sample mean converges in probability to the population mean:

$$\bar{X} \xrightarrow{p} \mathbb{E}[X]$$

We know $\hat{\Theta}_n = g(\bar{X})$, where $g(t) = 12t - 6$. Since $g(t)$ is a continuous function, we can apply the Continuous Mapping Theorem:

$$\hat{\Theta}_n = g(\bar{X}) \xrightarrow{p} g(\mathbb{E}[X])$$

Substituting the value of $\mathbb{E}[X]$ we found:

$$g(\mathbb{E}[X]) = 12\mathbb{E}[X] - 6 = 12\left(\frac{\theta + 6}{12}\right) - 6 = (\theta + 6) - 6 = \theta$$

Therefore, $\hat{\Theta}_n \xrightarrow{p} \theta$. Yes, $\hat{\Theta}_n$ is a consistent estimator of θ .

c. Mean Squared Error (MSE)

The Mean Squared Error (MSE) is defined as $MSE(\hat{\Theta}_n) = \mathbb{E}[(\hat{\Theta}_n - \theta)^2]$. This can be decomposed as:

$$MSE(\hat{\Theta}_n) = \text{Var}(\hat{\Theta}_n) + (\text{Bias}(\hat{\Theta}_n))^2$$

From part (a), we know the estimator is unbiased, so $\text{Bias}(\hat{\Theta}_n) = \mathbb{E}[\hat{\Theta}_n] - \theta = 0$. Therefore, $MSE(\hat{\Theta}_n) = \text{Var}(\hat{\Theta}_n)$.

$$\begin{aligned}\text{Var}(\hat{\Theta}_n) &= \text{Var}(12\bar{X} - 6) \\ &= 12^2 \text{Var}(\bar{X}) \\ &= 144 \text{Var}\left(\frac{1}{n} \sum X_i\right) \\ &= \frac{144}{n^2} \sum \text{Var}(X_i) \quad (\text{since } X_i \text{ are independent}) \\ &= \frac{144}{n^2} (n \cdot \text{Var}(X)) \quad (\text{since } X_i \text{ are identically distributed}) \\ &= \frac{144}{n} \text{Var}(X)\end{aligned}$$

Now, we must find $\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$. First, we find $\mathbb{E}[X^2]$:

$$\begin{aligned}\mathbb{E}[X^2] &= \int_0^1 x^2 f_X(x) dx \\ &= \int_0^1 x^2 \left[\theta \left(x - \frac{1}{2} \right) + 1 \right] dx \\ &= \int_0^1 \left(\theta x^3 - \frac{\theta}{2} x^2 + x^2 \right) dx \\ &= \left[\frac{\theta x^4}{4} - \frac{\theta x^3}{6} + \frac{x^3}{3} \right]_0^1 \\ &= \frac{\theta}{4} - \frac{\theta}{6} + \frac{1}{3} \\ &= \frac{3\theta - 2\theta}{12} + \frac{4}{12} = \frac{\theta + 4}{12}\end{aligned}$$

Now, we calculate the variance:

$$\begin{aligned}\text{Var}(X) &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2 \\ &= \left(\frac{\theta + 4}{12} \right) - \left(\frac{\theta + 6}{12} \right)^2 \\ &= \frac{\theta + 4}{12} - \frac{\theta^2 + 12\theta + 36}{144} \\ &= \frac{12(\theta + 4)}{144} - \frac{\theta^2 + 12\theta + 36}{144} \\ &= \frac{(12\theta + 48) - (\theta^2 + 12\theta + 36)}{144} \\ &= \frac{12\theta + 48 - \theta^2 - 12\theta - 36}{144} \\ &= \frac{12 - \theta^2}{144}\end{aligned}$$

(Note: This variance is non-negative since $\theta \in [-2, 2]$, so $\theta^2 \in [0, 4]$, and $12 - \theta^2 \geq 8$.)

Finally, we substitute this back into our MSE equation:

$$\begin{aligned}\text{MSE}(\hat{\Theta}_n) &= \frac{144}{n} \text{Var}(X) \\ &= \frac{144}{n} \left(\frac{12 - \theta^2}{144} \right) \\ &= \frac{12 - \theta^2}{n}\end{aligned}$$

The Mean Squared Error is $\text{MSE}(\hat{\Theta}_n) = \frac{12 - \theta^2}{n}$.

Q2: Let X_1, \dots, X_4 be a random sample from a *Geometric(p)* distribution. Suppose we observed $(x_1, x_2, x_3, x_4) = (2, 3, 3, 5)$. Find the likelihood function using

$$P(X_i = x_i; p) = p(1 - p)^{x_i - 1}$$

as the PMF.

The likelihood function, $L(p)$, is the joint probability mass function (PMF) of the observed data, treated as a function of the parameter p .

Since X_1, \dots, X_4 is a random sample, the observations are independent and identically distributed (i.i.d.). Therefore, the likelihood function is the product of the individual PMFs:

$$\begin{aligned} L(p) &= P(X_1 = x_1, X_2 = x_2, X_3 = x_3, X_4 = x_4 \mid p) \\ &= \prod_{i=1}^4 P(X_i = x_i \mid p) \end{aligned}$$

Now, we substitute the given PMF for the Geometric distribution, $P(X_i = x_i; p) = p(1 - p)^{x_i - 1}$, into the product:

$$\begin{aligned} L(p) &= \prod_{i=1}^4 [p(1 - p)^{x_i - 1}] \\ &= [p(1 - p)^{x_1 - 1}] \cdot [p(1 - p)^{x_2 - 1}] \cdot [p(1 - p)^{x_3 - 1}] \cdot [p(1 - p)^{x_4 - 1}] \end{aligned}$$

We can group the p terms and the $(1 - p)$ terms together:

$$\begin{aligned} L(p) &= (p \cdot p \cdot p \cdot p) \cdot (1 - p)^{(x_1 - 1) + (x_2 - 1) + (x_3 - 1) + (x_4 - 1)} \\ &= p^4 \cdot (1 - p)^{(x_1 + x_2 + x_3 + x_4) - 4} \\ &= p^4 \cdot (1 - p)^{\sum_{i=1}^4 x_i - 4} \end{aligned}$$

We are given the observed data $(x_1, x_2, x_3, x_4) = (2, 3, 3, 5)$. We can now plug these values into our expression.

The sample size is $n = 4$.

The sum of the observations is:

$$\sum_{i=1}^4 x_i = 2 + 3 + 3 + 5 = 13$$

Plugging $n = 4$ and the sum $\sum x_i = 13$ into the likelihood function:

$$\begin{aligned} L(p) &= p^4(1 - p)^{13 - 4} \\ &= p^4(1 - p)^9 \end{aligned}$$

The likelihood function is $L(p) = p^4(1 - p)^9$, for $p \in (0, 1)$.

- Q3: A data scientist is analyzing a large stream of incoming data, X_1, X_2, \dots, X_n , where each X_i is an independent draw from a distribution with mean μ and variance $\sigma^2 < \infty$.

Two different estimators are proposed to estimate the mean μ as the sample size n grows:

- (a) **Estimator A (The Sample Mean):** $\hat{\mu}_{A,n} = \frac{1}{n} \sum_{i=1}^n X_i$
- (b) **Estimator B (A "Weighted" Mean):** $\hat{\mu}_{B,n} = \frac{1}{n} \sum_{i=1}^n \left(1 + \frac{1}{i}\right) X_i$

Which estimator, if any, is a consistent estimator for μ ? Justify your answer by checking if the limit of its Mean Squared Error (MSE) converges to 0.

Solution: We will analyze the consistency of each estimator by calculating the limit of its MSE. Recall that $\text{MSE}(\hat{\Theta}) = \text{Var}(\hat{\Theta}) + (\text{Bias}(\hat{\Theta}))^2$. We also use the facts that $E[X_i] = \mu$ and $\text{Var}(X_i) = \sigma^2$.

1. Analysis of Estimator A (Sample Mean)

This is the standard proof for the sample mean.

Bias:

$$E[\hat{\mu}_{A,n}] = E\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n E[X_i] = \frac{1}{n}(n\mu) = \mu$$

$$\text{Bias}(\hat{\mu}_{A,n}) = E[\hat{\mu}_{A,n}] - \mu = \mu - \mu = 0 \quad (\text{Unbiased})$$

Variance:

$$\begin{aligned} \text{Var}(\hat{\mu}_{A,n}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \quad (\text{due to independence}) \\ &= \frac{1}{n^2}(n\sigma^2) = \frac{\sigma^2}{n} \end{aligned}$$

MSE Limit:

$$\begin{aligned} \text{MSE}(\hat{\mu}_{A,n}) &= \frac{\sigma^2}{n} + (0)^2 = \frac{\sigma^2}{n} \\ \lim_{n \rightarrow \infty} \text{MSE}(\hat{\mu}_{A,n}) &= \lim_{n \rightarrow \infty} \frac{\sigma^2}{n} = 0 \end{aligned}$$

Conclusion for A: Since its MSE converges to 0, Estimator A is a **consistent estimator** of μ .

2. Analysis of Estimator B (Weighted Mean)

Bias:

$$\begin{aligned} E[\hat{\mu}_{B,n}] &= E\left[\frac{1}{n} \sum_{i=1}^n \left(1 + \frac{1}{i}\right) X_i\right] = \frac{1}{n} \sum_{i=1}^n \left(1 + \frac{1}{i}\right) E[X_i] \\ &= \frac{1}{n} \sum_{i=1}^n \left(1 + \frac{1}{i}\right) \mu = \frac{\mu}{n} \sum_{i=1}^n \left(1 + \frac{1}{i}\right) \\ &= \frac{\mu}{n} \left(\sum_{i=1}^n 1 + \sum_{i=1}^n \frac{1}{i}\right) = \frac{\mu}{n}(n + H_n) \\ &= \mu \left(1 + \frac{H_n}{n}\right) \end{aligned}$$

where $H_n = \sum_{i=1}^n \frac{1}{i}$ is the n -th harmonic number.

Bias Limit: The bias is $\text{Bias}(\hat{\mu}_{B,n}) = E[\hat{\mu}_{B,n}] - \mu = \mu \left(\frac{H_n}{n} \right)$. As $n \rightarrow \infty$, $H_n \approx \ln(n) + \gamma$.

$$\lim_{n \rightarrow \infty} \text{Bias}(\hat{\mu}_{B,n}) = \lim_{n \rightarrow \infty} \mu \frac{\ln(n)}{n} = 0 \quad (\text{by L'Hôpital's rule})$$

This estimator is *asymptotically unbiased*, but it is biased for any finite n .

Variance:

$$\begin{aligned} \text{Var}(\hat{\mu}_{B,n}) &= \text{Var} \left(\frac{1}{n} \sum_{i=1}^n \left(1 + \frac{1}{i} \right) X_i \right) \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var} \left(\left(1 + \frac{1}{i} \right) X_i \right) \quad (\text{due to independence}) \\ &= \frac{1}{n^2} \sum_{i=1}^n \left(1 + \frac{1}{i} \right)^2 \text{Var}(X_i) = \frac{\sigma^2}{n^2} \sum_{i=1}^n \left(1 + \frac{2}{i} + \frac{1}{i^2} \right) \\ &= \frac{\sigma^2}{n^2} \left(\sum_{i=1}^n 1 + 2 \sum_{i=1}^n \frac{1}{i} + \sum_{i=1}^n \frac{1}{i^2} \right) \\ &= \frac{\sigma^2}{n^2} \left(n + 2H_n + \sum_{i=1}^n \frac{1}{i^2} \right) \end{aligned}$$

MSE Limit: As $n \rightarrow \infty$, we know:

- $H_n \approx \ln(n)$
- $\sum_{i=1}^{\infty} \frac{1}{i^2} = \frac{\pi^2}{6}$ (a finite constant)

The MSE is $\text{Var}(\hat{\mu}_{B,n}) + (\text{Bias}(\hat{\mu}_{B,n}))^2$. Let's look at the limit of the variance:

$$\lim_{n \rightarrow \infty} \text{Var}(\hat{\mu}_{B,n}) = \lim_{n \rightarrow \infty} \frac{\sigma^2}{n^2} \left(n + 2\ln(n) + \frac{\pi^2}{6} \right)$$

Distributing the $\frac{1}{n^2}$:

$$= \lim_{n \rightarrow \infty} \sigma^2 \left(\frac{1}{n} + \frac{2\ln(n)}{n^2} + \frac{\pi^2}{6n^2} \right)$$

All terms go to 0 as $n \rightarrow \infty$. Thus, $\lim_{n \rightarrow \infty} \text{Var}(\hat{\mu}_{B,n}) = 0$.

Since $\lim_{n \rightarrow \infty} \text{Var}(\hat{\mu}_{B,n}) = 0$ and $\lim_{n \rightarrow \infty} \text{Bias}(\hat{\mu}_{B,n}) = 0$, we have:

$$\lim_{n \rightarrow \infty} \text{MSE}(\hat{\mu}_{B,n}) = \lim_{n \rightarrow \infty} \text{Var}(\hat{\mu}_{B,n}) + \lim_{n \rightarrow \infty} (\text{Bias}(\hat{\mu}_{B,n}))^2 = 0 + 0^2 = 0$$

Conclusion for B: Since its MSE also converges to 0, Estimator B is also a **consistent estimator** of μ .

Final Answer

Both Estimator A and Estimator B are consistent estimators for μ , as the MSE of both estimators converges to 0 as $n \rightarrow \infty$.

(Note: Although both are consistent, Estimator A converges much faster ($\text{MSE} = O(1/n)$) while Estimator B converges much slower ($\text{MSE} \approx O(\ln(n)/n^2)$ for the variance and $O((\ln(n)/n)^2)$ for the bias squared). Therefore, in practice, Estimator A is far superior.)

- Q4: (For conceptual understanding) Often when working with maximum likelihood functions, out of ease we maximize the log-likelihood rather than the likelihood to find the maximum likelihood estimator. Why is maximizing $L(\mathbf{x}; \theta)$ as a function of θ equivalent to maximizing $\log L(\mathbf{x}; \theta)$?

Solution:

Maximizing $L(\mathbf{x}; \theta)$ is equivalent to maximizing $\log L(\mathbf{x}; \theta)$ because the **natural logarithm**, $f(x) = \log(x)$, is a **strictly increasing function** (also called a monotonically increasing function).

(A function $f(x)$ is strictly increasing if for any two values a and b , if $a > b$, then $f(a) > f(b)$)

Detailed proof:

Assume that maximising the likelihood is not equivalent to maximising the log-likelihood. This would imply that there exist two parameter values θ_1 and θ_2 , such that:

$$L(\mathbf{x}; \theta_1) > L(\mathbf{x}; \theta_2)$$

and

$$\log L(\mathbf{x}; \theta_1) \leq \log L(\mathbf{x}; \theta_2)$$

But since the logarithm is strictly increasing, this leads to a contradiction. Therefore, our assumption must be false, and maximizing the likelihood is indeed equivalent to maximizing the log-likelihood. i.e $\theta_1 = \theta_2$.

- Q5: Let X be one observation from a $N(0, \sigma^2)$ distribution.

- Find an unbiased estimator of σ^2 .
- Find the log likelihood, $\log(L(x; \sigma^2))$, using

$$f_X(x; \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left\{-\frac{x^2}{2\sigma^2}\right\}$$

as the PDF.

- Find the Maximum Likelihood Estimate (MLE) for the standard deviation σ , $\hat{\sigma}_{ML}$.

Solution:

a. Find an unbiased estimator of σ^2

An estimator $T(X)$ is unbiased for σ^2 if $E[T(X)] = \sigma^2$.

We are given that $X \sim N(0, \sigma^2)$, which means:

- $E[X] = 0$
- $\text{Var}(X) = \sigma^2$

By the definition of variance:

$$\text{Var}(X) = E[X^2] - (E[X])^2$$

Substituting the known values:

$$\sigma^2 = E[X^2] - (0)^2$$

$$\sigma^2 = E[X^2]$$

This shows that the expected value of the statistic X^2 is exactly σ^2 . Therefore, an unbiased estimator for σ^2 is $T(X) = X^2$.

b. Find the log likelihood, $\log(L(x; \sigma^2))$

Since we have only one observation $X = x$, the likelihood function L is equal to the PDF:

$$L(x; \sigma^2) = f_X(x; \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left\{-\frac{x^2}{2\sigma^2}\right\}$$

Let $\ell = \log L$ be the log-likelihood.

$$\begin{aligned} \ell(x; \sigma^2) &= \log\left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left\{-\frac{x^2}{2\sigma^2}\right\}\right) \\ \ell(x; \sigma^2) &= \log(1) - \log(\sqrt{2\pi}) - \log(\sigma) + \log\left(\exp\left\{-\frac{x^2}{2\sigma^2}\right\}\right) \\ \ell(x; \sigma^2) &= 0 - \frac{1}{2}\log(2\pi) - \log(\sigma) - \frac{x^2}{2\sigma^2} \end{aligned}$$

The question asks for the log-likelihood as a function of σ^2 . Let $\theta = \sigma^2$, which means $\sigma = \sqrt{\theta} = \theta^{1/2}$. We substitute this into the expression:

$$\begin{aligned} \ell(x; \theta) &= -\frac{1}{2}\log(2\pi) - \log(\theta^{1/2}) - \frac{x^2}{2\theta} \\ \ell(x; \theta) &= -\frac{1}{2}\log(2\pi) - \frac{1}{2}\log(\theta) - \frac{x^2}{2\theta} \end{aligned}$$

Substituting σ^2 back in for θ , we get the final answer:

$$\log L(x; \sigma^2) = -\frac{1}{2}\log(2\pi) - \frac{1}{2}\log(\sigma^2) - \frac{x^2}{2\sigma^2}$$

c. Find the MLE for the standard deviation σ , $\hat{\sigma}_{ML}$

To find the MLE, we maximize the log-likelihood function. It is easier to first find the MLE for the variance, $\theta = \sigma^2$, and then use the invariance property of MLEs.

Let $\ell(\theta) = -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\theta) - \frac{x^2}{2\theta}$. We take the derivative with respect to θ and set it to 0.

$$\begin{aligned}\frac{d\ell}{d\theta} &= \frac{d}{d\theta} \left[-\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\theta) - \frac{x^2}{2} \theta^{-1} \right] \\ \frac{d\ell}{d\theta} &= 0 - \frac{1}{2\theta} - \frac{x^2}{2} (-1)\theta^{-2} \\ \frac{d\ell}{d\theta} &= -\frac{1}{2\theta} + \frac{x^2}{2\theta^2}\end{aligned}$$

Set the derivative to 0 to find the maximum:

$$\begin{aligned}-\frac{1}{2\theta} + \frac{x^2}{2\theta^2} &= 0 \\ \frac{x^2}{2\theta^2} &= \frac{1}{2\theta} \\ x^2(2\theta) &= 2\theta^2 \\ 2\theta^2 - 2\theta x^2 &= 0 \\ 2\theta(\theta - x^2) &= 0\end{aligned}$$

Since $\theta = \sigma^2$ must be greater than 0, the only valid solution is $\theta = x^2$. So, the MLE for the variance σ^2 is:

$$\hat{\sigma}_{ML}^2 = X^2$$

By the **invariance property of MLEs**, if $\hat{\theta}$ is the MLE for θ , then the MLE for $g(\theta)$ is $g(\hat{\theta})$. Here, we want the MLE for the standard deviation σ . The function is $g(\theta) = \sqrt{\theta}$.

$$\begin{aligned}\hat{\sigma}_{ML} &= g(\hat{\sigma}_{ML}^2) = \sqrt{\hat{\sigma}_{ML}^2} \\ \hat{\sigma}_{ML} &= \sqrt{X^2} = |X|\end{aligned}$$

The Maximum Likelihood Estimate for σ is $|X|$.

Note: The invariance property of the Maximum Likelihood Estimator (MLE) states that if θ is the MLE of a parameter θ , then $g(\theta)$ is the MLE of the function $g(\theta)$

Q6: Consider the simple linear regression model $Y_i = \beta_0 + \beta_1 x_i + \epsilon_i$, where ϵ_i 's are independent $N(0, \sigma^2)$ random variables. The estimators are:

$$\hat{\beta}_1 = \frac{S_{xy}}{S_{xx}}, \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{x}$$

where

$$S_{xx} = \sum_{i=1}^n (x_i - \bar{x})^2, \quad S_{xy} = \sum_{i=1}^n (x_i - \bar{x})(Y_i - \bar{Y})$$

(a) Show that $\hat{\beta}_1$ is a normal random variable.

- (b) **Show that** $\mathbb{E}[\hat{\beta}_1] = \beta_1$.
(c) **Show that** $\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{S_{xx}}$.

Solution

- (a) **Show that** $\hat{\beta}_1$ **is a normal random variable.**

First, we simplify S_{xy} . Since $\sum(x_i - \bar{x}) = 0$:

$$S_{xy} = \sum(x_i - \bar{x})(Y_i - \bar{Y}) = \sum(x_i - \bar{x})Y_i - \bar{Y} \sum(x_i - \bar{x}) = \sum_{i=1}^n (x_i - \bar{x})Y_i$$

Now, substitute this into $\hat{\beta}_1$:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})Y_i}{S_{xx}} = \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{S_{xx}} \right) Y_i$$

Let $k_i = (x_i - \bar{x})/S_{xx}$. Since x_i , \bar{x} , and S_{xx} are treated as fixed constants, k_i is a constant.

$$\hat{\beta}_1 = \sum_{i=1}^n k_i Y_i$$

Each Y_i is an independent normal random variable. Since $\hat{\beta}_1$ is a linear combination of independent normal variables, $\hat{\beta}_1$ is also a normal random variable.

- (b) **Show that** $\mathbb{E}[\hat{\beta}_1] = \beta_1$.

Using the linear form from part (a) and the linearity of expectation:

$$\mathbb{E}[\hat{\beta}_1] = \mathbb{E} \left[\sum_{i=1}^n k_i Y_i \right] = \sum_{i=1}^n k_i \mathbb{E}[Y_i]$$

We know $\mathbb{E}[Y_i] = \mathbb{E}[\beta_0 + \beta_1 x_i + \epsilon_i] = \beta_0 + \beta_1 x_i + \mathbb{E}[\epsilon_i] = \beta_0 + \beta_1 x_i$.

$$\mathbb{E}[\hat{\beta}_1] = \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{S_{xx}} \right) (\beta_0 + \beta_1 x_i)$$

$$\mathbb{E}[\hat{\beta}_1] = \frac{1}{S_{xx}} \left[\sum(x_i - \bar{x})\beta_0 + \sum(x_i - \bar{x})\beta_1 x_i \right]$$

The first sum is $\beta_0 \sum(x_i - \bar{x}) = 0$. The second sum is $\beta_1 \sum(x_i - \bar{x})x_i = \beta_1 \sum(x_i - \bar{x})^2 = \beta_1 S_{xx}$.

$$\mathbb{E}[\hat{\beta}_1] = \frac{1}{S_{xx}} [0 + \beta_1 S_{xx}] = \beta_1$$

- (c) **Show that** $\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{S_{xx}}$.

Using the linear form $\hat{\beta}_1 = \sum k_i Y_i$. Since the Y_i 's are independent:

$$\text{Var}(\hat{\beta}_1) = \text{Var} \left(\sum_{i=1}^n k_i Y_i \right) = \sum_{i=1}^n k_i^2 \text{Var}(Y_i)$$

We are given $\text{Var}(Y_i) = \sigma^2$.

$$\text{Var}(\hat{\beta}_1) = \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{S_{xx}} \right)^2 \sigma^2 = \frac{\sigma^2}{(S_{xx})^2} \sum_{i=1}^n (x_i - \bar{x})^2$$

By definition, $\sum(x_i - \bar{x})^2 = S_{xx}$.

$$\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{(S_{xx})^2} \cdot S_{xx} = \frac{\sigma^2}{S_{xx}}$$

Q7: Again consider the simple linear regression model from Problem 1.

- (a) **Show that $\hat{\beta}_0$ is a normal random variable.**
- (b) **Show that $\mathbb{E}[\hat{\beta}_0] = \beta_0$.**
- (c) **Show that $\text{Cov}(\hat{\beta}_1, Y_i) = \frac{x_i - \bar{x}}{S_{xx}} \sigma^2$.**
- (d) **Show that $\text{Cov}(\hat{\beta}_1, \bar{Y}) = 0$.**
- (e) **Show that $\text{Var}(\hat{\beta}_0) = \frac{\sum_{i=1}^n x_i^2}{n S_{xx}} \sigma^2$.**

Solution

- (a) **Show that $\hat{\beta}_0$ is a normal random variable.**

We have $\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{x}$. Substitute the linear forms for \bar{Y} and $\hat{\beta}_1$:

$$\begin{aligned} \hat{\beta}_0 &= \left(\frac{1}{n} \sum Y_i \right) - \left(\sum k_i Y_i \right) \bar{x} \quad \text{where } k_i = \frac{x_i - \bar{x}}{S_{xx}} \\ \hat{\beta}_0 &= \sum \frac{1}{n} Y_i - \sum (k_i \bar{x}) Y_i = \sum_{i=1}^n \left(\frac{1}{n} - k_i \bar{x} \right) Y_i \\ \hat{\beta}_0 &= \sum_{i=1}^n \left(\frac{1}{n} - \frac{\bar{x}(x_i - \bar{x})}{S_{xx}} \right) Y_i \end{aligned}$$

Let $c_i = \left(\frac{1}{n} - \frac{\bar{x}(x_i - \bar{x})}{S_{xx}} \right)$. Since c_i is a constant, $\hat{\beta}_0$ is a linear combination of the independent normal variables Y_i . Therefore, $\hat{\beta}_0$ is also a normal random variable.

- (b) **Show that $\mathbb{E}[\hat{\beta}_0] = \beta_0$.**

By linearity of expectation:

$$\mathbb{E}[\hat{\beta}_0] = \mathbb{E}[\bar{Y} - \hat{\beta}_1 \bar{x}] = \mathbb{E}[\bar{Y}] - \mathbb{E}[\hat{\beta}_1] \bar{x}$$

From Problem 1(b), we know $\mathbb{E}[\hat{\beta}_1] = \beta_1$. We find $\mathbb{E}[\bar{Y}]$:

$$\begin{aligned} \mathbb{E}[\bar{Y}] &= \mathbb{E} \left[\frac{1}{n} \sum Y_i \right] = \frac{1}{n} \sum \mathbb{E}[Y_i] = \frac{1}{n} \sum (\beta_0 + \beta_1 x_i) \\ \mathbb{E}[\bar{Y}] &= \frac{1}{n} (n\beta_0 + \beta_1 \sum x_i) = \beta_0 + \beta_1 \left(\frac{\sum x_i}{n} \right) = \beta_0 + \beta_1 \bar{x} \end{aligned}$$

Substitute these results back:

$$\mathbb{E}[\hat{\beta}_0] = (\beta_0 + \beta_1 \bar{x}) - (\beta_1) \bar{x} = \beta_0$$

(c) **Show that** $\text{Cov}(\hat{\beta}_1, Y_i) = \frac{x_i - \bar{x}}{S_{xx}}\sigma^2$.

Using the linear form $\hat{\beta}_1 = \sum_{j=1}^n k_j Y_j$ (with index j):

$$\text{Cov}(\hat{\beta}_1, Y_i) = \text{Cov}\left(\sum_{j=1}^n k_j Y_j, Y_i\right) = \sum_{j=1}^n k_j \text{Cov}(Y_j, Y_i)$$

Since the Y_j 's are independent, $\text{Cov}(Y_j, Y_i) = 0$ for $j \neq i$. The only non-zero term is when $j = i$:

$$\text{Cov}(Y_i, Y_i) = \text{Var}(Y_i) = \sigma^2$$

The sum collapses to the $j = i$ term:

$$\text{Cov}(\hat{\beta}_1, Y_i) = k_i \text{Var}(Y_i) = k_i \sigma^2 = \frac{x_i - \bar{x}}{S_{xx}} \sigma^2$$

(d) **Show that** $\text{Cov}(\hat{\beta}_1, \bar{Y}) = 0$.

Using the result from part (c):

$$\begin{aligned} \text{Cov}(\hat{\beta}_1, \bar{Y}) &= \text{Cov}\left(\hat{\beta}_1, \frac{1}{n} \sum_{i=1}^n Y_i\right) = \frac{1}{n} \sum_{i=1}^n \text{Cov}(\hat{\beta}_1, Y_i) \\ \text{Cov}(\hat{\beta}_1, \bar{Y}) &= \frac{1}{n} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{S_{xx}} \sigma^2 \right) \end{aligned}$$

Pull the constants out of the sum:

$$\text{Cov}(\hat{\beta}_1, \bar{Y}) = \frac{\sigma^2}{n S_{xx}} \sum_{i=1}^n (x_i - \bar{x})$$

Since $\sum(x_i - \bar{x}) = 0$, the entire expression is zero.

$$\text{Cov}(\hat{\beta}_1, \bar{Y}) = 0$$

(e) **Show that** $\text{Var}(\hat{\beta}_0) = \frac{\sum_{i=1}^n x_i^2}{n S_{xx}} \sigma^2$.

Use the variance formula for a difference:

$$\text{Var}(\hat{\beta}_0) = \text{Var}(\bar{Y} - \hat{\beta}_1 \bar{x}) = \text{Var}(\bar{Y}) + \text{Var}(\hat{\beta}_1 \bar{x}) - 2\text{Cov}(\bar{Y}, \hat{\beta}_1 \bar{x})$$

Since \bar{x} is a constant:

$$\text{Var}(\hat{\beta}_0) = \text{Var}(\bar{Y}) + \bar{x}^2 \text{Var}(\hat{\beta}_1) - 2\bar{x} \text{Cov}(\bar{Y}, \hat{\beta}_1)$$

From part (d), $\text{Cov}(\bar{Y}, \hat{\beta}_1) = 0$, so the last term vanishes.

$$\text{Var}(\hat{\beta}_0) = \text{Var}(\bar{Y}) + \bar{x}^2 \text{Var}(\hat{\beta}_1)$$

We need $\text{Var}(\bar{Y})$ and $\text{Var}(\hat{\beta}_1)$.

- From Problem 1(c): $\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{S_{xx}}$.
- $\text{Var}(\bar{Y}) = \text{Var}\left(\frac{1}{n} \sum Y_i\right) = \frac{1}{n^2} \sum \text{Var}(Y_i) = \frac{1}{n^2} (n\sigma^2) = \frac{\sigma^2}{n}$.

Substitute these in:

$$\text{Var}(\hat{\beta}_0) = \frac{\sigma^2}{n} + \bar{x}^2 \left(\frac{\sigma^2}{S_{xx}} \right) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right)$$

Find a common denominator:

$$\text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{S_{xx} + n\bar{x}^2}{nS_{xx}} \right)$$

Recall the identity $S_{xx} = \sum(x_i - \bar{x})^2 = \sum x_i^2 - n\bar{x}^2$. This implies $S_{xx} + n\bar{x}^2 = \sum x_i^2$. Substitute this into the numerator:

$$\text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{\sum_{i=1}^n x_i^2}{nS_{xx}} \right) = \frac{\sum_{i=1}^n x_i^2}{nS_{xx}} \sigma^2$$

Q8: Let $X_1, X_2, X_3, \dots, X_n$ be a random sample with unknown mean $\mathbb{E}[X_i] = \mu$, and unknown variance $\text{Var}(X_i) = \sigma^2$. Suppose that we would like to estimate $\theta = \mu^2$. We define the estimator $\hat{\Theta}$ as

$$\hat{\Theta} = (\bar{X})^2 = \left[\frac{1}{n} \sum_{k=1}^n X_k \right]^2$$

to estimate θ . Is $\hat{\Theta}$ an unbiased estimator of θ ? Why?

Solution An estimator is unbiased if $E[\hat{\Theta}] = \theta$. We must calculate $E[\hat{\Theta}] = E[(\bar{X})^2]$.

Recall the relationship for any random variable Y : $E[Y^2] = \text{Var}(Y) + (E[Y])^2$. Let's apply this to $Y = \bar{X}$:

$$E[\hat{\Theta}] = E[(\bar{X})^2] = \text{Var}(\bar{X}) + (E[\bar{X}])^2$$

We need to find the two components:

- $E[\bar{X}] = E \left[\frac{1}{n} \sum X_i \right] = \frac{1}{n} \sum E[X_i] = \frac{1}{n} \sum \mu = \frac{1}{n}(n\mu) = \mu$. So, $(E[\bar{X}])^2 = \mu^2$.
- $\text{Var}(\bar{X}) = \text{Var} \left(\frac{1}{n} \sum X_i \right) = \frac{1}{n^2} \sum \text{Var}(X_i)$ (since X_i are independent). $\text{Var}(\bar{X}) = \frac{1}{n^2} \sum \sigma^2 = \frac{1}{n^2}(n\sigma^2) = \frac{\sigma^2}{n}$.

Substitute these back into the equation:

$$E[\hat{\Theta}] = \frac{\sigma^2}{n} + \mu^2$$

The parameter we are estimating is $\theta = \mu^2$. Since $E[\hat{\Theta}] = \mu^2 + \frac{\sigma^2}{n} \neq \mu^2$ (unless $\sigma^2 = 0$), the estimator is biased.

Answer: No, $\hat{\Theta}$ is **not an unbiased estimator** of θ . Its bias is $E[\hat{\Theta}] - \theta = \frac{\sigma^2}{n}$.

Q9: **Estimating the parameter of a uniform random variable I.** We are given i.i.d. observations X_1, \dots, X_n that are uniformly distributed over the interval $[0, \theta]$. What is the ML estimator of θ ? Is it consistent? Is it unbiased or asymptotically unbiased?

Solution Let $X_{(n)} = \max(X_1, \dots, X_n)$.

- (1) **ML Estimator:** The PDF for a single observation is $f(x_i|\theta) = \frac{1}{\theta}$ for $0 \leq x_i \leq \theta$, and 0 otherwise. The joint likelihood function is:

$$L(\theta) = \prod_{i=1}^n f(x_i|\theta) = \left(\frac{1}{\theta}\right)^n$$

This likelihood is only non-zero if all $x_i \leq \theta$. This introduces the constraint that $\theta \geq x_i$ for all i , which means $\theta \geq \max(x_i) = x_{(n)}$.

We want to maximize $L(\theta) = \theta^{-n}$ subject to $\theta \geq x_{(n)}$. Since $L(\theta)$ is a decreasing function of θ , we make it as large as possible by choosing the smallest possible θ that satisfies the constraint. The smallest valid θ is $x_{(n)}$.

Therefore, the ML estimator is $\hat{\theta}_{ML} = X_{(n)}$.

- (2) **Consistency: Yes, it is consistent.** An estimator is consistent if it converges in probability to θ . We check if $\lim_{n \rightarrow \infty} P(|\hat{\theta}_{ML} - \theta| > \epsilon) = 0$ for any $\epsilon > 0$. Since $\hat{\theta}_{ML} = X_{(n)} \leq \theta$, this simplifies to $P(\theta - X_{(n)} > \epsilon) = P(X_{(n)} < \theta - \epsilon)$.

$$P(X_{(n)} < \theta - \epsilon) = P(X_1 < \theta - \epsilon, \dots, X_n < \theta - \epsilon)$$

Since they are i.i.d.:

$$= P(X_1 < \theta - \epsilon)^n = \left(\frac{\theta - \epsilon}{\theta}\right)^n = \left(1 - \frac{\epsilon}{\theta}\right)^n$$

As $n \rightarrow \infty$, $\left(1 - \frac{\epsilon}{\theta}\right)^n \rightarrow 0$ because $1 - \frac{\epsilon}{\theta} < 1$. Thus, the estimator is consistent.

- (3) **Bias: It is biased, but asymptotically unbiased.** We must find $E[\hat{\theta}_{ML}] = E[X_{(n)}]$. Let $Y = X_{(n)}$. The CDF of Y is:

$$F_Y(y) = P(Y \leq y) = P(X_1 \leq y, \dots, X_n \leq y) = P(X_1 \leq y)^n = \left(\frac{y}{\theta}\right)^n$$

The PDF of Y is the derivative: $f_Y(y) = F'_Y(y) = \frac{ny^{n-1}}{\theta^n}$ for $0 \leq y \leq \theta$. The expectation is:

$$\begin{aligned} E[Y] &= \int_0^\theta y f_Y(y) dy = \int_0^\theta y \left(\frac{ny^{n-1}}{\theta^n}\right) dy = \frac{n}{\theta^n} \int_0^\theta y^n dy \\ E[Y] &= \frac{n}{\theta^n} \left[\frac{y^{n+1}}{n+1} \right]_0^\theta = \frac{n}{\theta^n} \left(\frac{\theta^{n+1}}{n+1} \right) = \frac{n}{n+1} \theta \end{aligned}$$

- **Unbiased?** No. $E[\hat{\theta}_{ML}] = \frac{n}{n+1} \theta \neq \theta$.
- **Asymptotically Unbiased?** Yes. $\lim_{n \rightarrow \infty} E[\hat{\theta}_{ML}] = \lim_{n \rightarrow \infty} \frac{n}{n+1} \theta = \theta$.

- Q10: **Estimating the parameter of a uniform random variable II.** We are given i.i.d. observations X_1, \dots, X_n that are uniformly distributed over the interval $[\theta, \theta + 1]$. Find a ML estimator of θ . Is it consistent? Is it unbiased or asymptotically unbiased?

Solution Let $X_{(1)} = \min(X_1, \dots, X_n)$ and $X_{(n)} = \max(X_1, \dots, X_n)$.

(1) **ML Estimator:** The PDF for a single observation is $f(x_i|\theta) = \frac{1}{(\theta+1)-\theta} = 1$ for $\theta \leq x_i \leq \theta + 1$. The joint likelihood is $L(\theta) = \prod_{i=1}^n 1 = 1$, but only if all observations satisfy the constraints. The constraints are:

- $x_i \geq \theta$ for all $i \implies \theta \leq \min(x_i) \implies \theta \leq x_{(1)}$.
- $x_i \leq \theta + 1$ for all $i \implies \theta \geq \max(x_i) - 1 \implies \theta \geq x_{(n)} - 1$.

So, $L(\theta) = 1$ if $X_{(n)} - 1 \leq \theta \leq X_{(1)}$, and $L(\theta) = 0$ otherwise.

The likelihood is 1 for any θ in the interval $[X_{(n)} - 1, X_{(1)}]$. This means the ML estimator is **not unique**. Any value in this interval maximizes the likelihood.

We can choose any value, for example an endpoint. Let's analyze $\hat{\theta}_{ML} = X_{(1)}$. (Another valid choice is $\hat{\theta}_{ML} = X_{(n)} - 1$).

(2) **Consistency:** (Analyzing $\hat{\theta} = X_{(1)}$) **Yes, it is consistent.** We check if $\lim_{n \rightarrow \infty} P(|X_{(1)} - \theta| > \epsilon) = 0$. Since $X_{(1)} \geq \theta$, this simplifies to $P(X_{(1)} - \theta > \epsilon) = P(X_{(1)} > \theta + \epsilon)$.

$$\begin{aligned} P(X_{(1)} > \theta + \epsilon) &= P(X_1 > \theta + \epsilon, \dots, X_n > \theta + \epsilon) \\ &= P(X_1 > \theta + \epsilon)^n \end{aligned}$$

The probability for a single X_i is $P(X_i > \theta + \epsilon) = \frac{(\theta+1)-(\theta+\epsilon)}{(\theta+1)-\theta} = \frac{1-\epsilon}{1} = 1 - \epsilon$.

$$P(X_{(1)} > \theta + \epsilon) = (1 - \epsilon)^n$$

As $n \rightarrow \infty$, $(1 - \epsilon)^n \rightarrow 0$ because $1 - \epsilon < 1$. Thus, the estimator is consistent.

(3) **Bias:** (Analyzing $\hat{\theta} = X_{(1)}$) **It is biased, but asymptotically unbiased.** We need $E[X_{(1)}]$. Let $Z = X_{(1)}$. The CDF of Z is:

$$F_Z(z) = P(Z \leq z) = 1 - P(Z > z) = 1 - P(X_1 > z, \dots, X_n > z)$$

$$F_Z(z) = 1 - P(X_1 > z)^n$$

For $z \in [\theta, \theta + 1]$, $P(X_1 > z) = \frac{(\theta+1)-z}{1} = 1 + \theta - z$.

$$F_Z(z) = 1 - (1 + \theta - z)^n$$

The PDF is $f_Z(z) = F'_Z(z) = 0 - n(1 + \theta - z)^{n-1} \cdot (-1) = n(1 + \theta - z)^{n-1}$. The expectation is:

$$E[Z] = \int_{\theta}^{\theta+1} z \cdot n(1 + \theta - z)^{n-1} dz$$

(Using substitution $u = 1 + \theta - z$, $du = -dz$, $z = 1 + \theta - u$)

$$E[Z] = \int_1^0 (1 + \theta - u)nu^{n-1}(-du) = n \int_0^1 (1 + \theta - u)u^{n-1} du$$

$$E[Z] = n \int_0^1 ((1 + \theta)u^{n-1} - u^n) du = n \left[(1 + \theta) \frac{u^n}{n} - \frac{u^{n+1}}{n+1} \right]_0^1$$

$$E[Z] = n \left(\frac{1 + \theta}{n} - \frac{1}{n+1} \right) = (1 + \theta) - \frac{n}{n+1} = \theta + 1 - \frac{n}{n+1} = \theta + \frac{1}{n+1}$$

- **Unbiased?** No. $E[X_{(1)}] = \theta + \frac{1}{n+1} \neq \theta$.

- **Asymptotically Unbiased?** Yes. $\lim_{n \rightarrow \infty} E[X_{(1)}] = \lim_{n \rightarrow \infty} (\theta + \frac{1}{n+1}) = \theta$.